

RABBITMQ IS MORE THAN
A SIDEKIQ REPLACEMENT



Stanko Krtalic Rusendic

 github.com/stankec

 [@monorkin](https://twitter.com/monorkin)

 hey@stanko.io

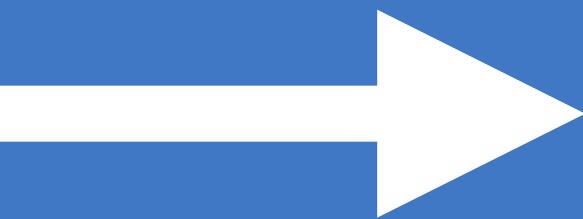
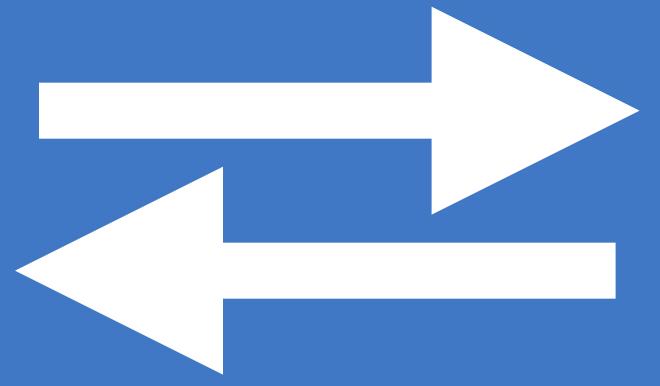
 stanko.io

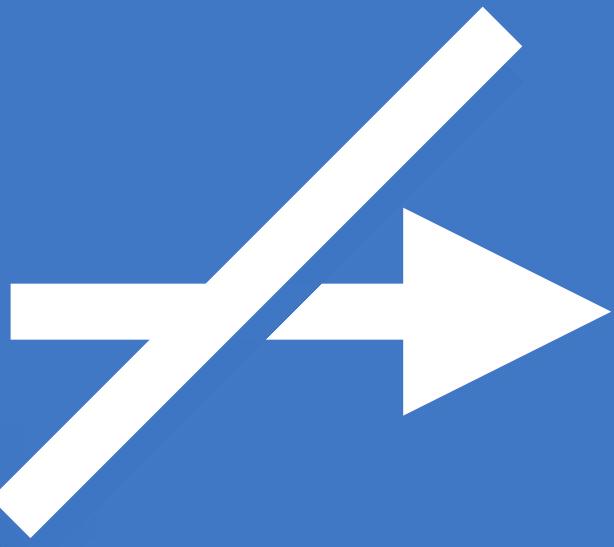
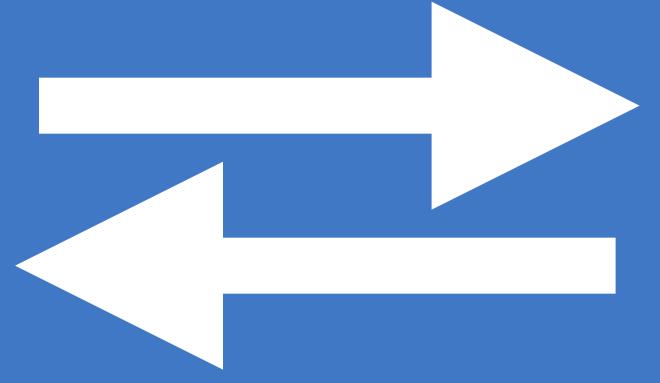


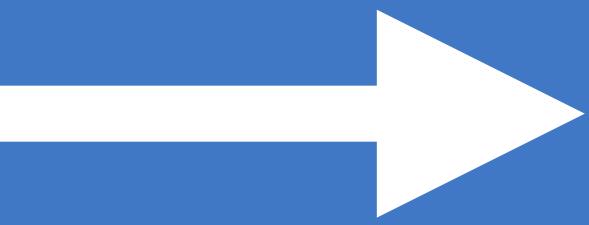
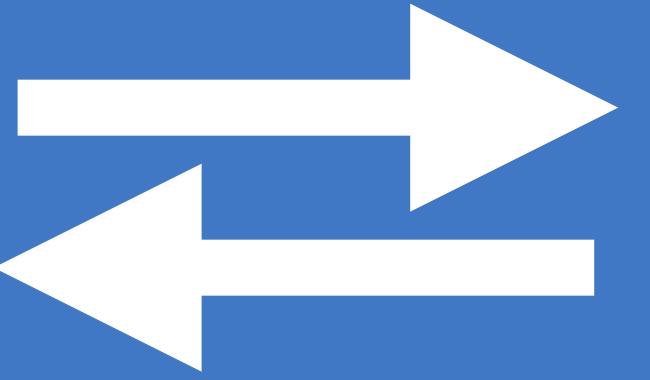
WORKERS

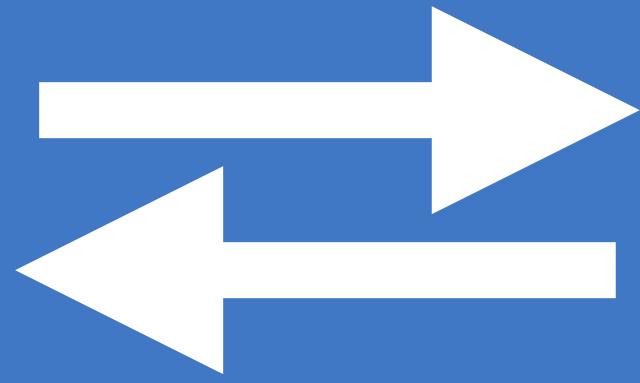


Asynchronous execution









???

```
3 require_relative 'chat'
4
5 class App < Roda
6   Chat.start_heartbeat
7
8   plugin :streaming
9   plugin :render, engine: 'slim'
10
11  route do |r|
12    r.root do
13      view('root', layout: false)
14    end
15
16    r.get 'stream' do
17      response['Content-Type'] = 'text/event-stream;charset=UTF-8'
18      user_queue = Chat.subscribe
19      callback = proc { Chat.unsubscribe(user_queue) }
20      stream(loop: true, callback: callback) do |out|
21        Chat.stream(out, user_queue)
22      end
23    end
24
25    r.on 'messages' do
26      r.post do
27        name = r.params['name']
28        message = r.params['message']
29        Chat.message(name, message)
30      end
31    end
32  end
33 end
```

→ demo (master) ✘

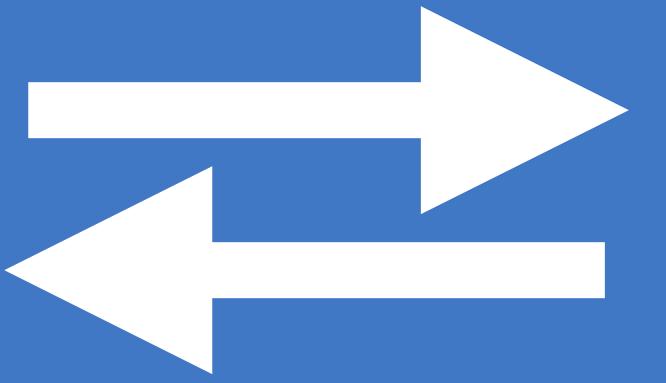
```
3 require_relative 'chat'
4
5 class App < Roda
6   JOBS = Queue.new
7   Thread.new do
8     loop do
9       JOBS.pop.call
10    end
11  end
12
13 Chat.start_heartbeat
14
15 plugin :streaming
16 plugin :render, engine: 'slim'
17
18 route do |r|
19   r.root do
20     view('root', layout: false)
21   end
22
23   r.get 'stream' do
24     response['Content-Type'] = 'text/event-stream;charset=UTF-8'
25     user_queue = Chat.subscribe
26     callback = proc { Chat.unsubscribe(user_queue) }
27     stream(loop: true, callback: callback) do |out|
28       Chat.stream(out, user_queue)
29     end
30   end
31
32   r.on 'messages' do
33     r.post do
34       name = r.params['name']
35       message = r.params['message']
36       if message.match?(%r{^\\broadcast\\s+.*$})
37         message = message.scan(%r{^\\broadcast\\s+(.*)$}).flatten.first
38         JOBS << proc do
39           Chat.broadcast(message)
40         end
41       else
42         Chat.message(name, message)
43       end
44     end
45   end
46 end
```

```
^Gracefully stopping ... (press Ctrl+C again to force)
Stopping demo_demo-01_1 ...
Killing demo_demo-01_1 ... done
→ demo (master) ✘
```

Downsides?

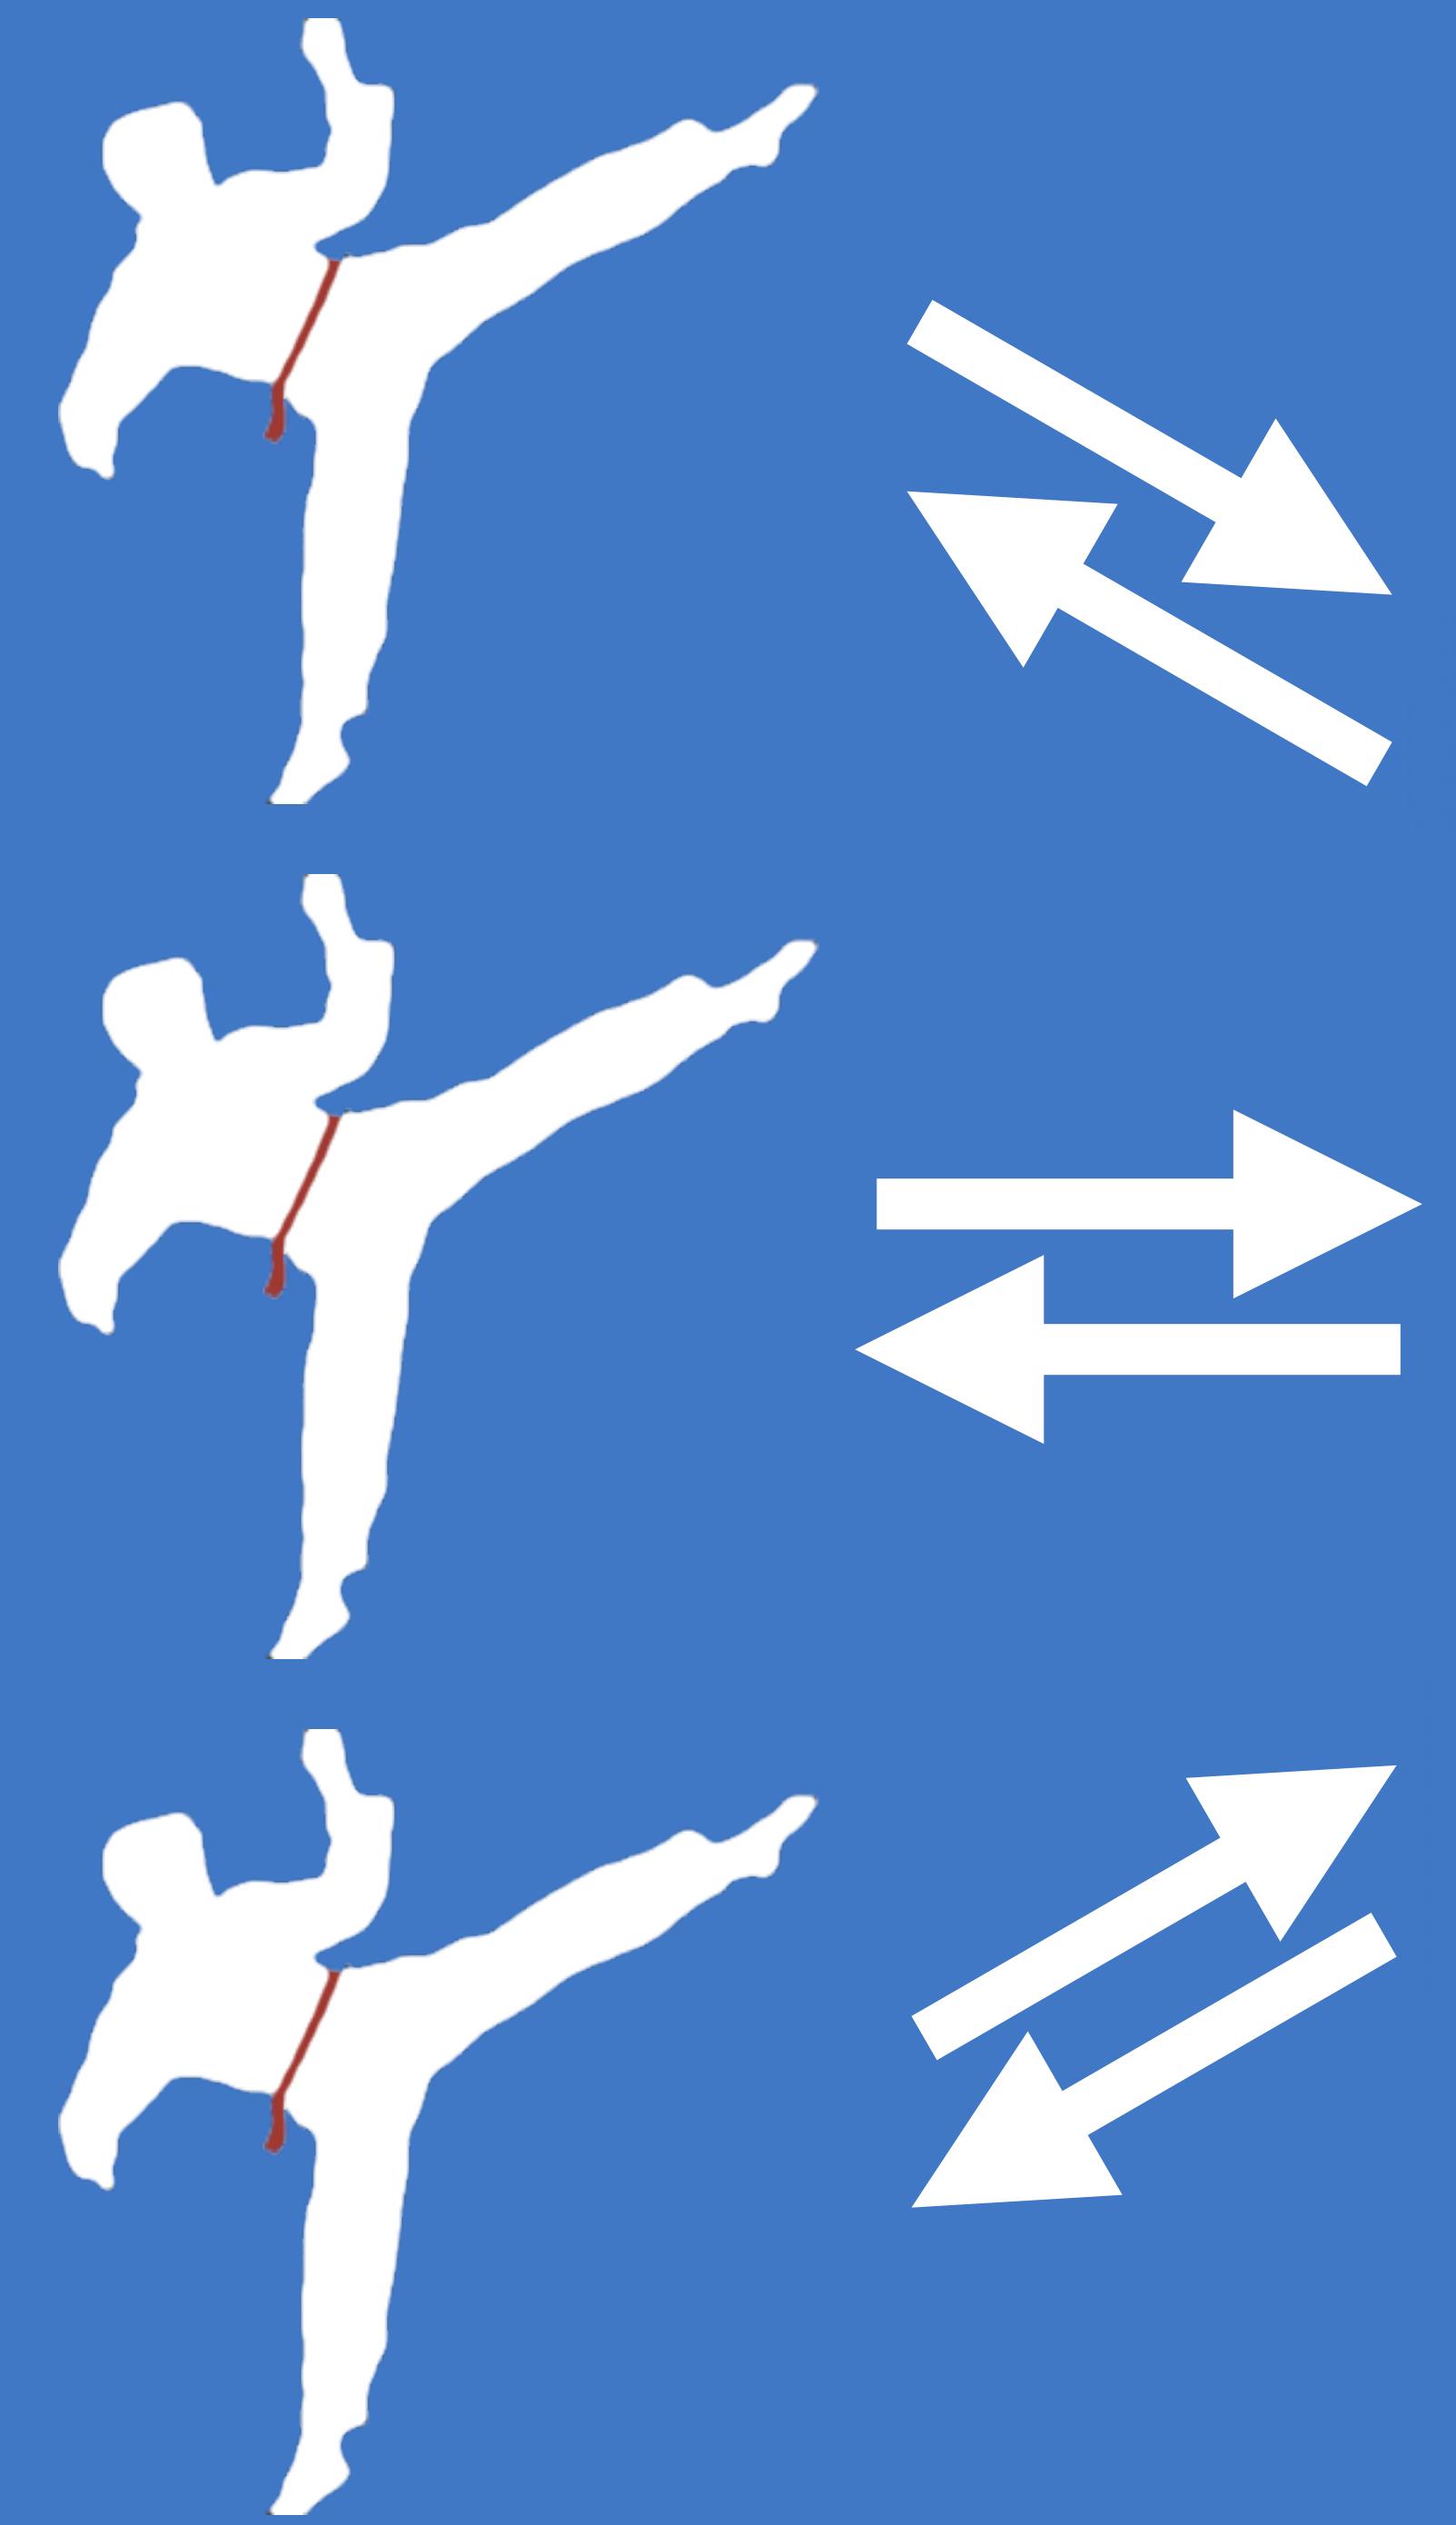
Debugging

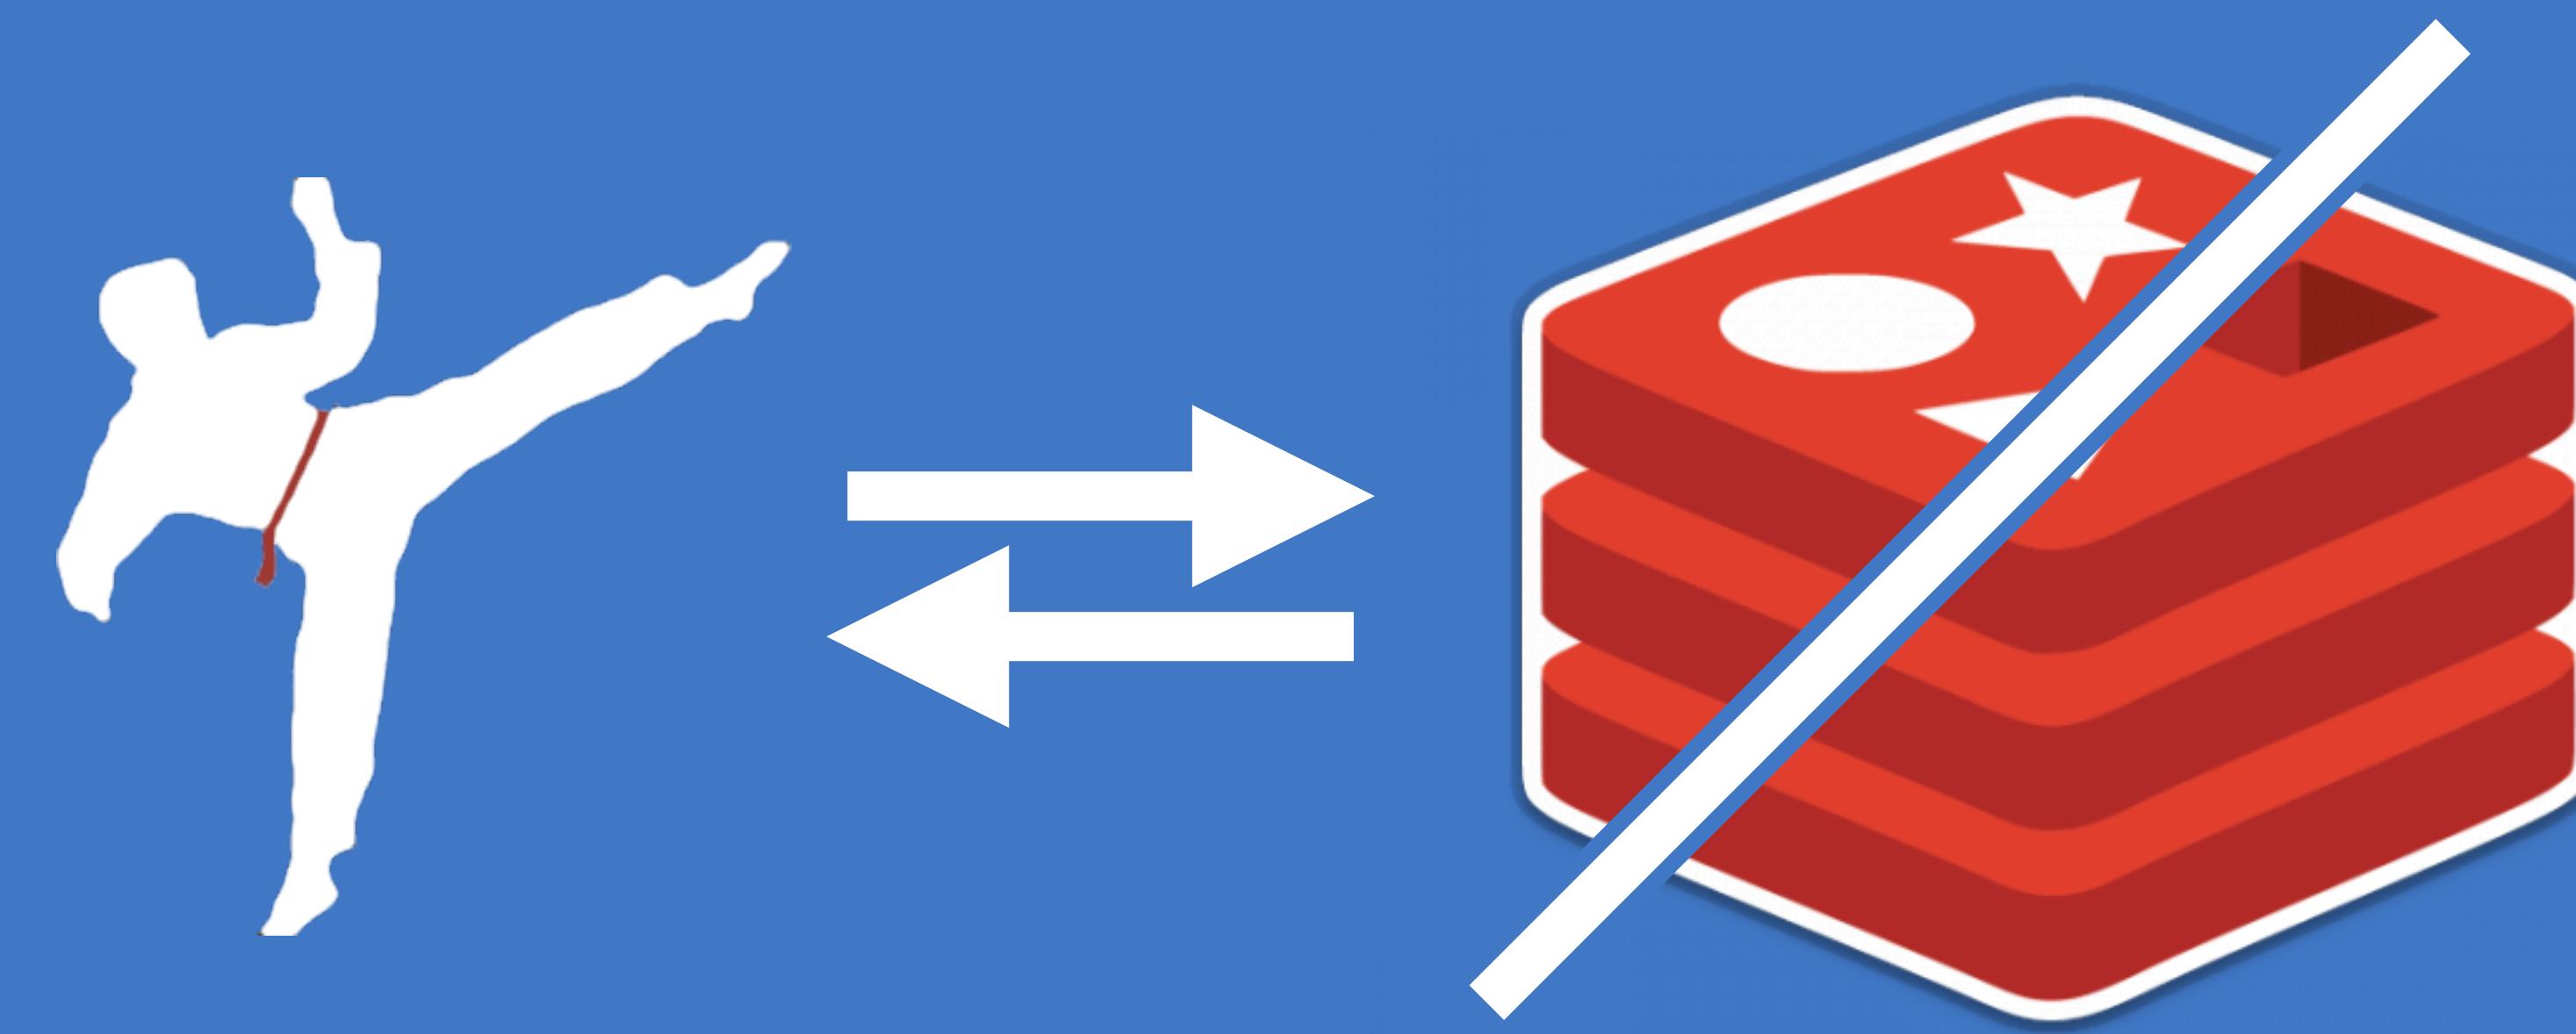
Load balancing



FloatingPoint ➤ 1 > papa ◀ 213.202.109.194 < 192.168.113.243 ◀ 12915/16384M < ⚡ 100% < 2018/02/26 < 22:26 :

```
127.0.0.1:6379>
127.0.0.1:6379> KEYS *
1) "stat:failed:push_notifs"
2) "dead"
3) "stat:failed"
4) "retry"
127.0.0.1:6379> ZRANGE retry 0 -1
1) "{\"retry_count\":10,\"retried_at\":1519677868,\"queue\":\"push_notifs\",\"max_retry_count\":10,\"jid\":\"37407568761197553900\",\"finished_at\":null,\"failed_at\":1519660887,\"error_message\":\"no match of right hand side value: nil\",\"error_backtrace\":\"(papa_pal) lib/papa_pal/workers/push_apns.ex:9: PapaPal.Workers.PushApns.perform/3\\n      (verk) lib/verk/worker.ex:35: Verk.Worker.handle_cast/2\\n      (stdlib) gen_server.erl:616: :gen_server.try_dispatch/4\\n      (stdlib) gen_server.erl:686: :gen_server.handle_msg/6\\n      (stdlib) proc_lib.erl:247: :proc_lib.init_p_do_apply/3\\n\",\"enqueued_at\":1519660886,\"class\":\"PapaPal.Workers.PushApns\",\"args\":[\"900185ef-5bc3-4481-ad0c-fb1ad4637d2a\"],\"Remember to complete your visit!\",\"{\\\\\"payload\\\\\":null}]}"
127.0.0.1:6379>
```



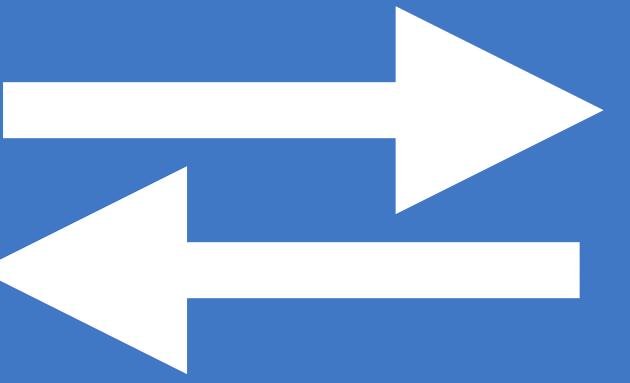
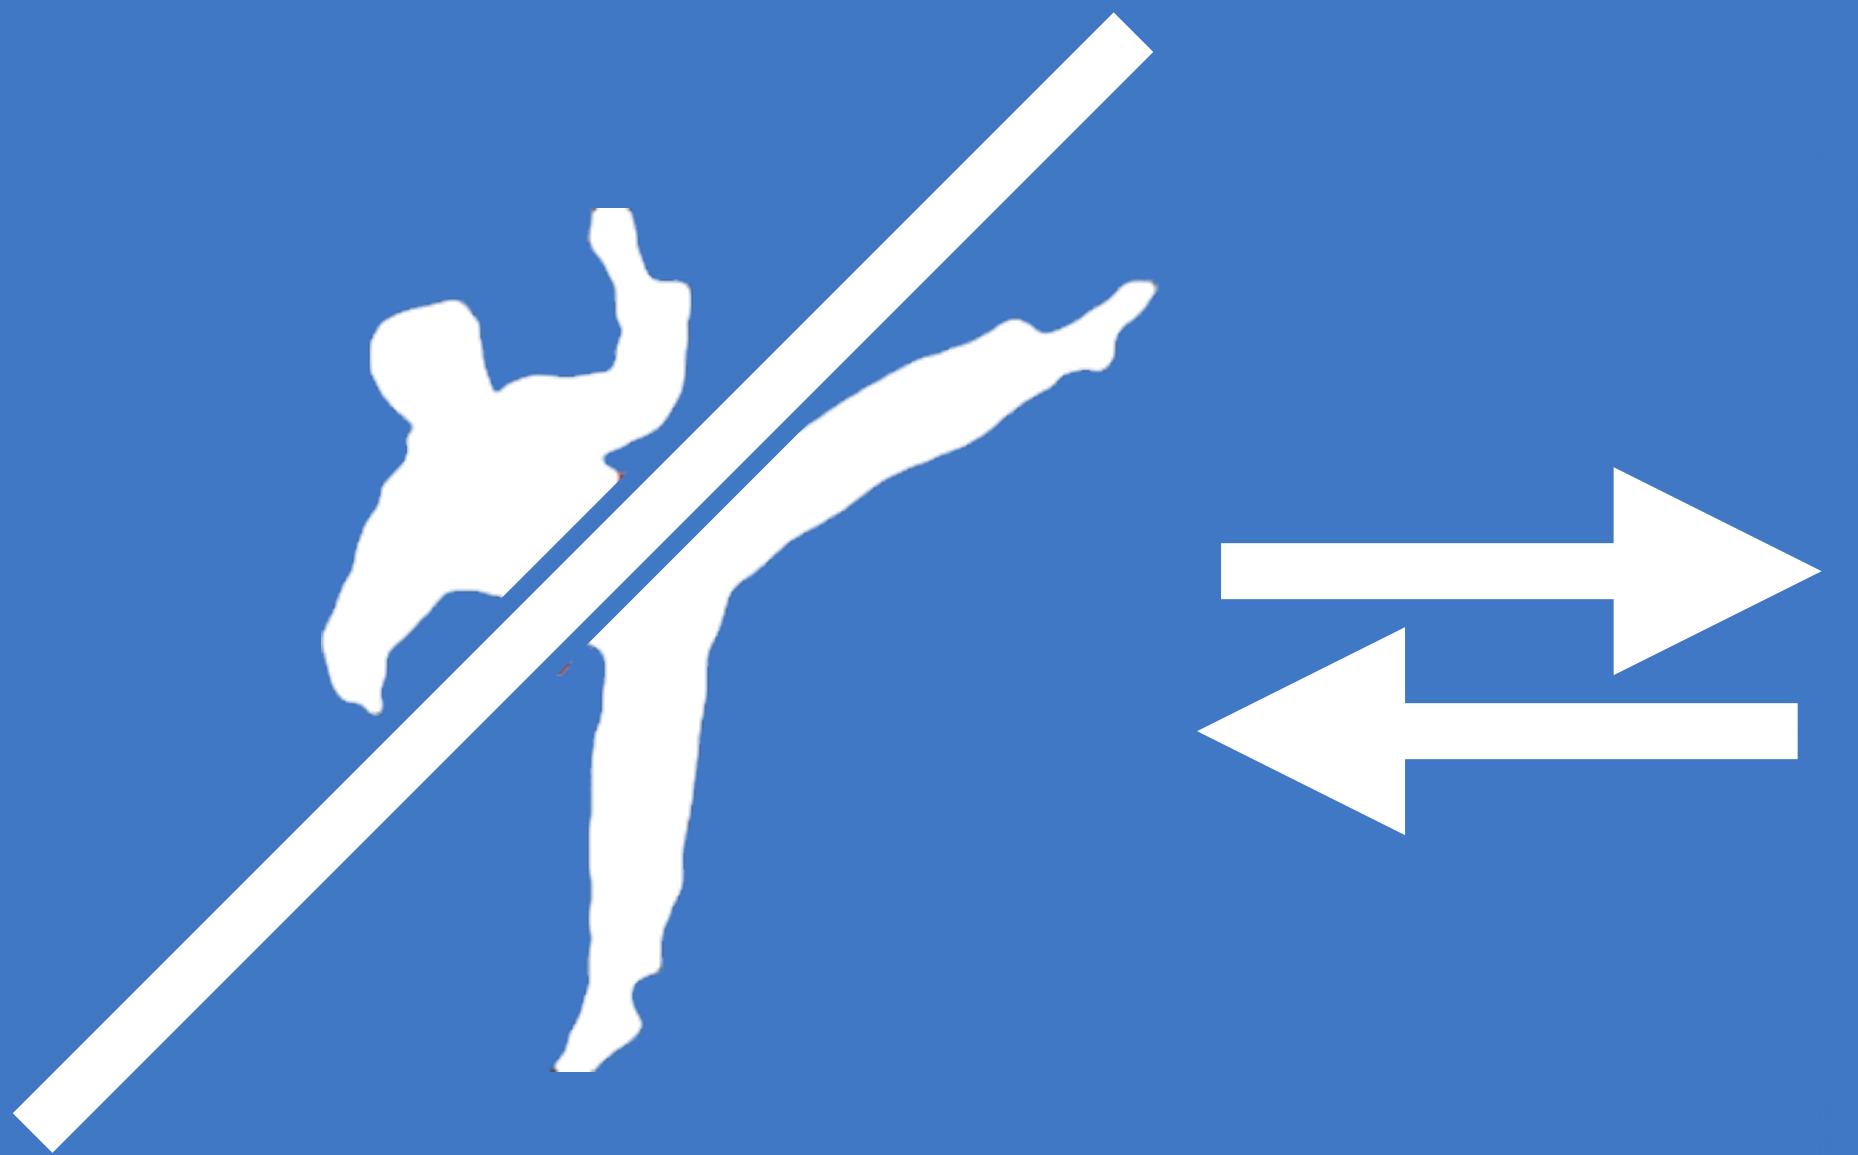


RDB disadvantages

- RDB is NOT good if you need to minimize the chance of data loss in case Redis stops working (for example after a power outage). You can configure different *save points* where an RDB is produced (for instance after at least five minutes and 100 writes against the data set, but you can have multiple save points). However you'll usually create an RDB snapshot every five minutes or more, so in case of Redis stopping working without a correct shutdown for any reason you should be prepared to lose the latest minutes of data.
- RDB needs to fork() often in order to persist on disk using a child process. Fork() can be time consuming if the dataset is big, and may result in Redis to stop serving clients for some millisecond or even for one second if the dataset is very big and the CPU performance not great. AOF also needs to fork() but you can tune how often you want to rewrite your logs without any trade-off on durability.

AOF disadvantages

- AOF files are usually bigger than the equivalent RDB files for the same dataset.
- AOF can be slower than RDB depending on the exact `fsync` policy. In general with `fsync` set to *every second* performances are still very high, and with `fsync` disabled it should be exactly as fast as RDB even under high load. Still RDB is able to provide more guarantees about the maximum latency even in the case of an huge write load.
- In the past we experienced rare bugs in specific commands (for instance there was one involving blocking commands like `BRPOPLPUSH`) causing the AOF produced to not reproduce exactly the same dataset on reloading. This bugs are rare and we have tests in the test suite creating random complex datasets automatically and reloading them to check everything is ok, but this kind of bugs are almost impossible with RDB persistence. To make this point more clear: the Redis AOF works incrementally updating an existing state, like MySQL or MongoDB does, while the RDB snapshotting creates everything from scratch again and again, that is conceptually more robust. However - 1) It should be noted that every time the AOF is rewritten by Redis it is recreated from scratch starting from the actual data contained in the data set, making resistance to bugs stronger compared to an always appending AOF file (or one rewritten reading the old AOF instead of reading the data in memory). 2) We never had a single report from users about an AOF corruption that was detected in the real world.



```
1 class TestWorker
2   include Sidekiq::Worker
3
4   def perform(*args)
5     # Do something
6   rescue
7     # Do something else
8   end
9 end
```

```
1 # frozen_string_literal: true
2
3 require_relative 'demo'
4 include Demo
5
6 begin
7   Demo.halt_and_catch_fire
8 rescue
9   puts 'All is good'
10 ensure
11   puts "The VM didn't crash this time"
12 end
```

```
~
```

Process Crashes

If the Sidekiq process segfaults or crashes the Ruby VM, any jobs that were being processed are lost. [Sidekiq Pro](#) offers a [reliable queueing](#) feature which does not lose those jobs.

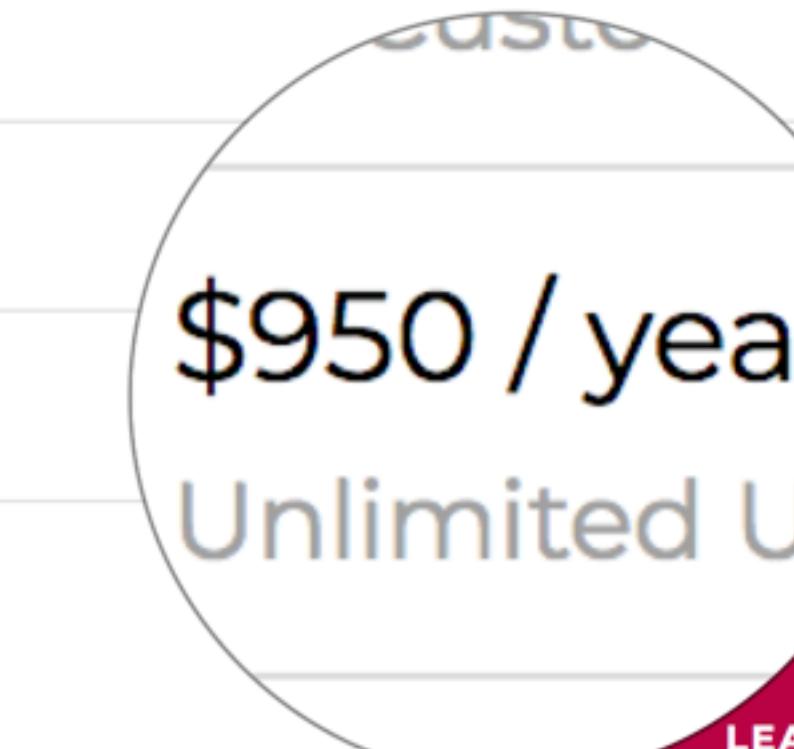
OSS

AUDIENCE	Hobbyists
DEDICATED SUPPORT	None
LICENSE	LGPL
PRICING	Free
PURCHASING	n/a

[GET STARTED ➔](#)

PRO

Small Business

[LEARN MORE ➔](#)[BUY ➔](#)

ENTERPRISE

Med-Large Companies

Email

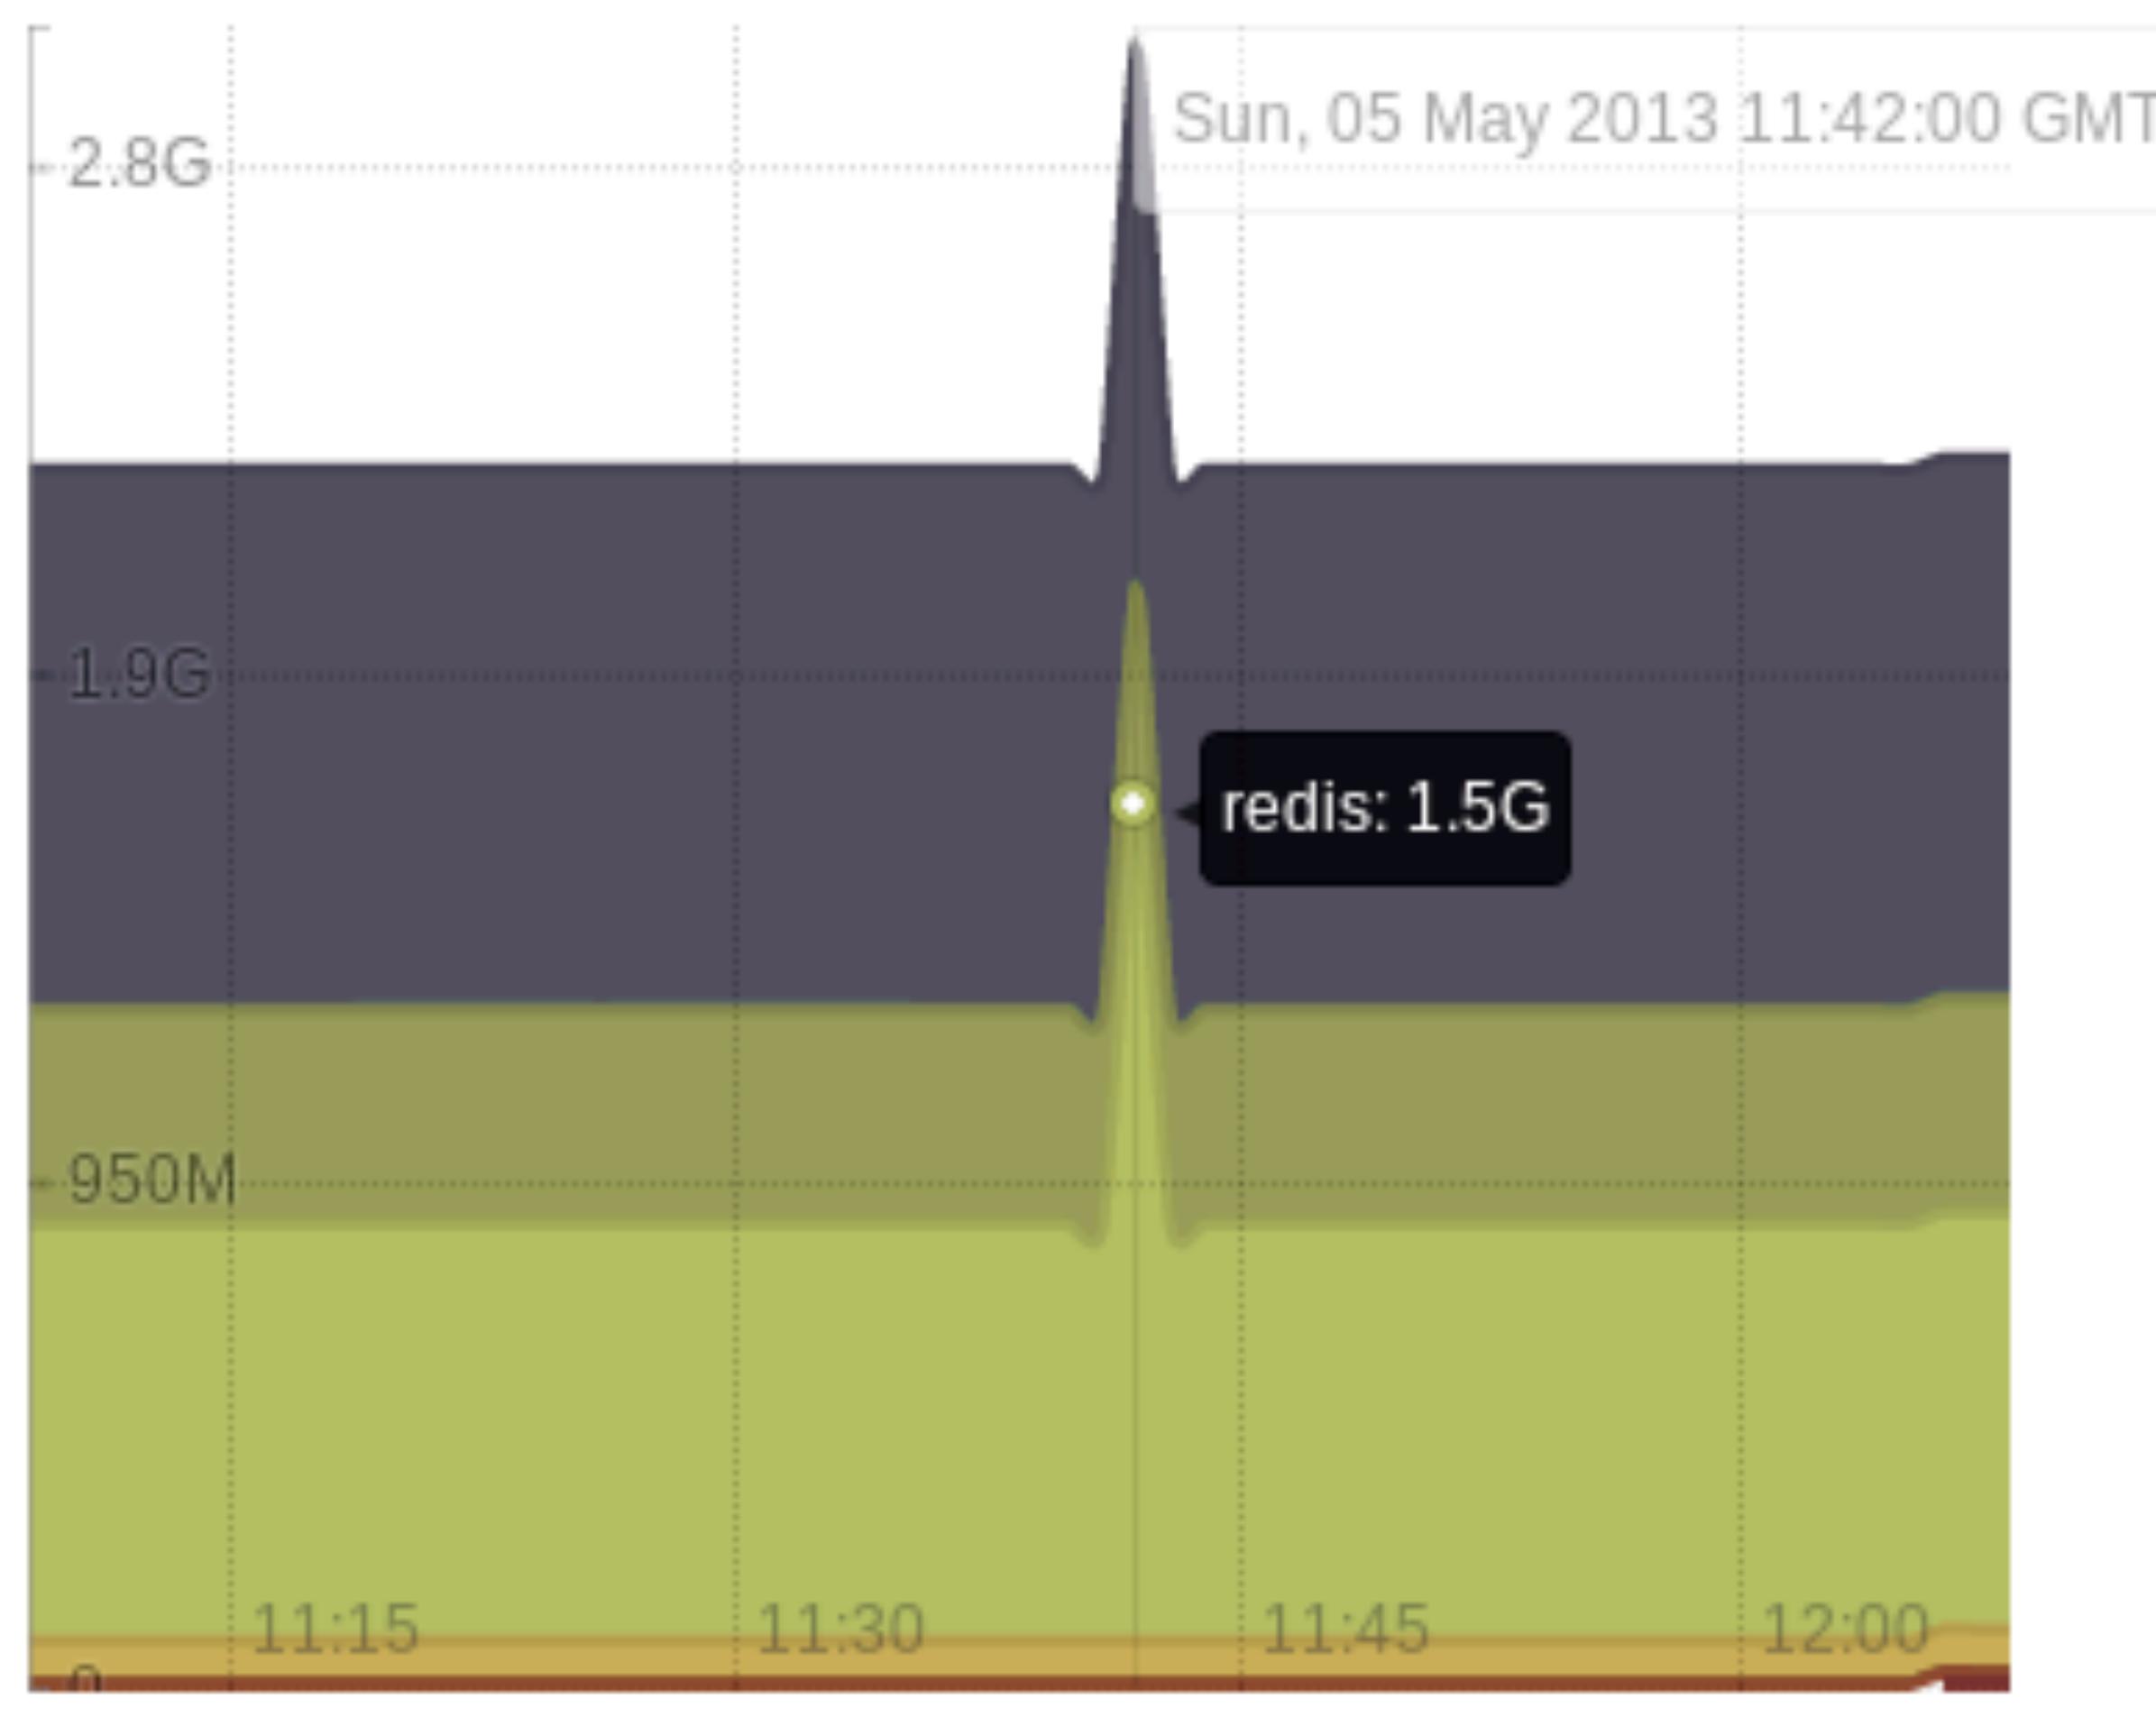
Commercial
with Custom Terms

\$1950 / year
per 100 Workers

Credit Card, Invoice

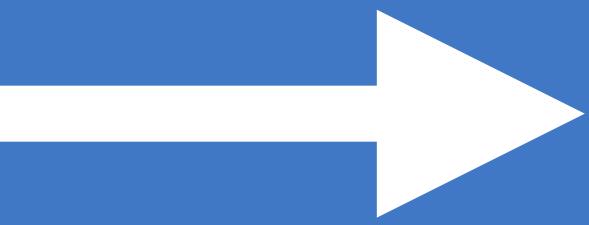
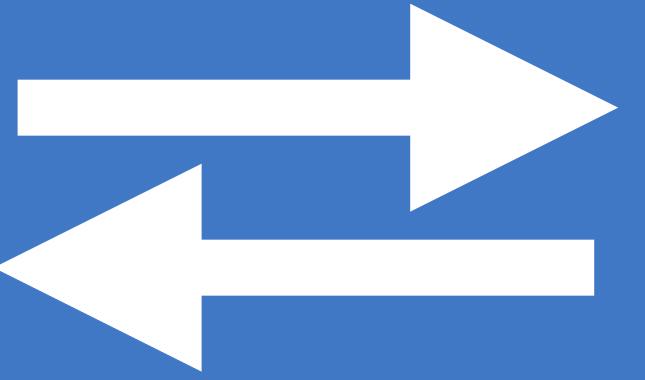
[LEARN MORE ➔](#)[BUY ➔](#)

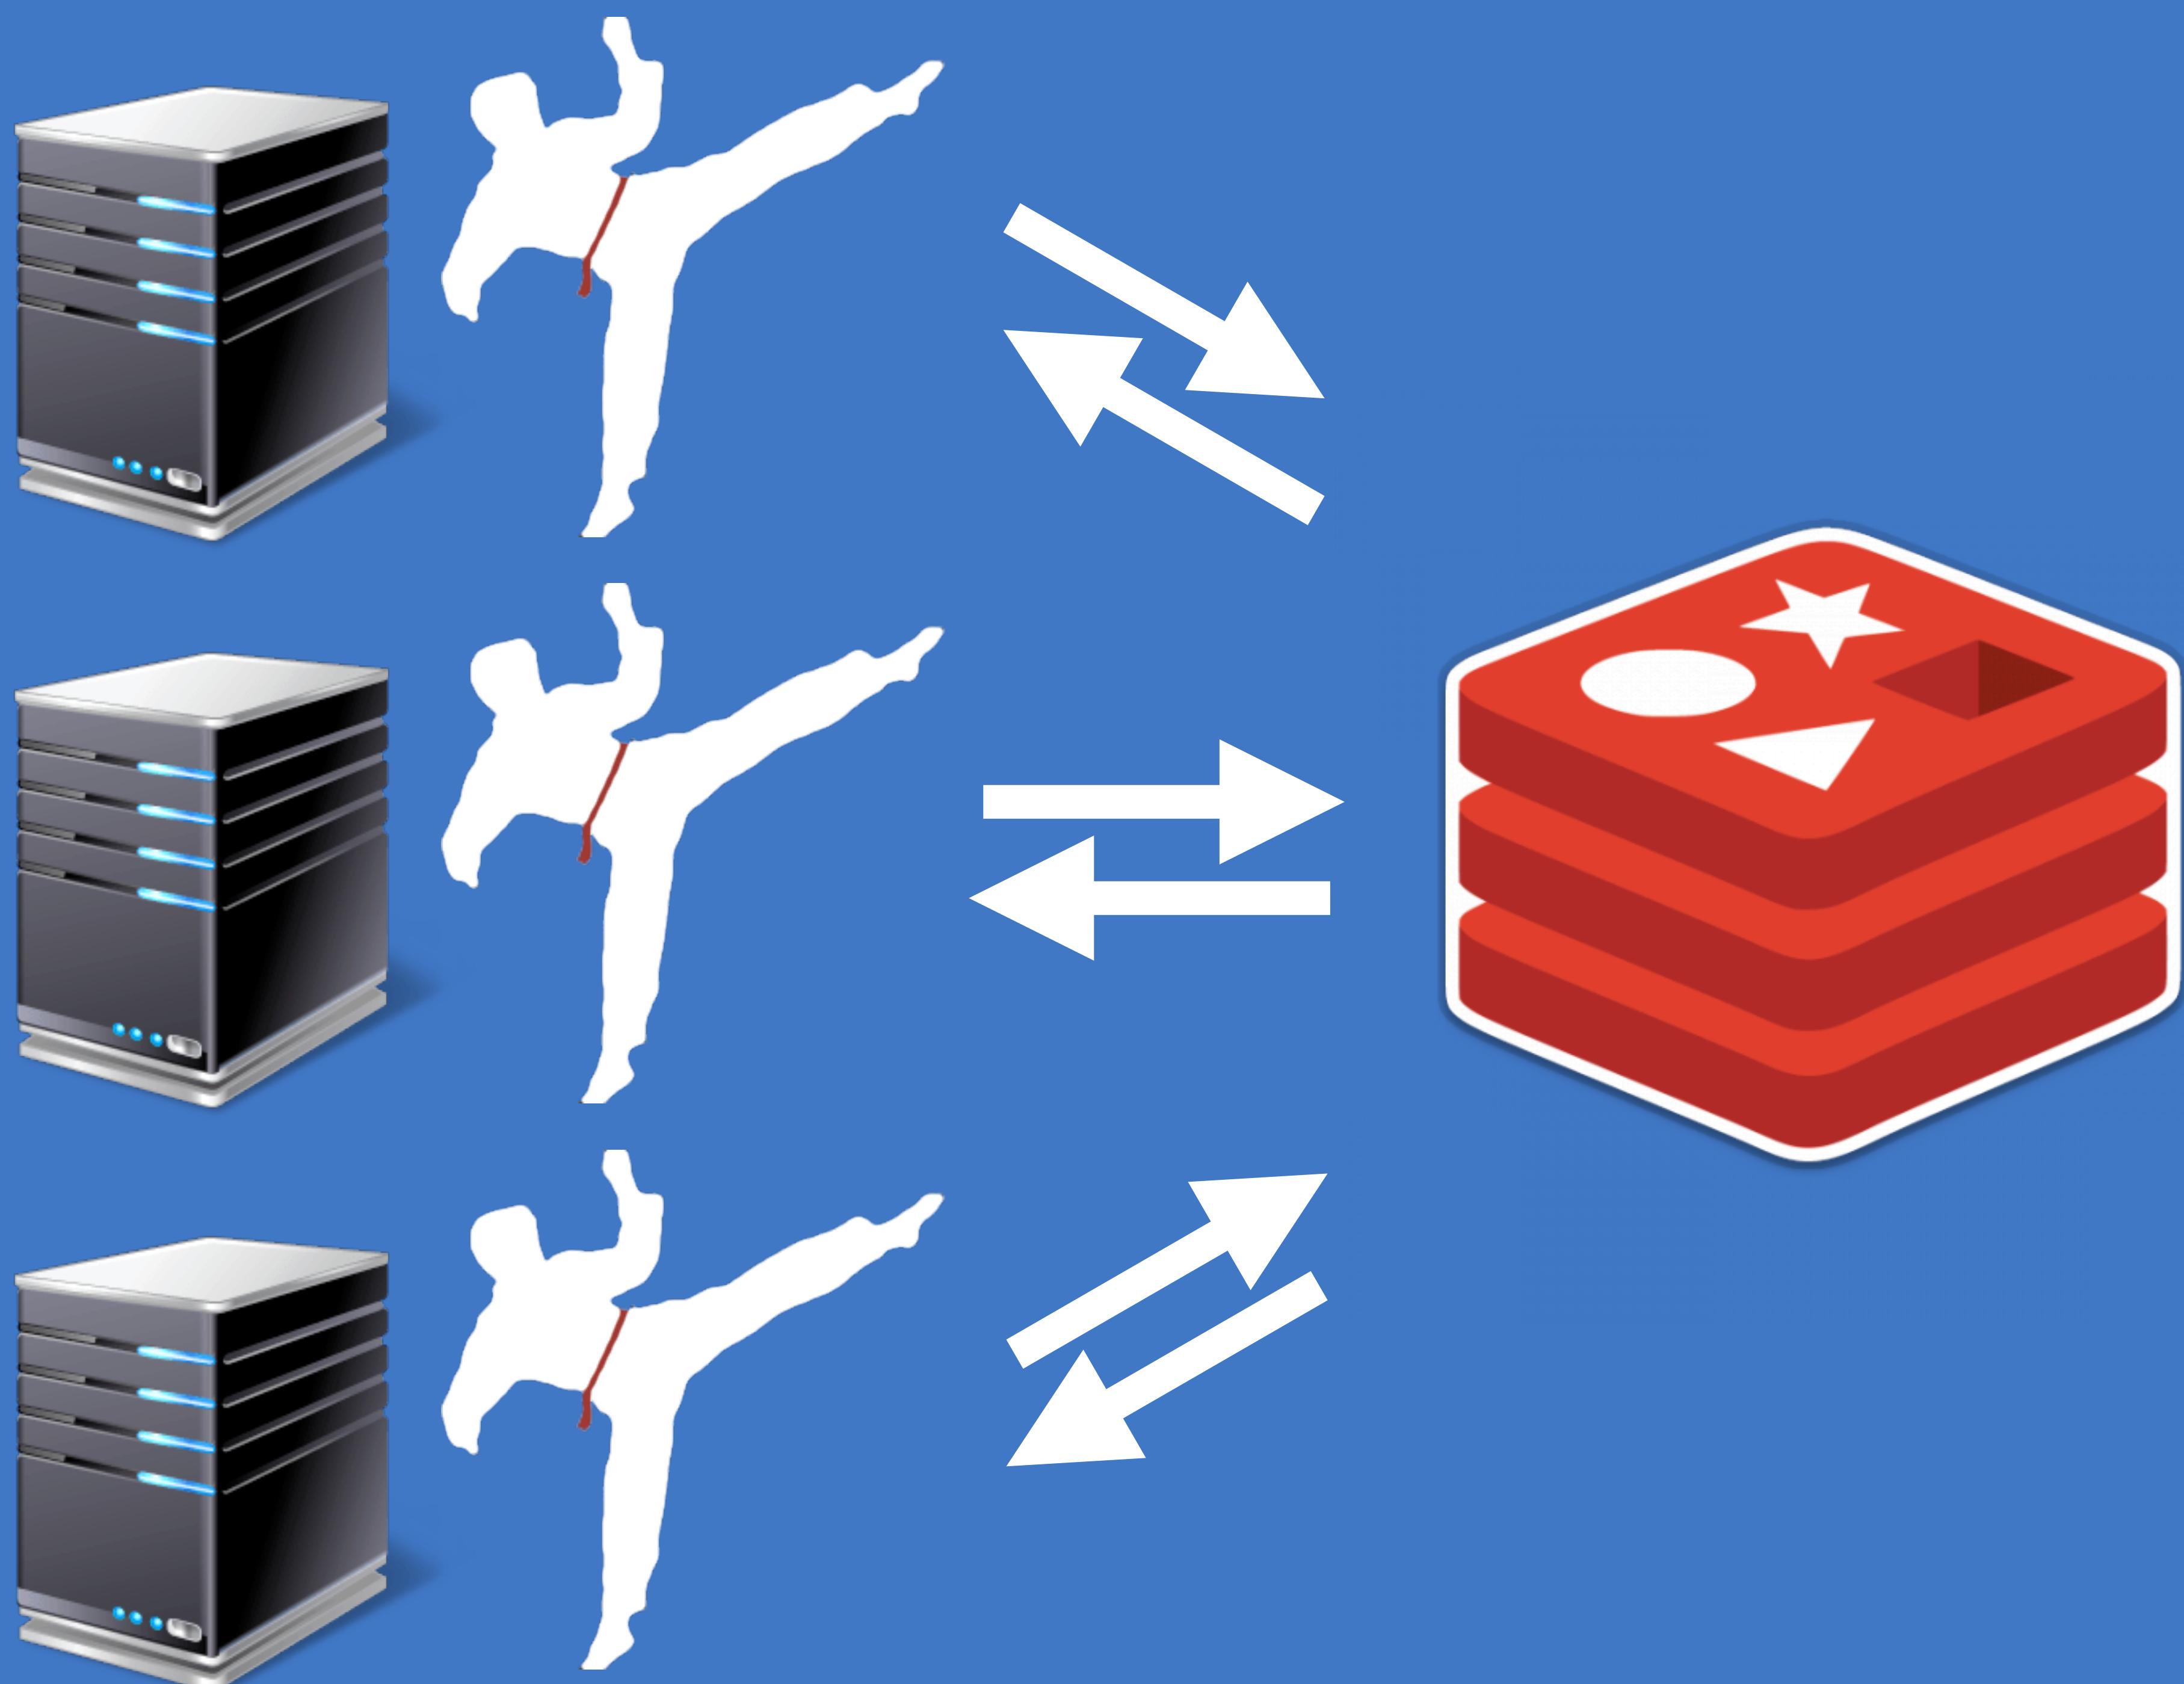
proc mem prod

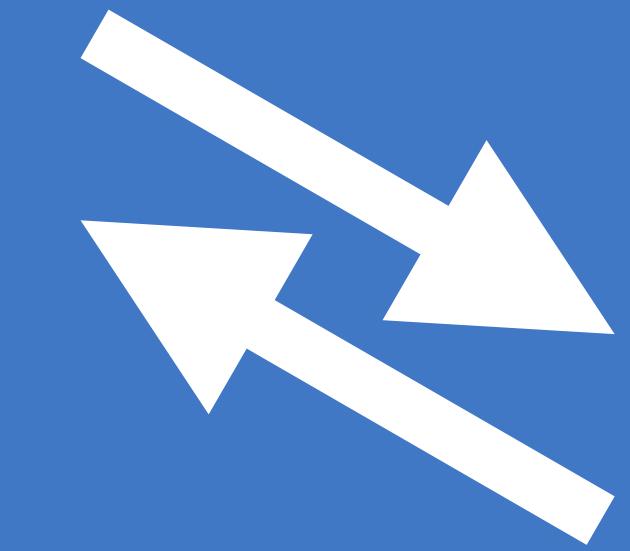
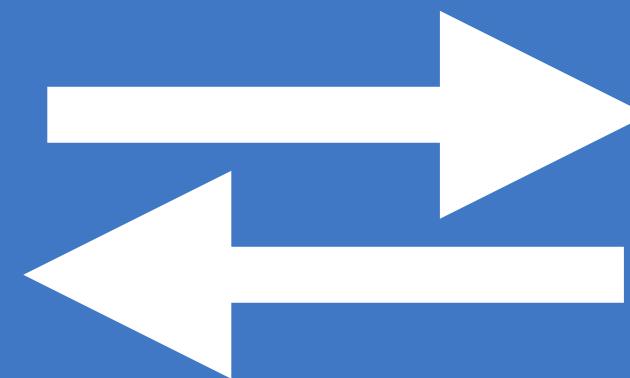
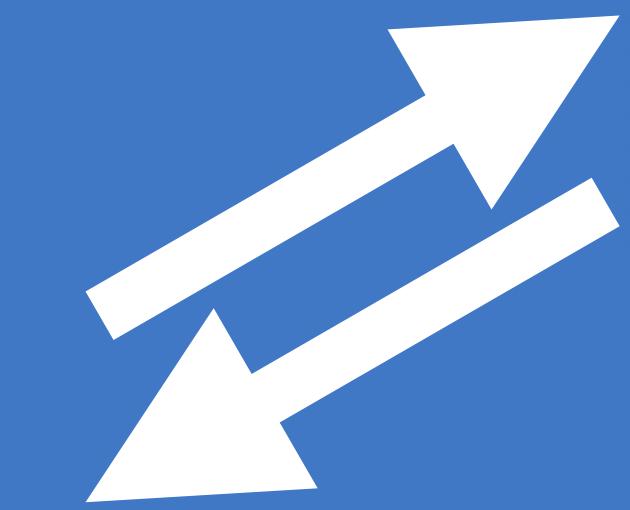


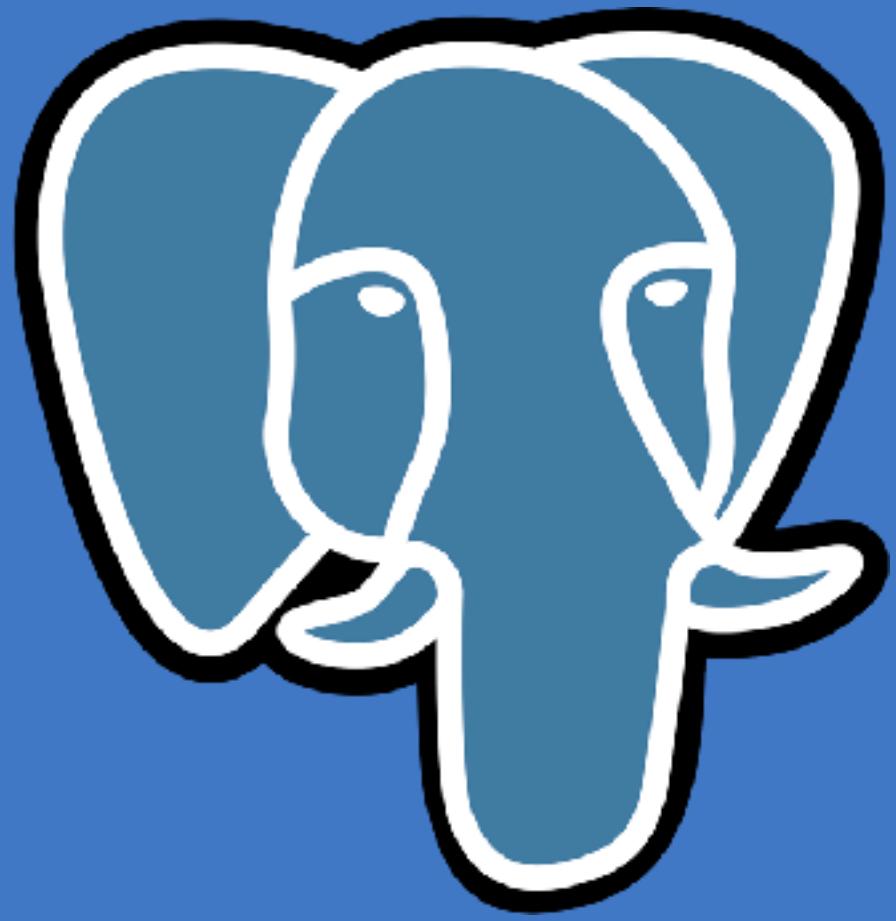
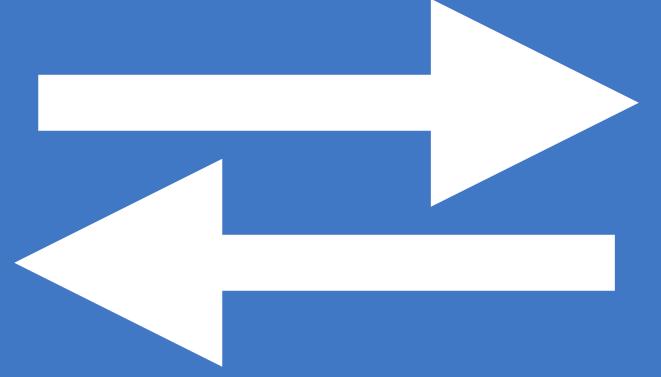
main process memory usage on production server

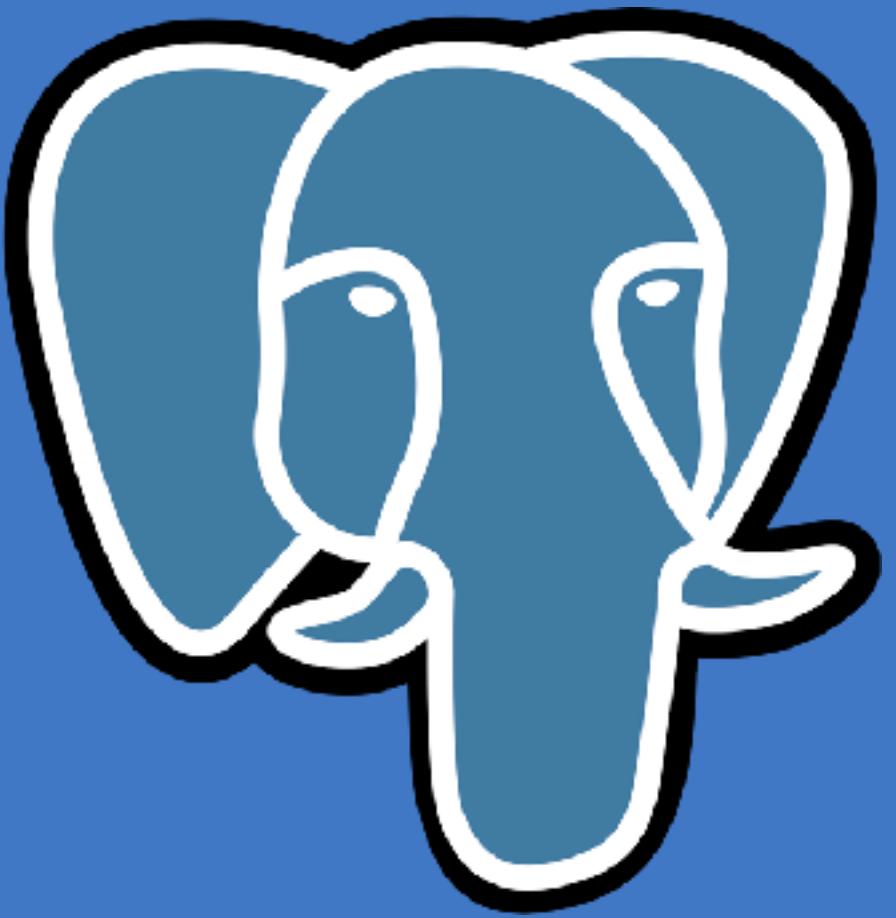
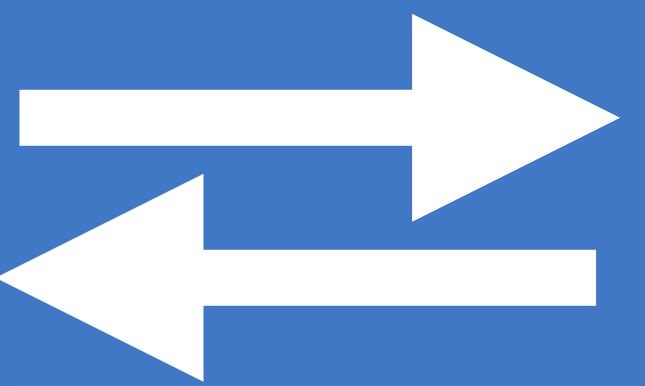
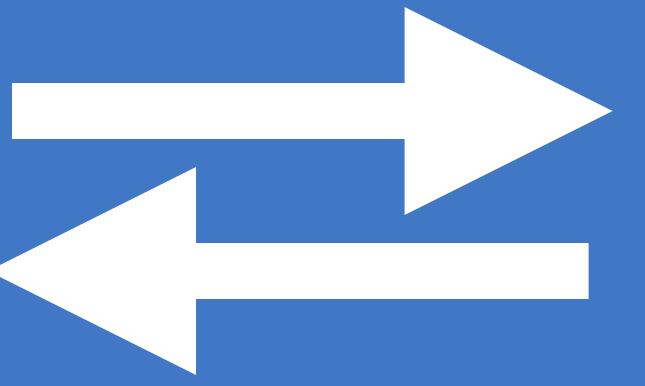
???













RABBITMQ





This repository

Search

Pull requests Issues Marketplace Explore



+



jondot / sneakers

Watch ▾

37

Unstar

1,597

Fork

238

Code

Issues 34

Pull requests 7

Projects 0

Wiki

Insights

A fast background processing framework for Ruby and RabbitMQ <http://sneakers.io>

394 commits

9 branches

19 releases

66 contributors

MIT

Branch: master ▾

New pull request

Create new file

Upload files

Find file

Clone or download ▾

 michaeklishin	committed on Jan 18 Merge pull request #335 from kke/patch-1	...	Latest commit c647e34 on Jan 18
 bin	Allow setting logger and provide server configuration convenience block.		3 years ago
 examples	Indicate that worker is a background task for NewRelic		11 months ago
 lib	Merge branch 'fix_connected_session_bug' of https://github.com/webit-...		4 months ago
 scripts	Add a fat and slim docker builds		2 years ago
 spec	Wording		4 months ago
 .gitignore	Maxretry Handler		4 years ago
 .travis.yml	Test against the same set of Rubies as Bunny		11 months ago
 Dockerfile	Add a fat and slim docker builds		2 years ago

```
1 class TitleScraper
2   include Sneakers::Worker
3
4   from_queue 'downloads'
5
6   def work(msg)
7     title = extract_title(msg)
8     logger.info "FOUND <#{title}>"
9     ack!
10  end
11
12  private
13
14  def extract_title(html)
15    html.scan(/<title>(.*)</title>/).flatten.first
16  end
17 end
```

```
i 1 class TestWorker
2   include Sidekiq::Worker
3
4   def perform(*args)
5     # Do something
6   end
7 end
+
8
+
9
+
10
+
11
+
12
+
13
+
14
+
15
+
16
+
17
```

```
1 class TitleScraper
2   include Sneakers::Worker
3
4   from_queue 'downloads'
5
6   def work(msg)
7     title = extract_title(msg)
8     logger.info "FOUND <#{title}>"
9     ack!
10  end
11
12  private
13
14  def extract_title(html)
15    html.scan(/<title>(.*)<\title>/).flatten.first
16  end
17 end
```

```
1 class TitleScraper
2   include Sneakers::Worker
3
4   from_queue 'downloads'
5
6   def work(msg)
7     title = extract_title(msg)
8     logger.info "FOUND <#{title}>"
9     ack!
10
11
12 private
13
14   def extract_title(html)
15     html.scan(/<title>(.*)<\title>/).flatten.first
16   end
17 end
```

Queue

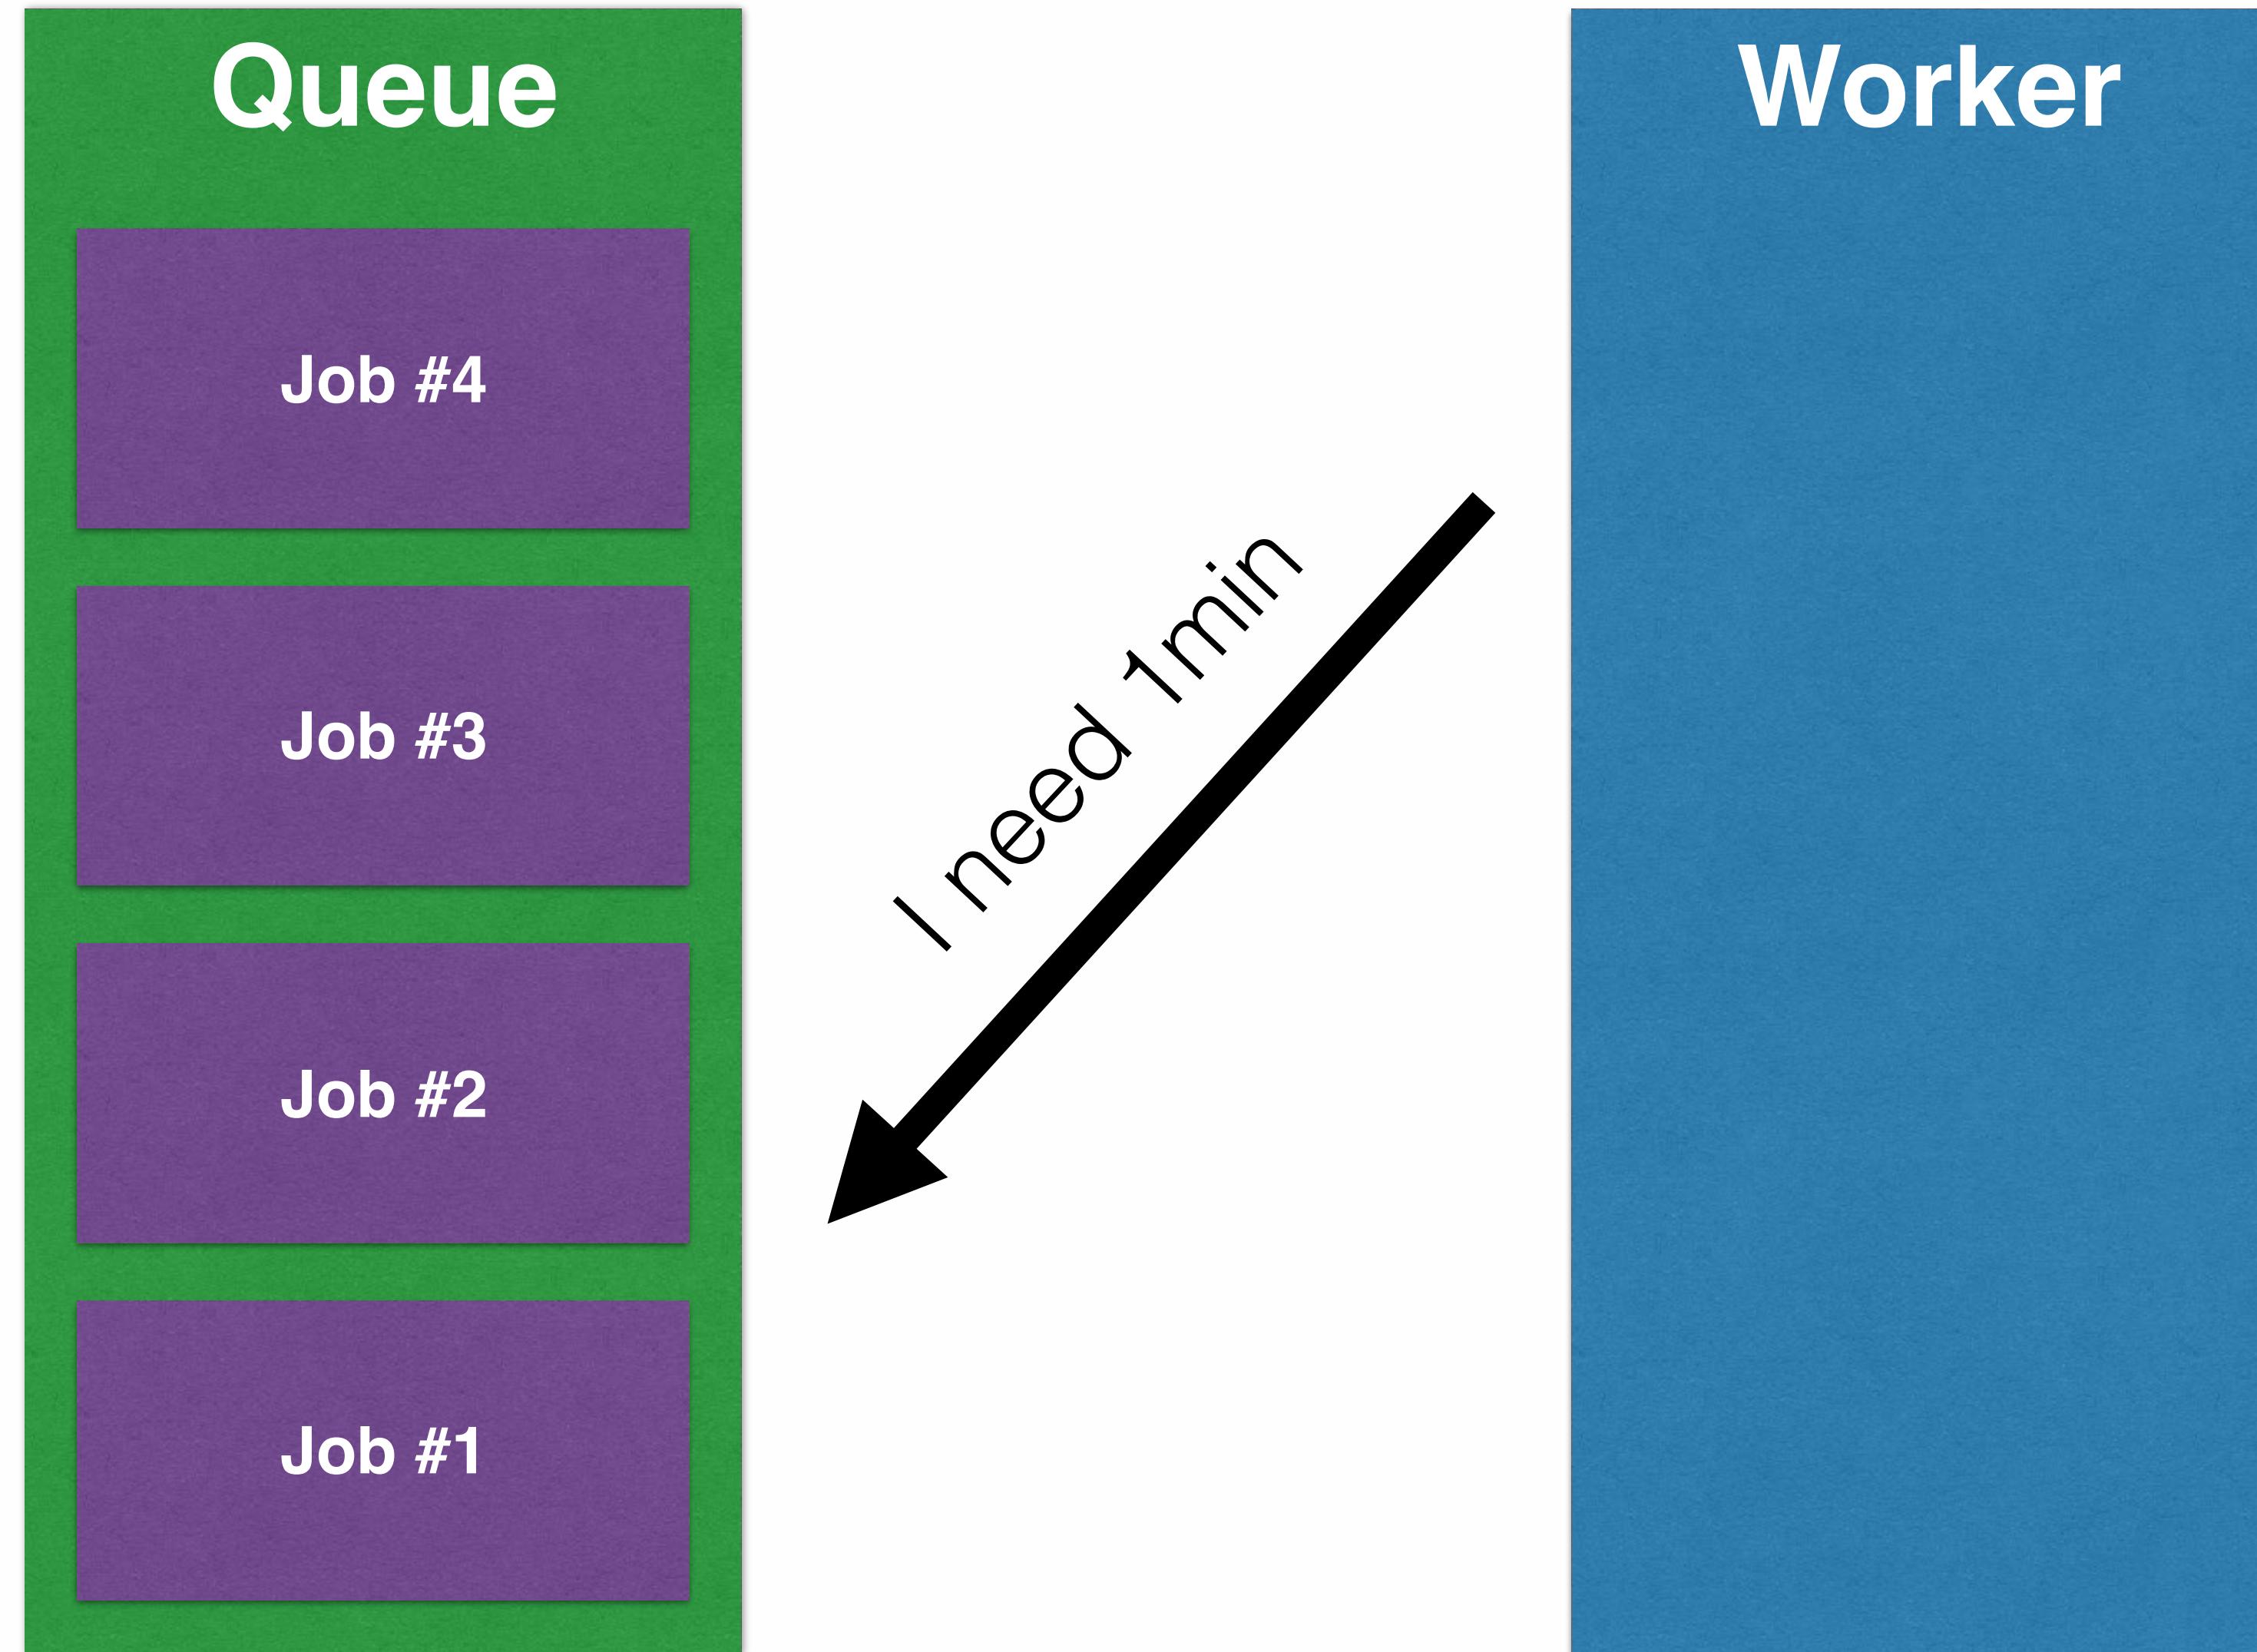
Job #4

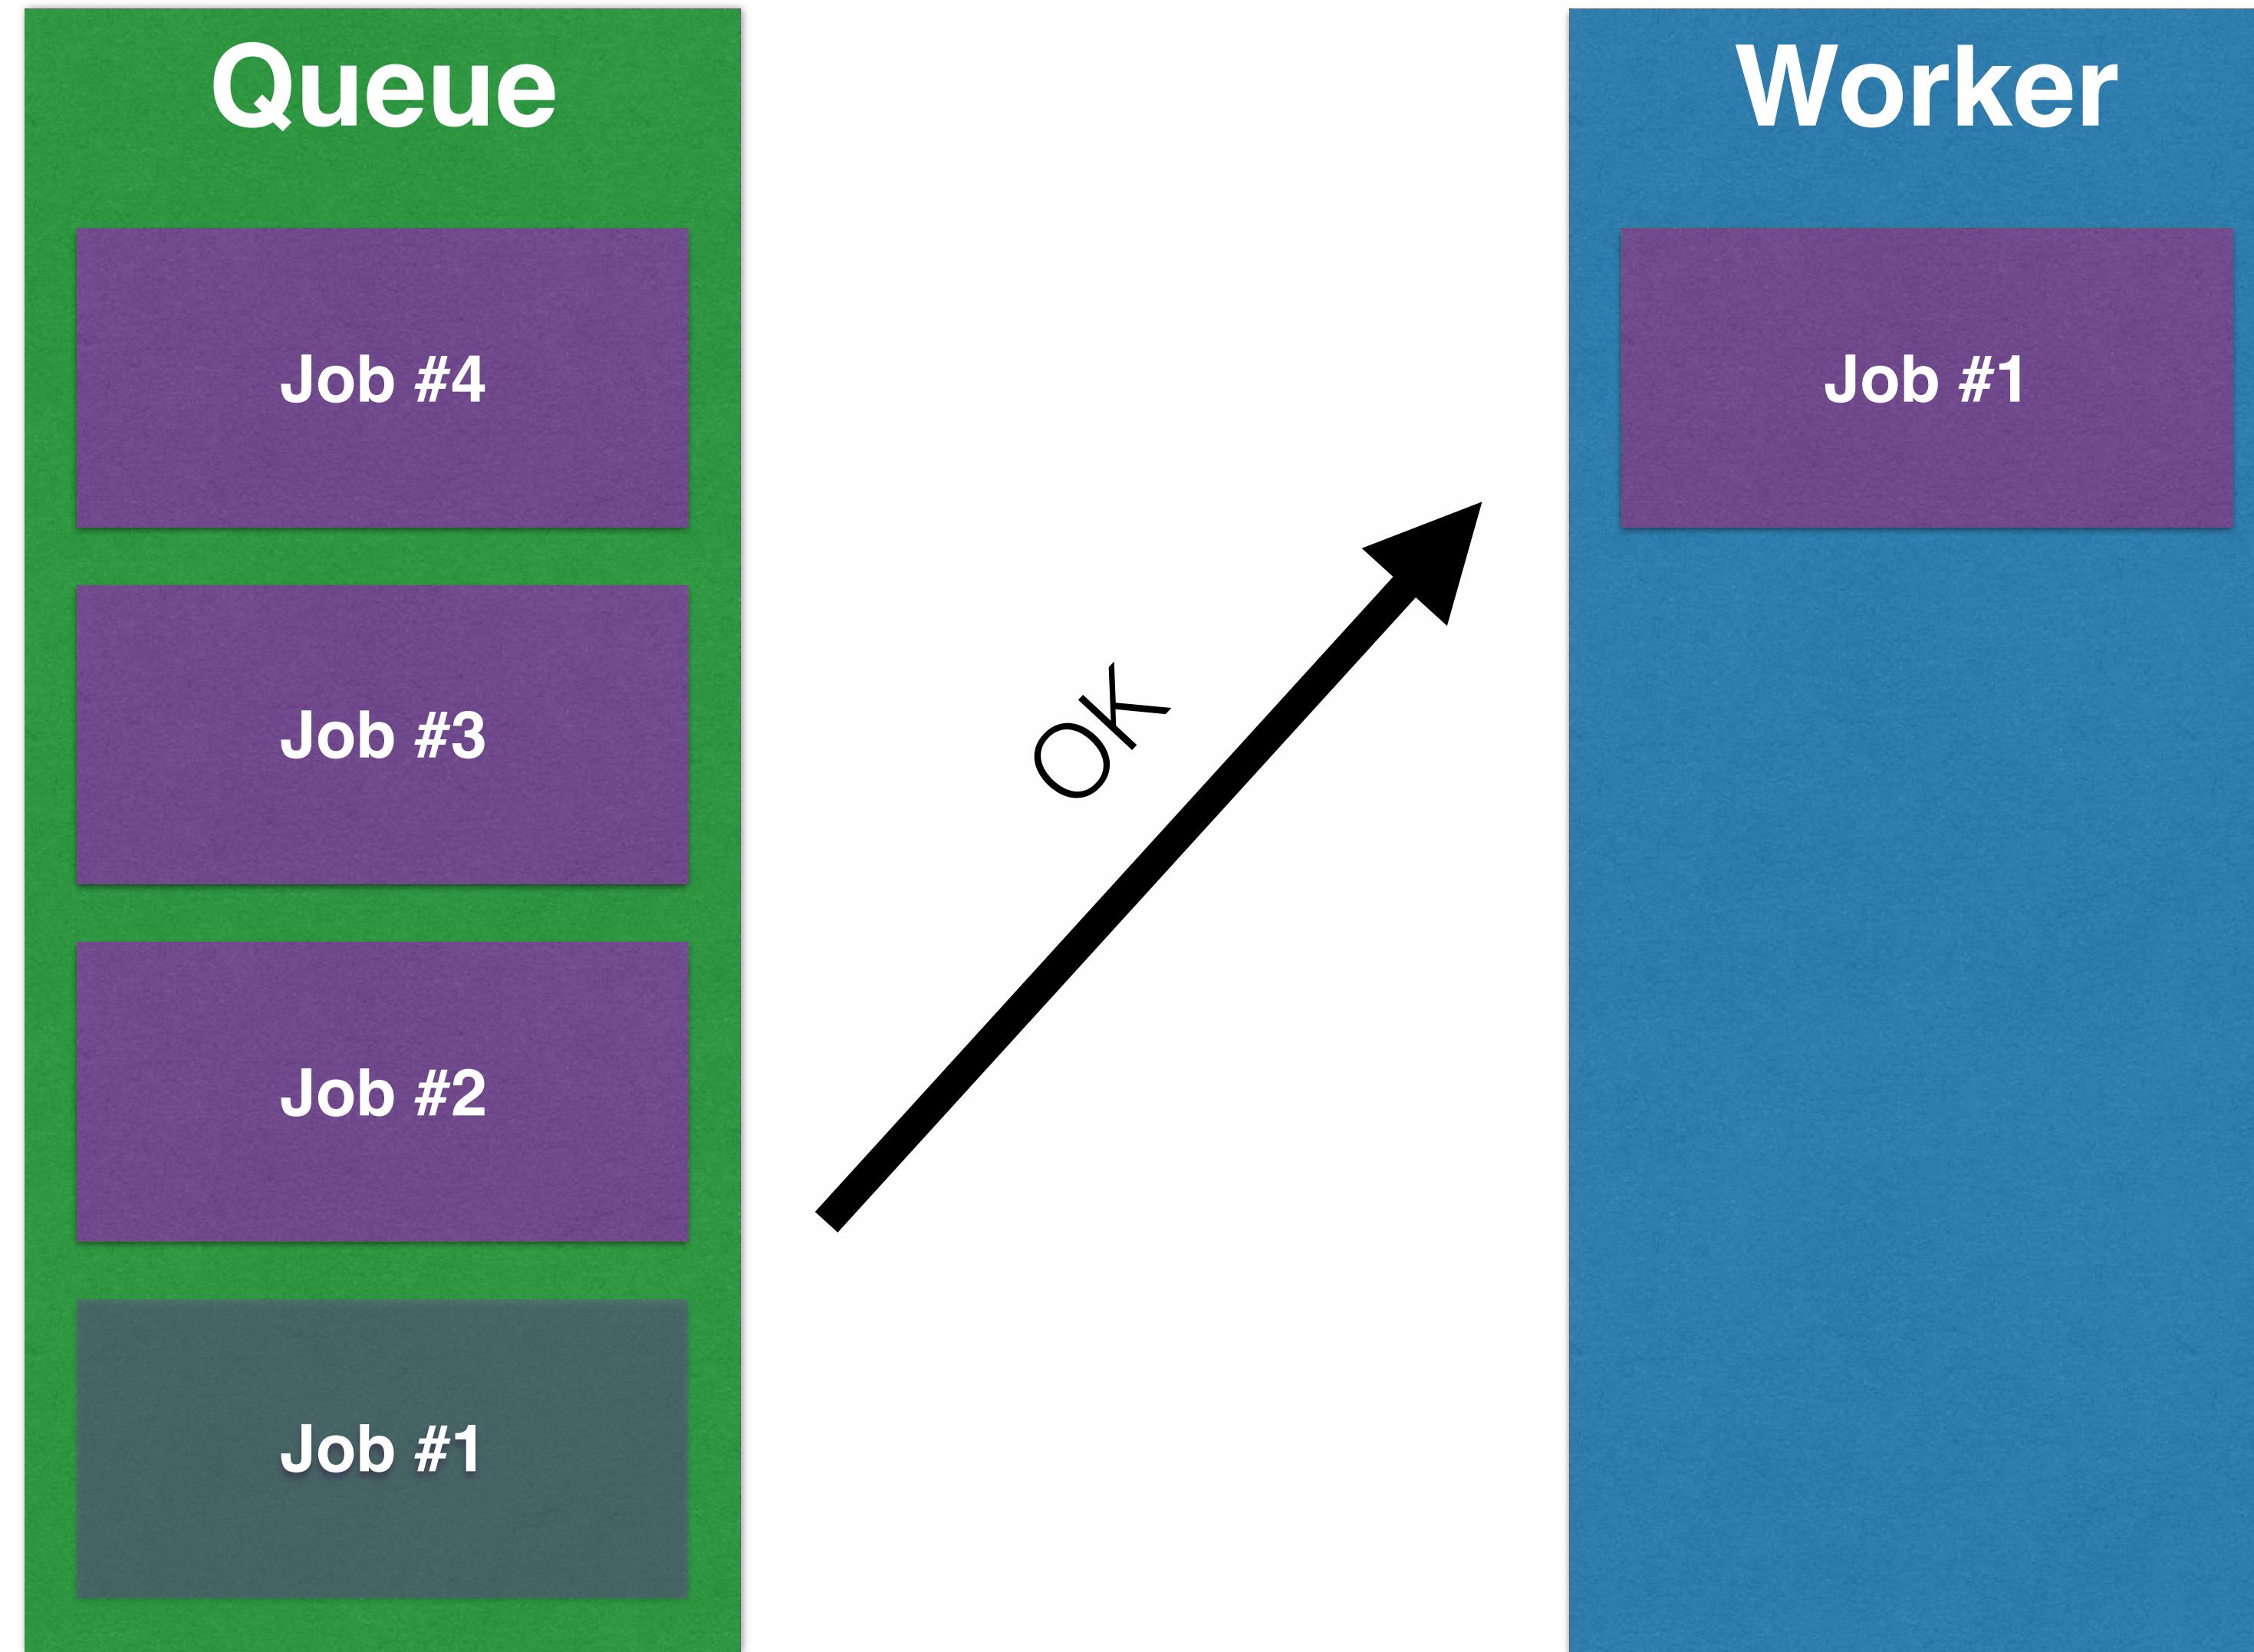
Job #3

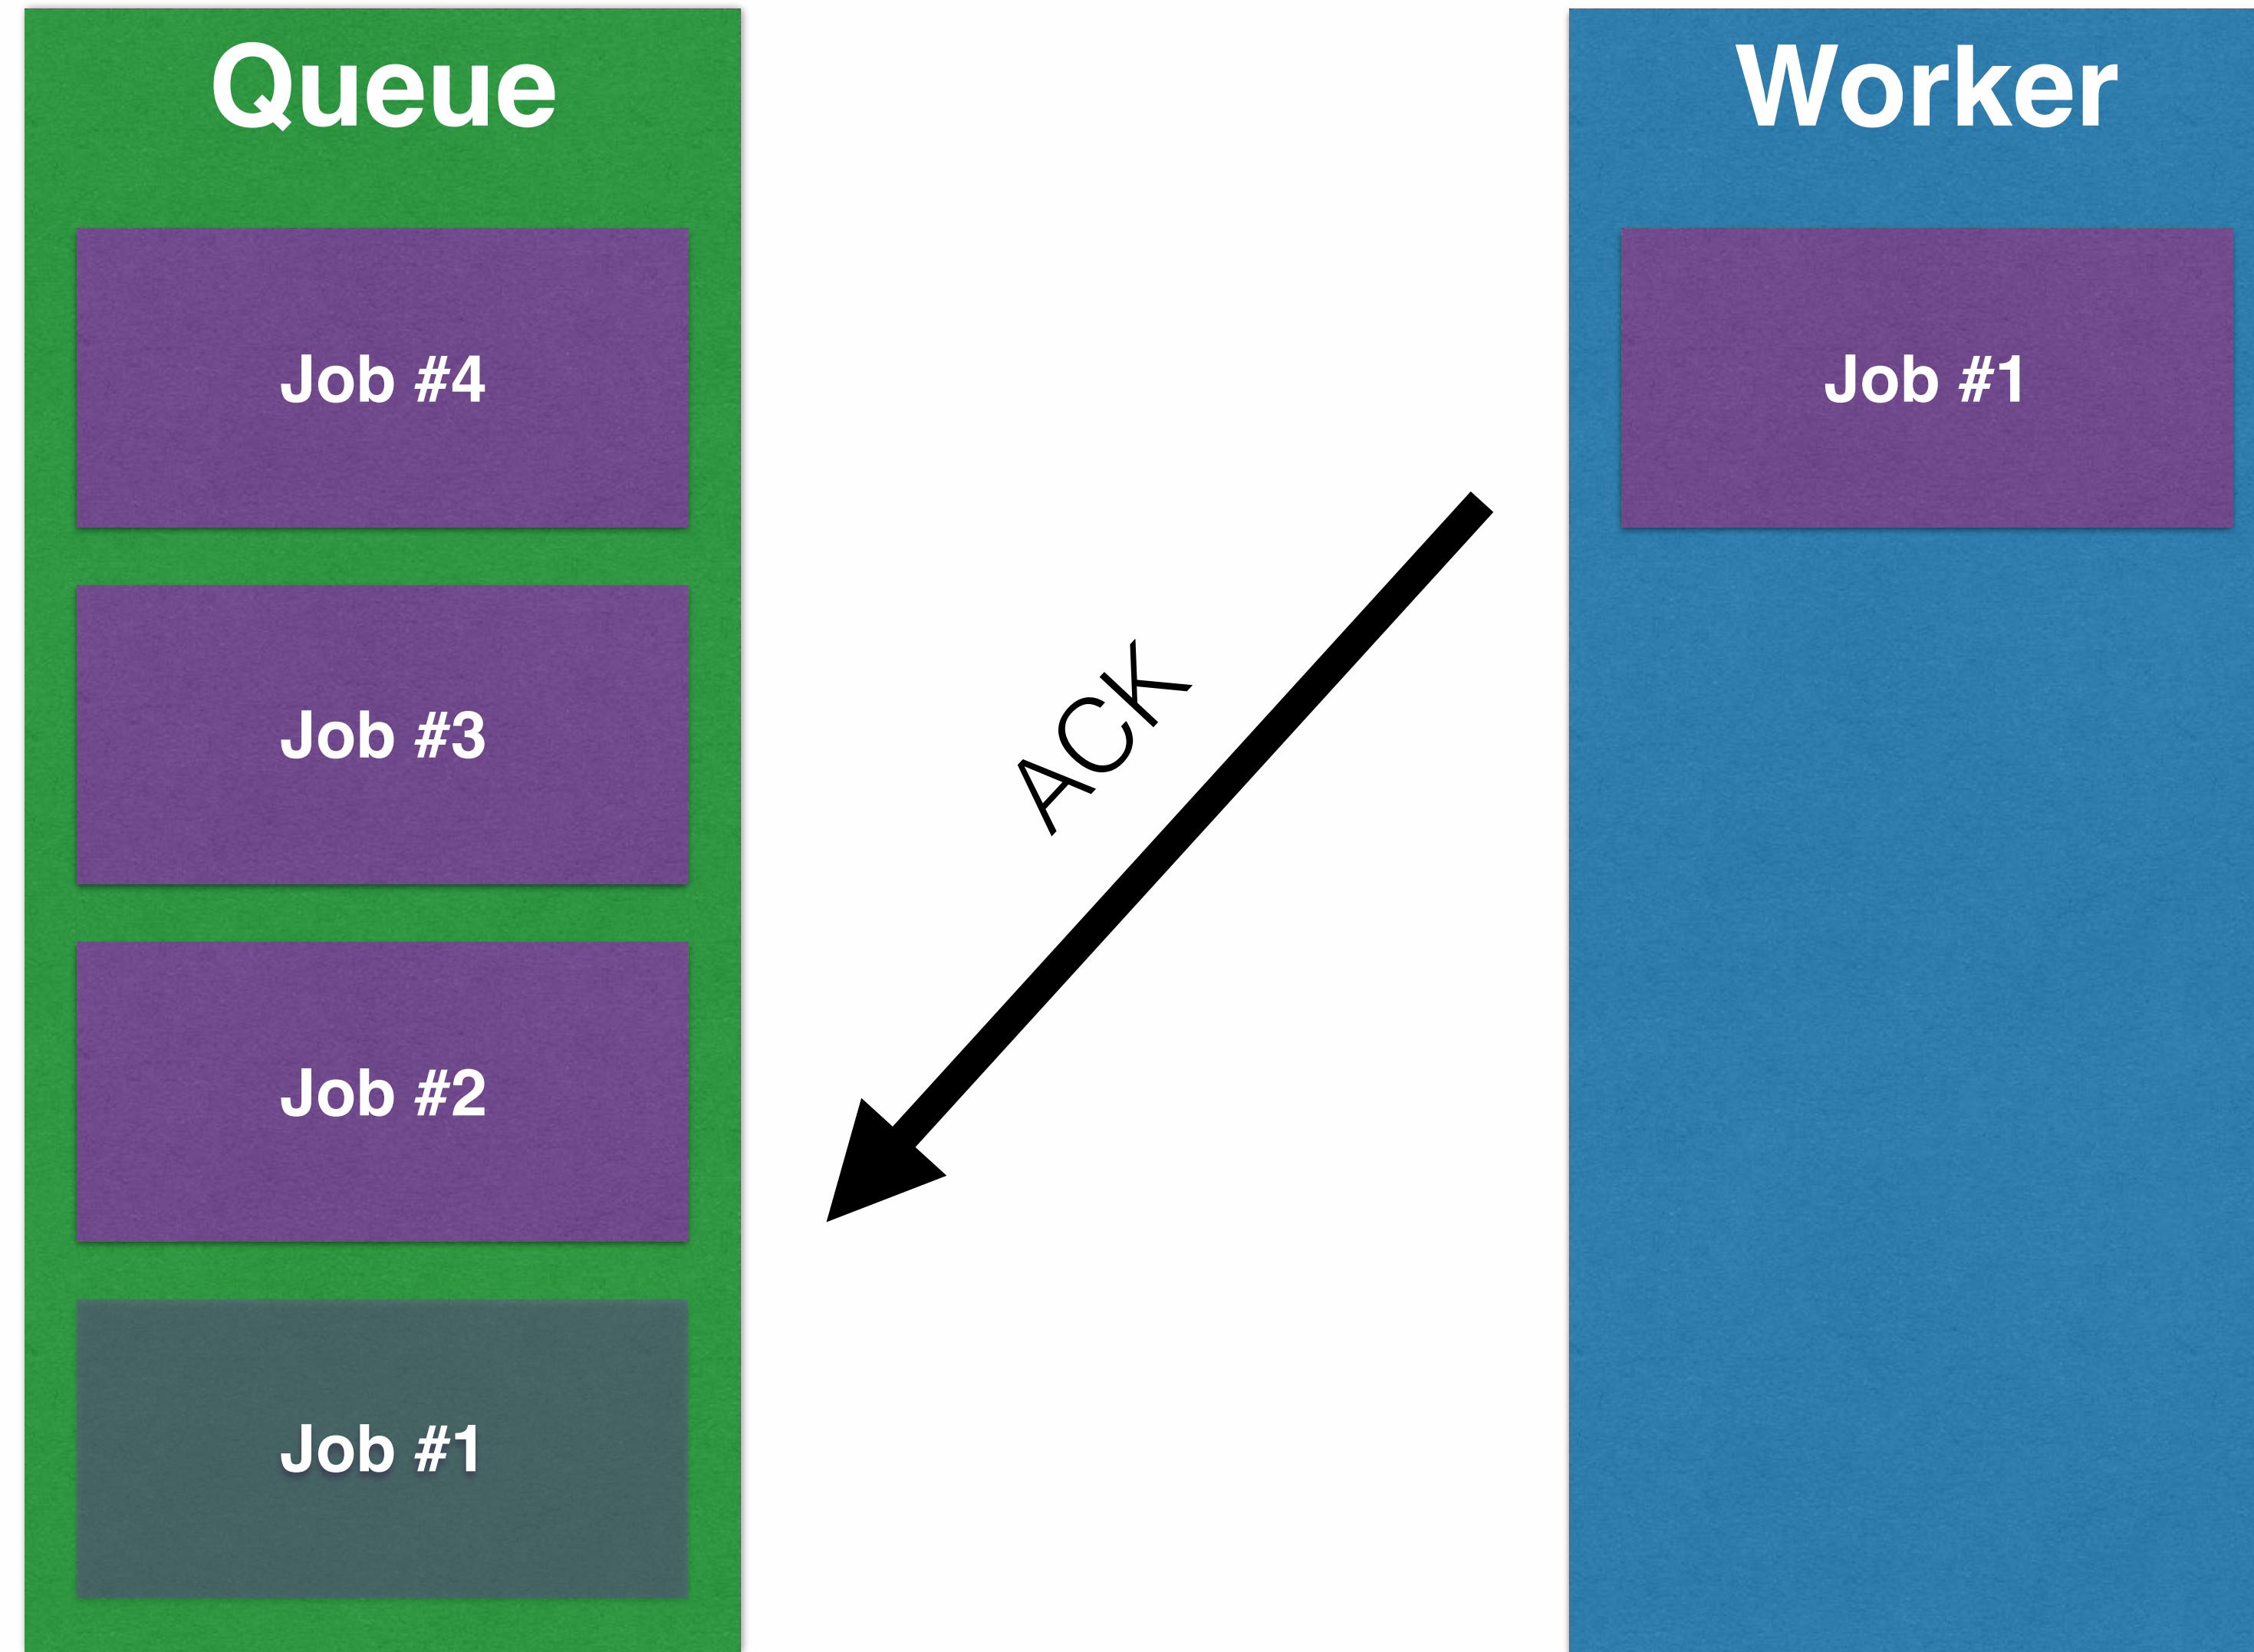
Job #2

Job #1

Worker







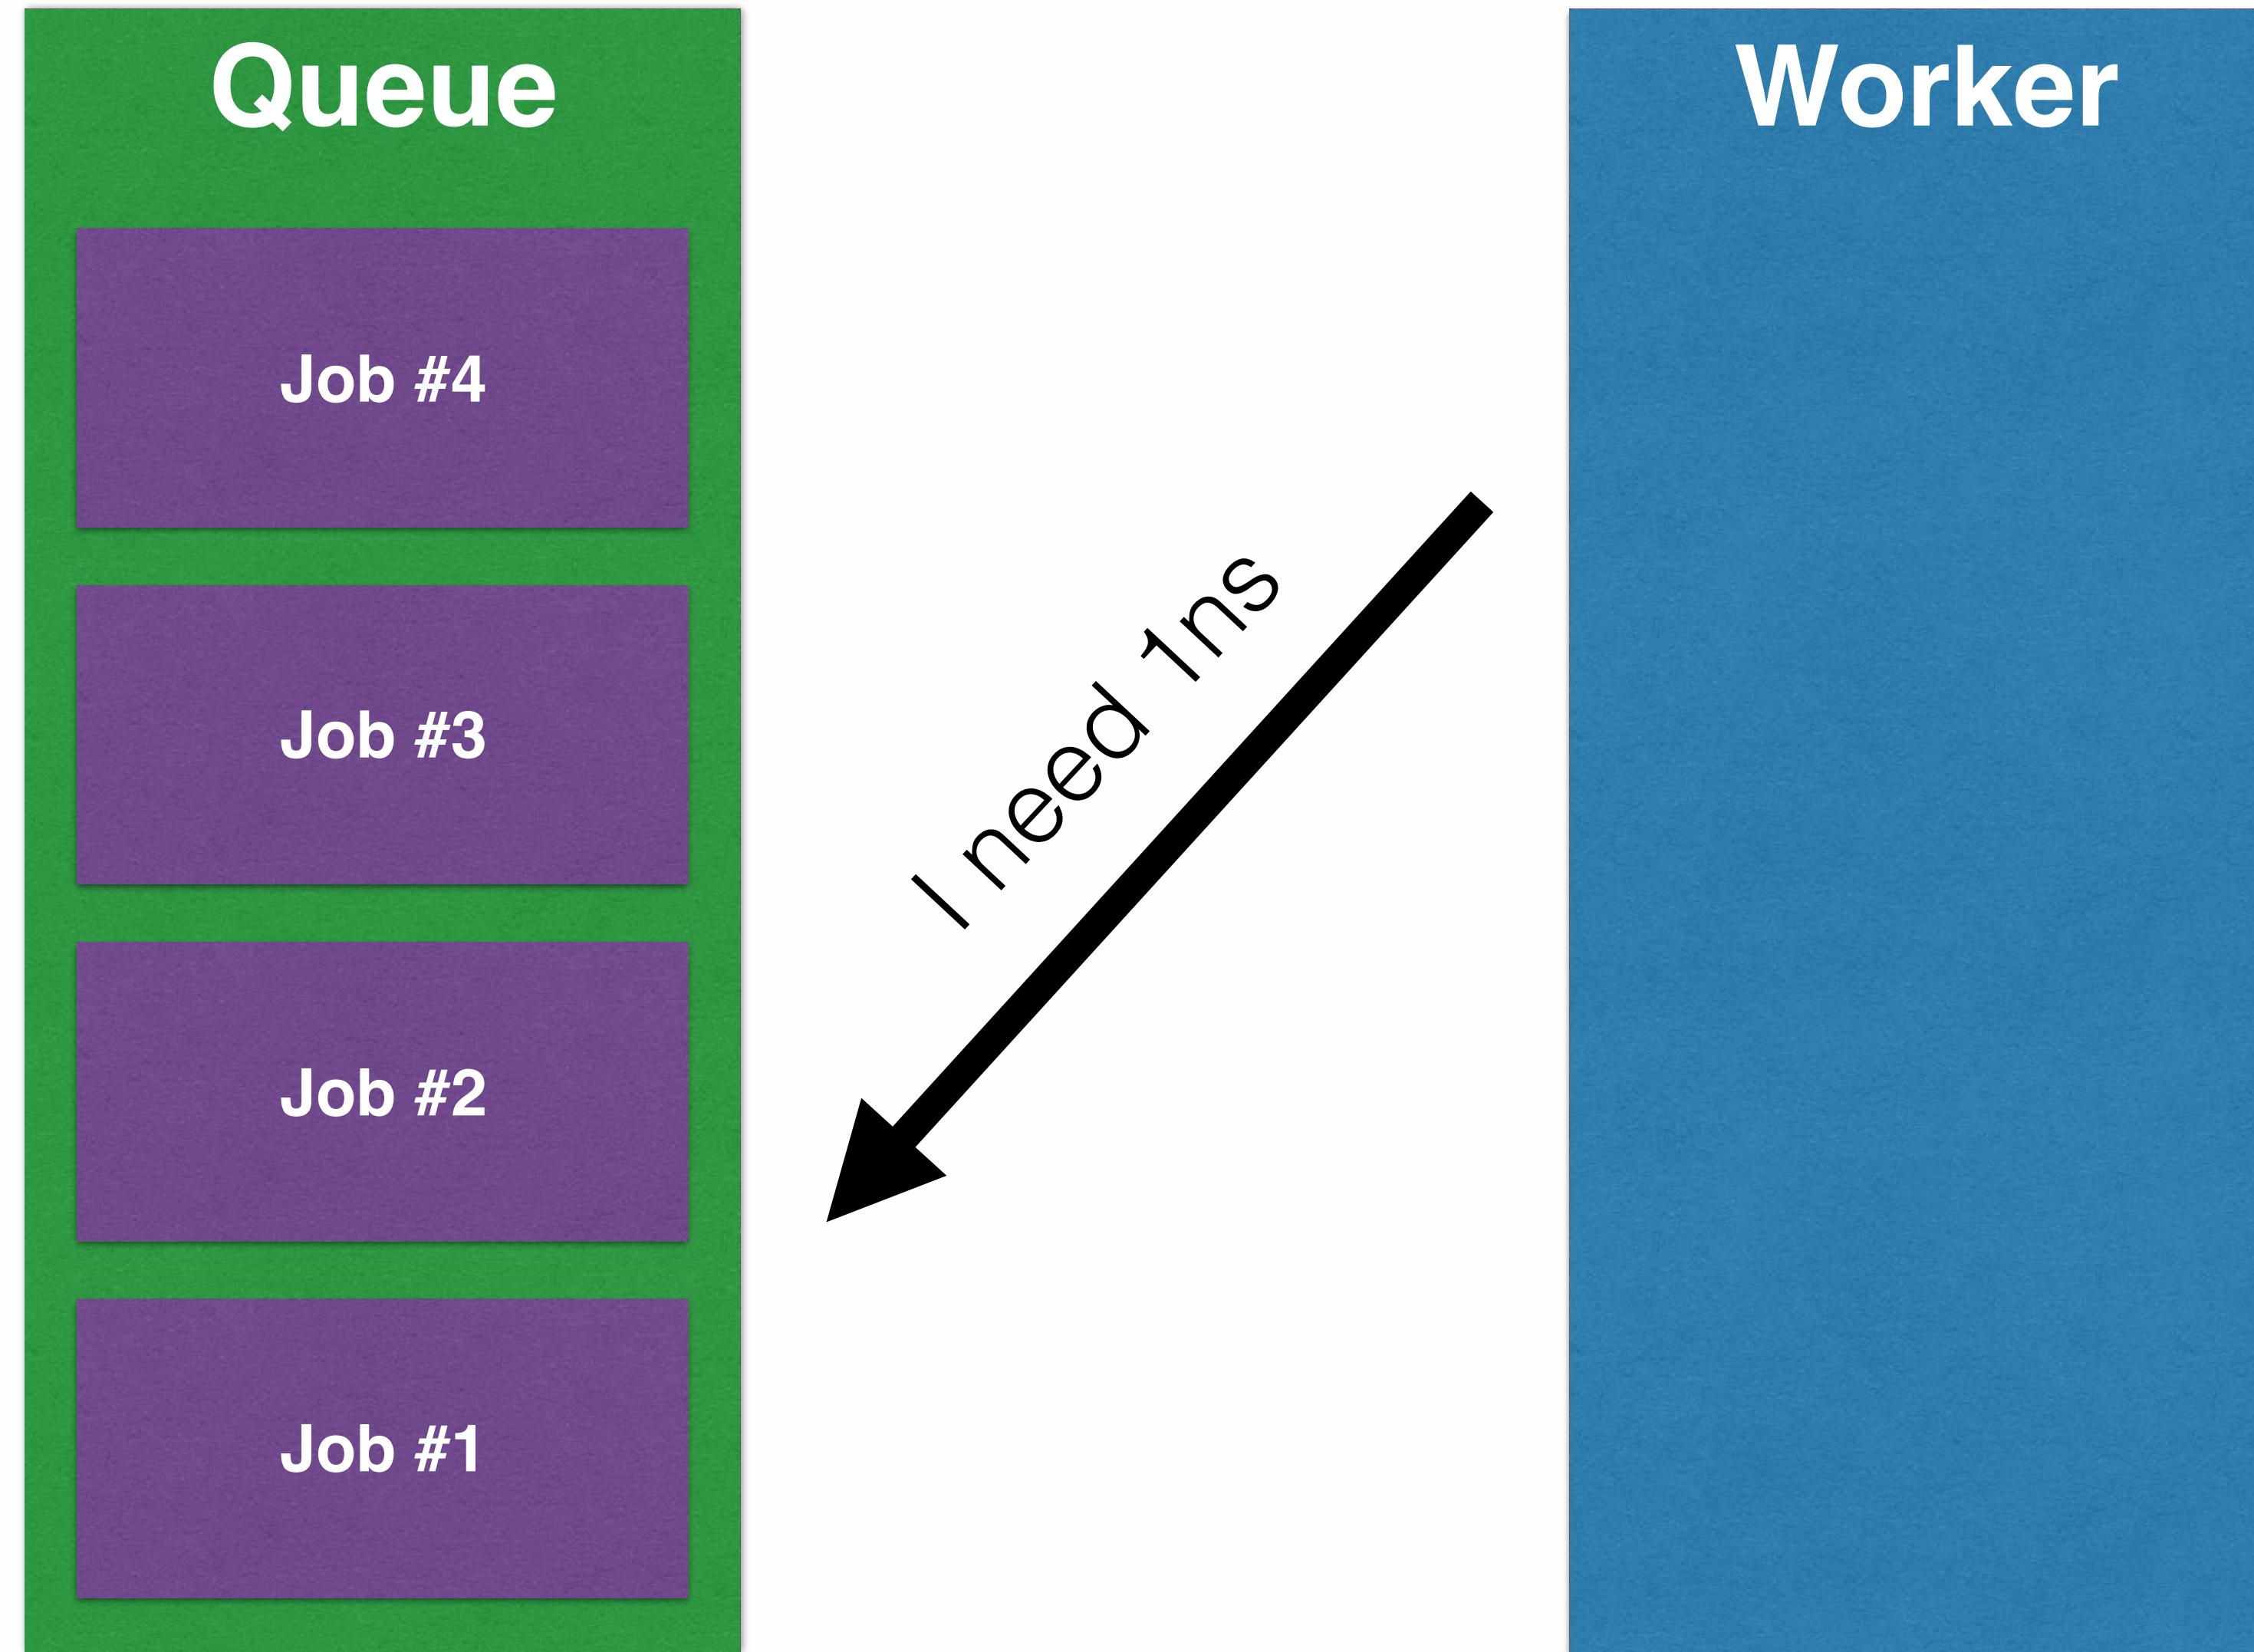
Queue

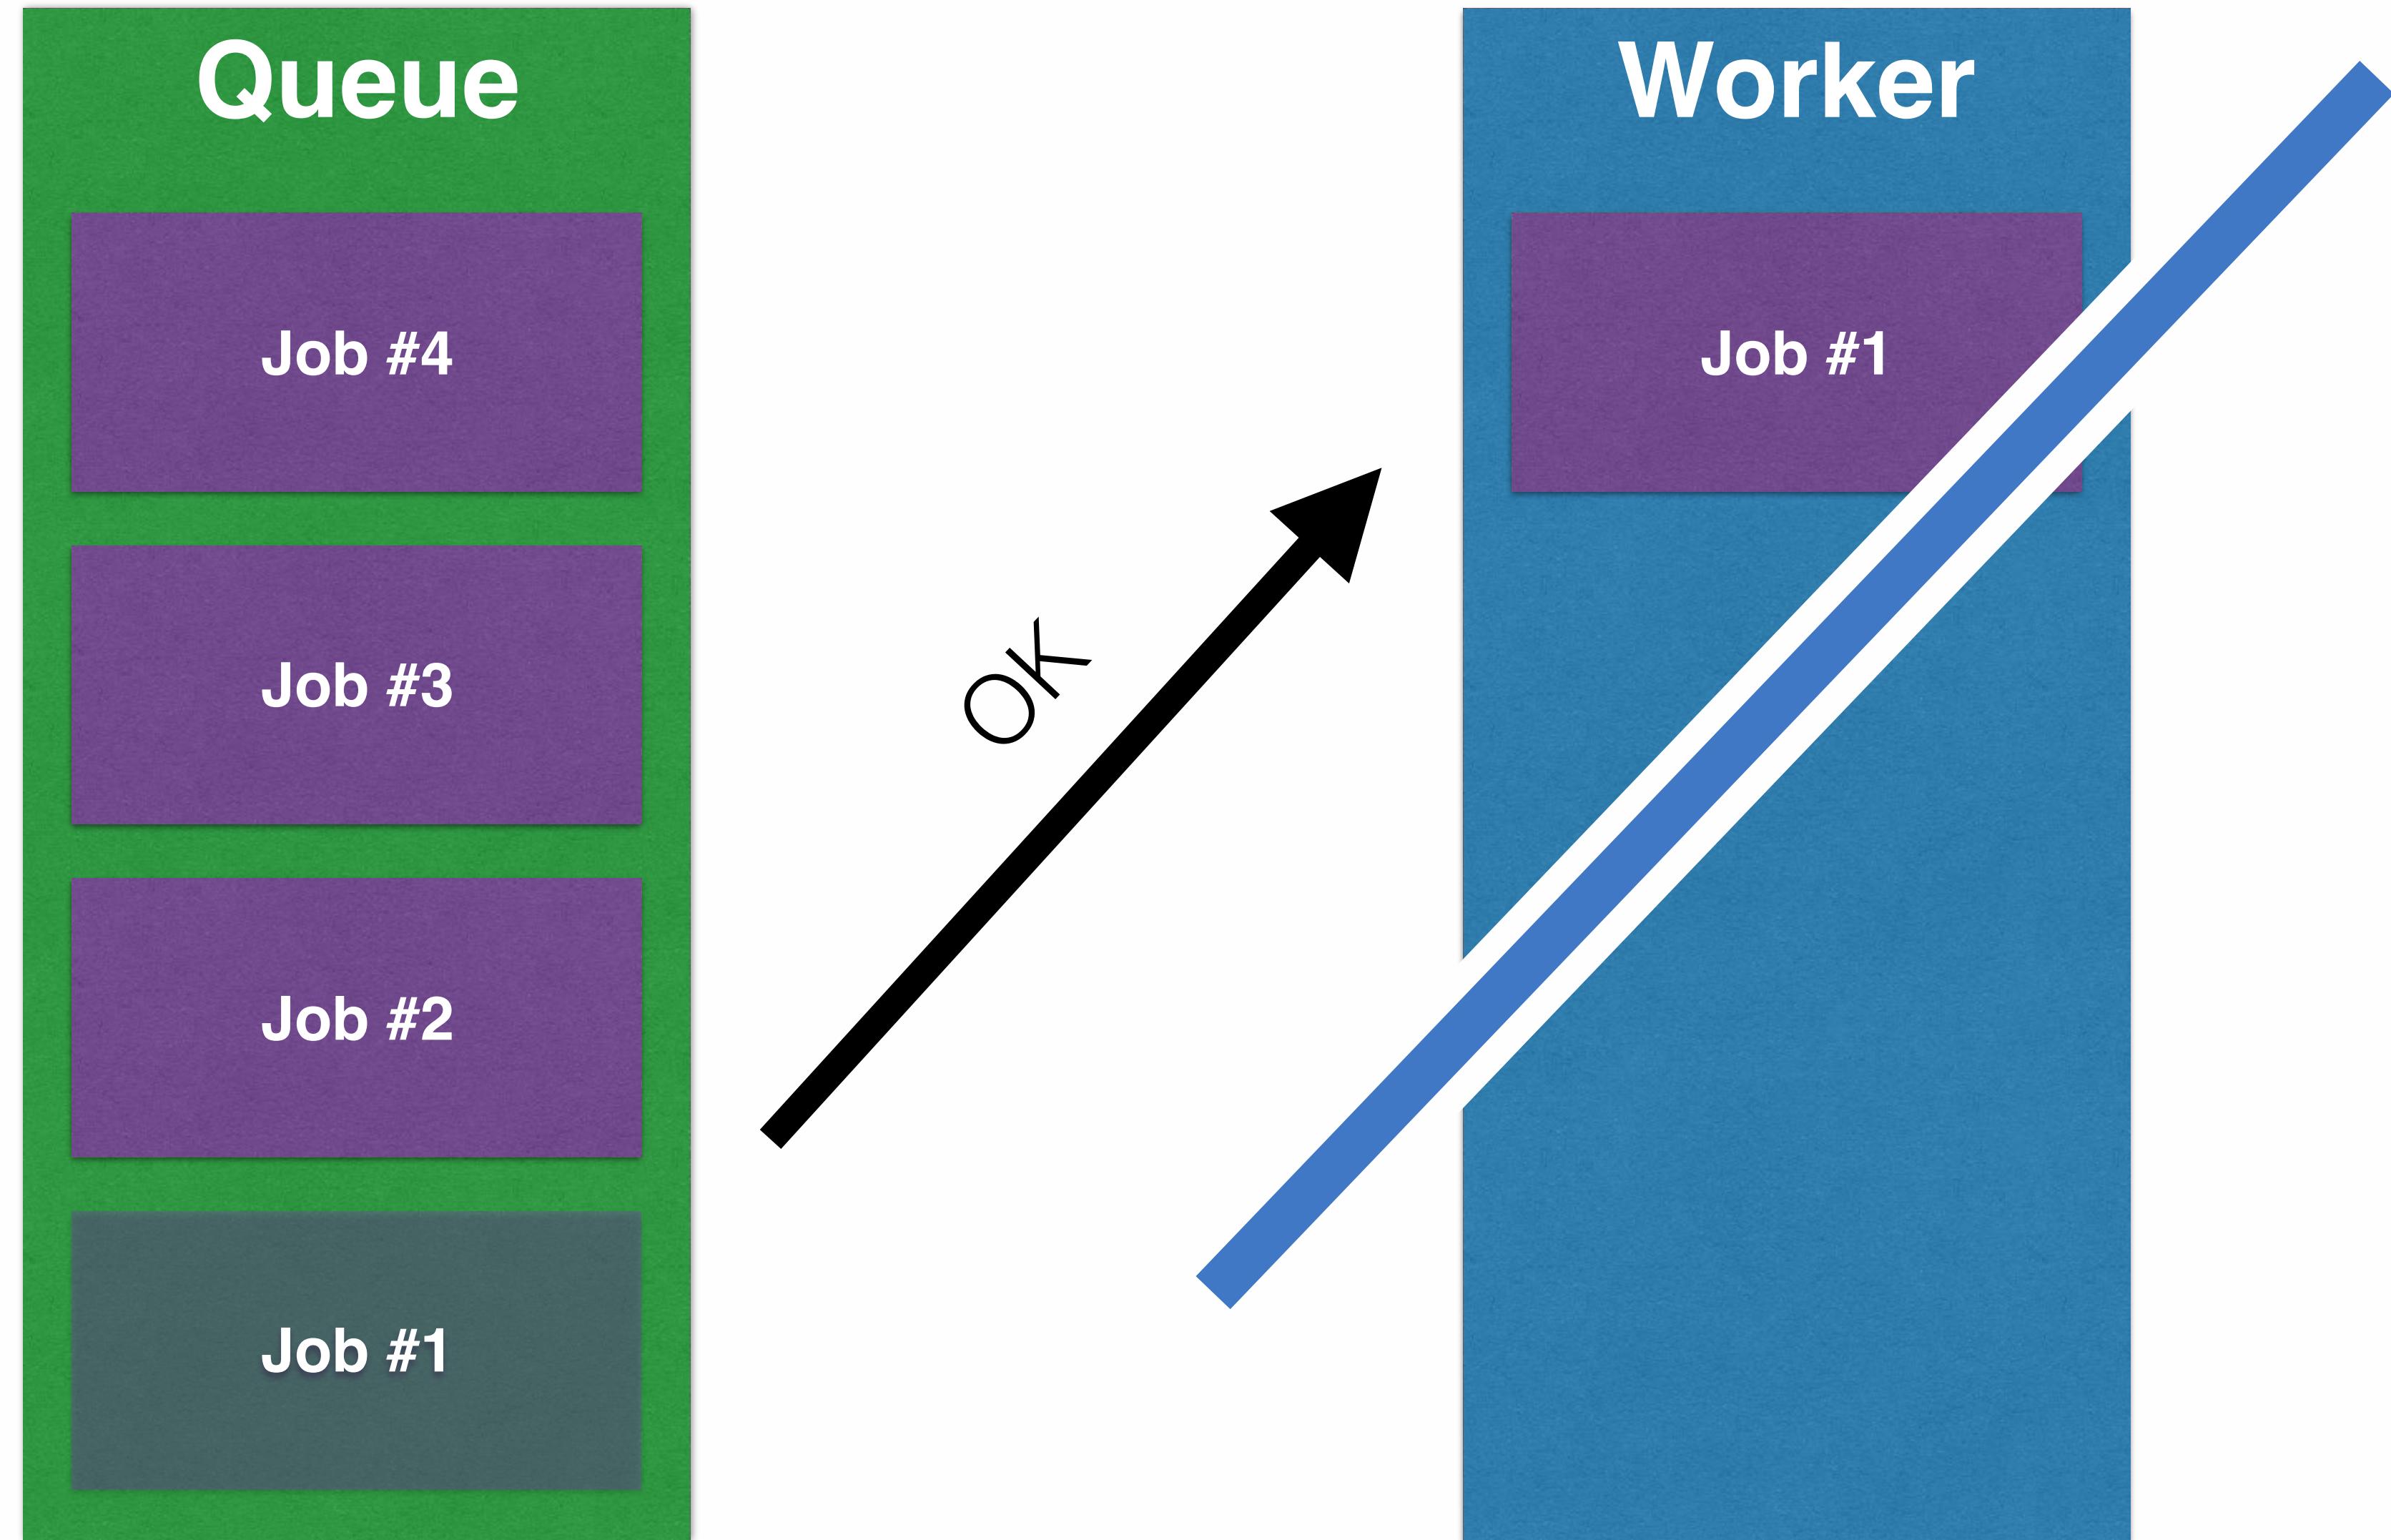
Job #4

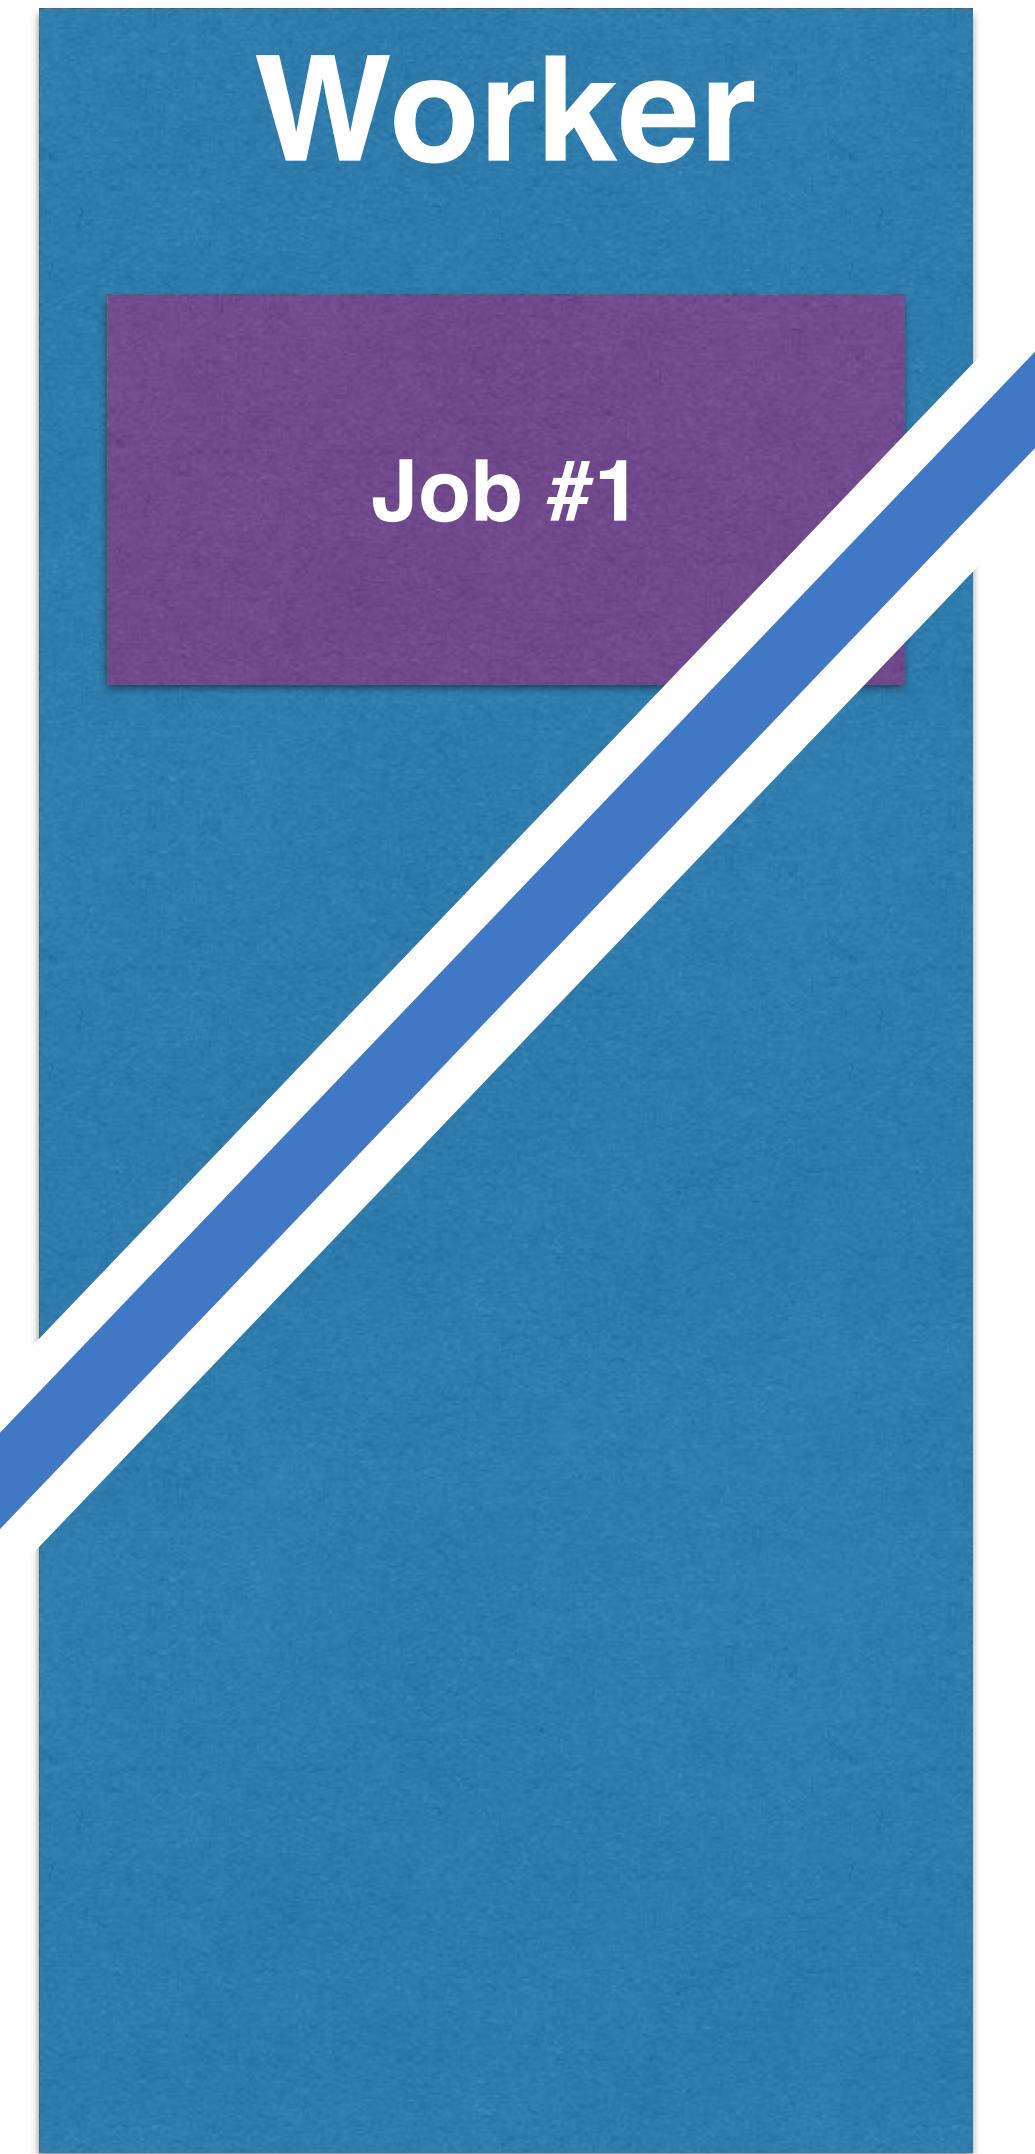
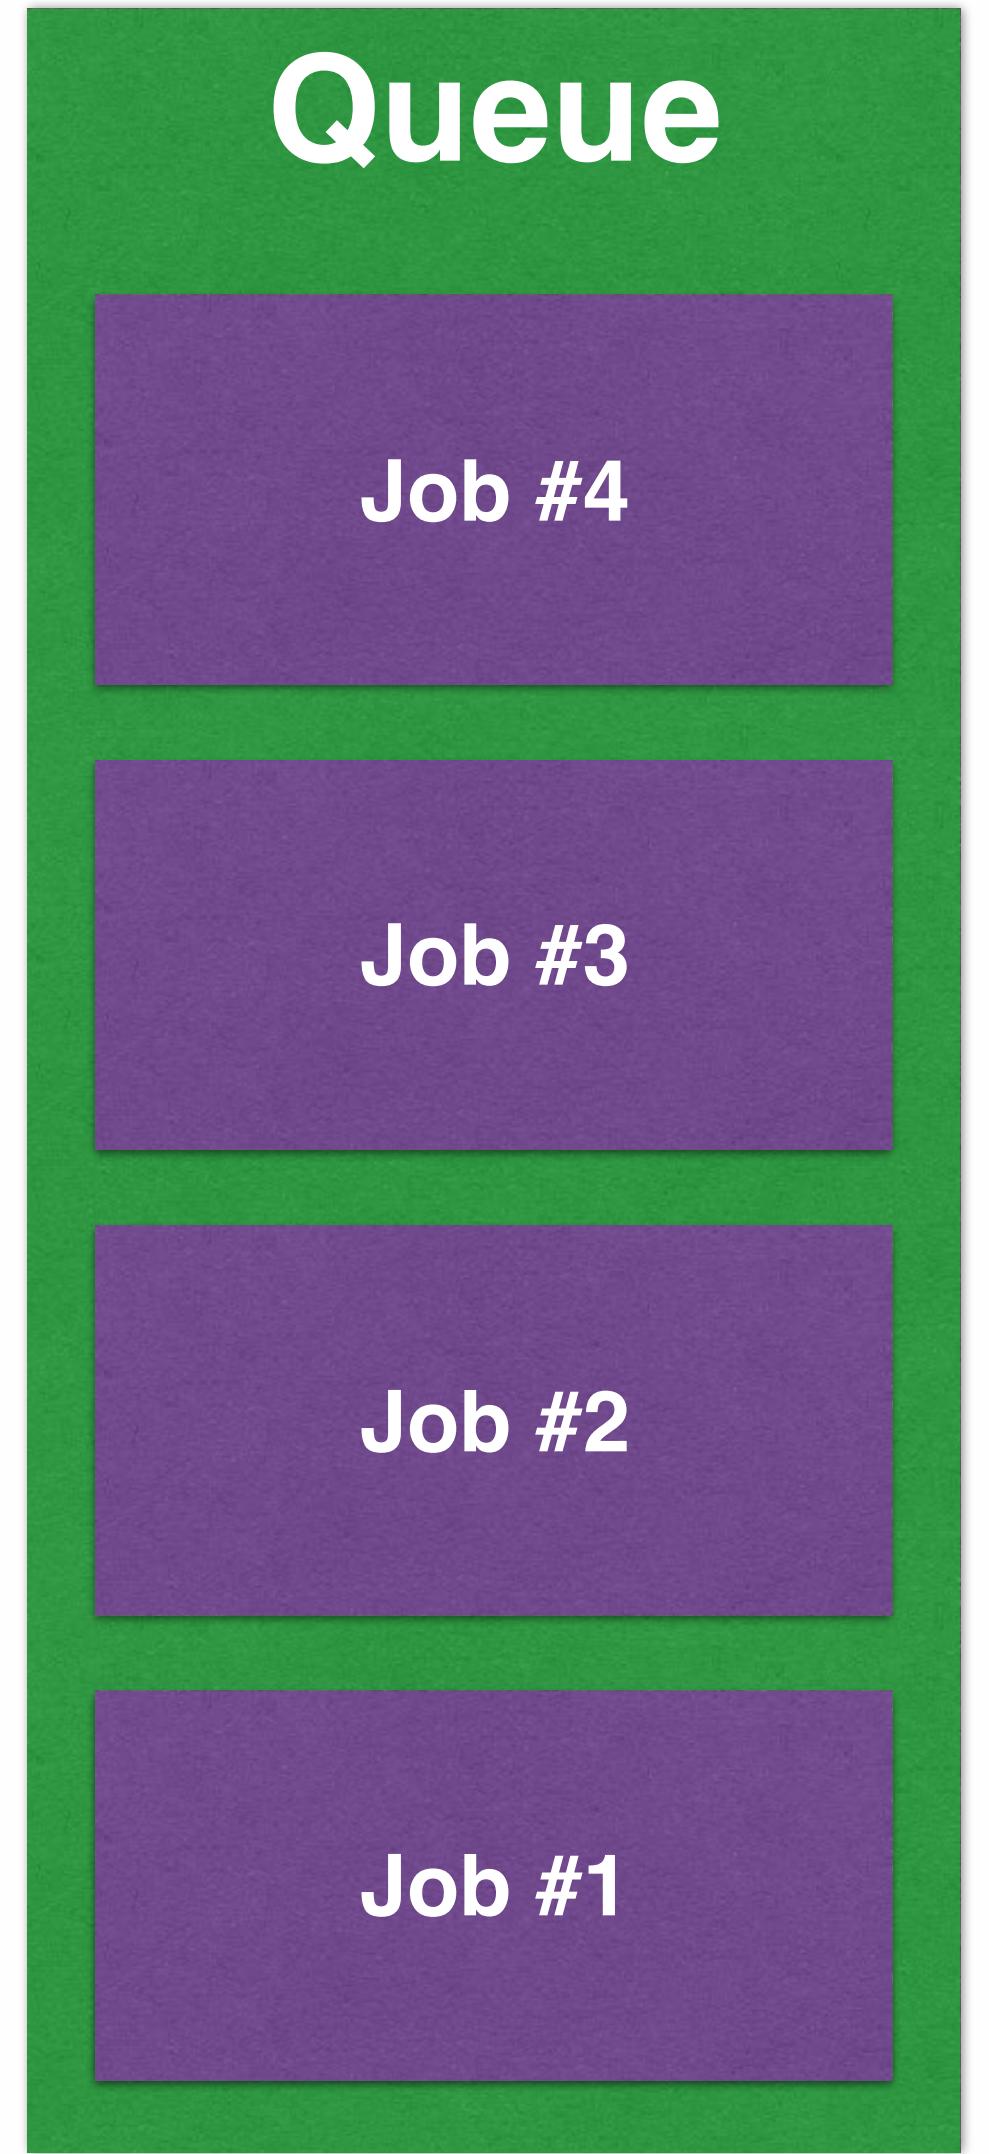
Job #3

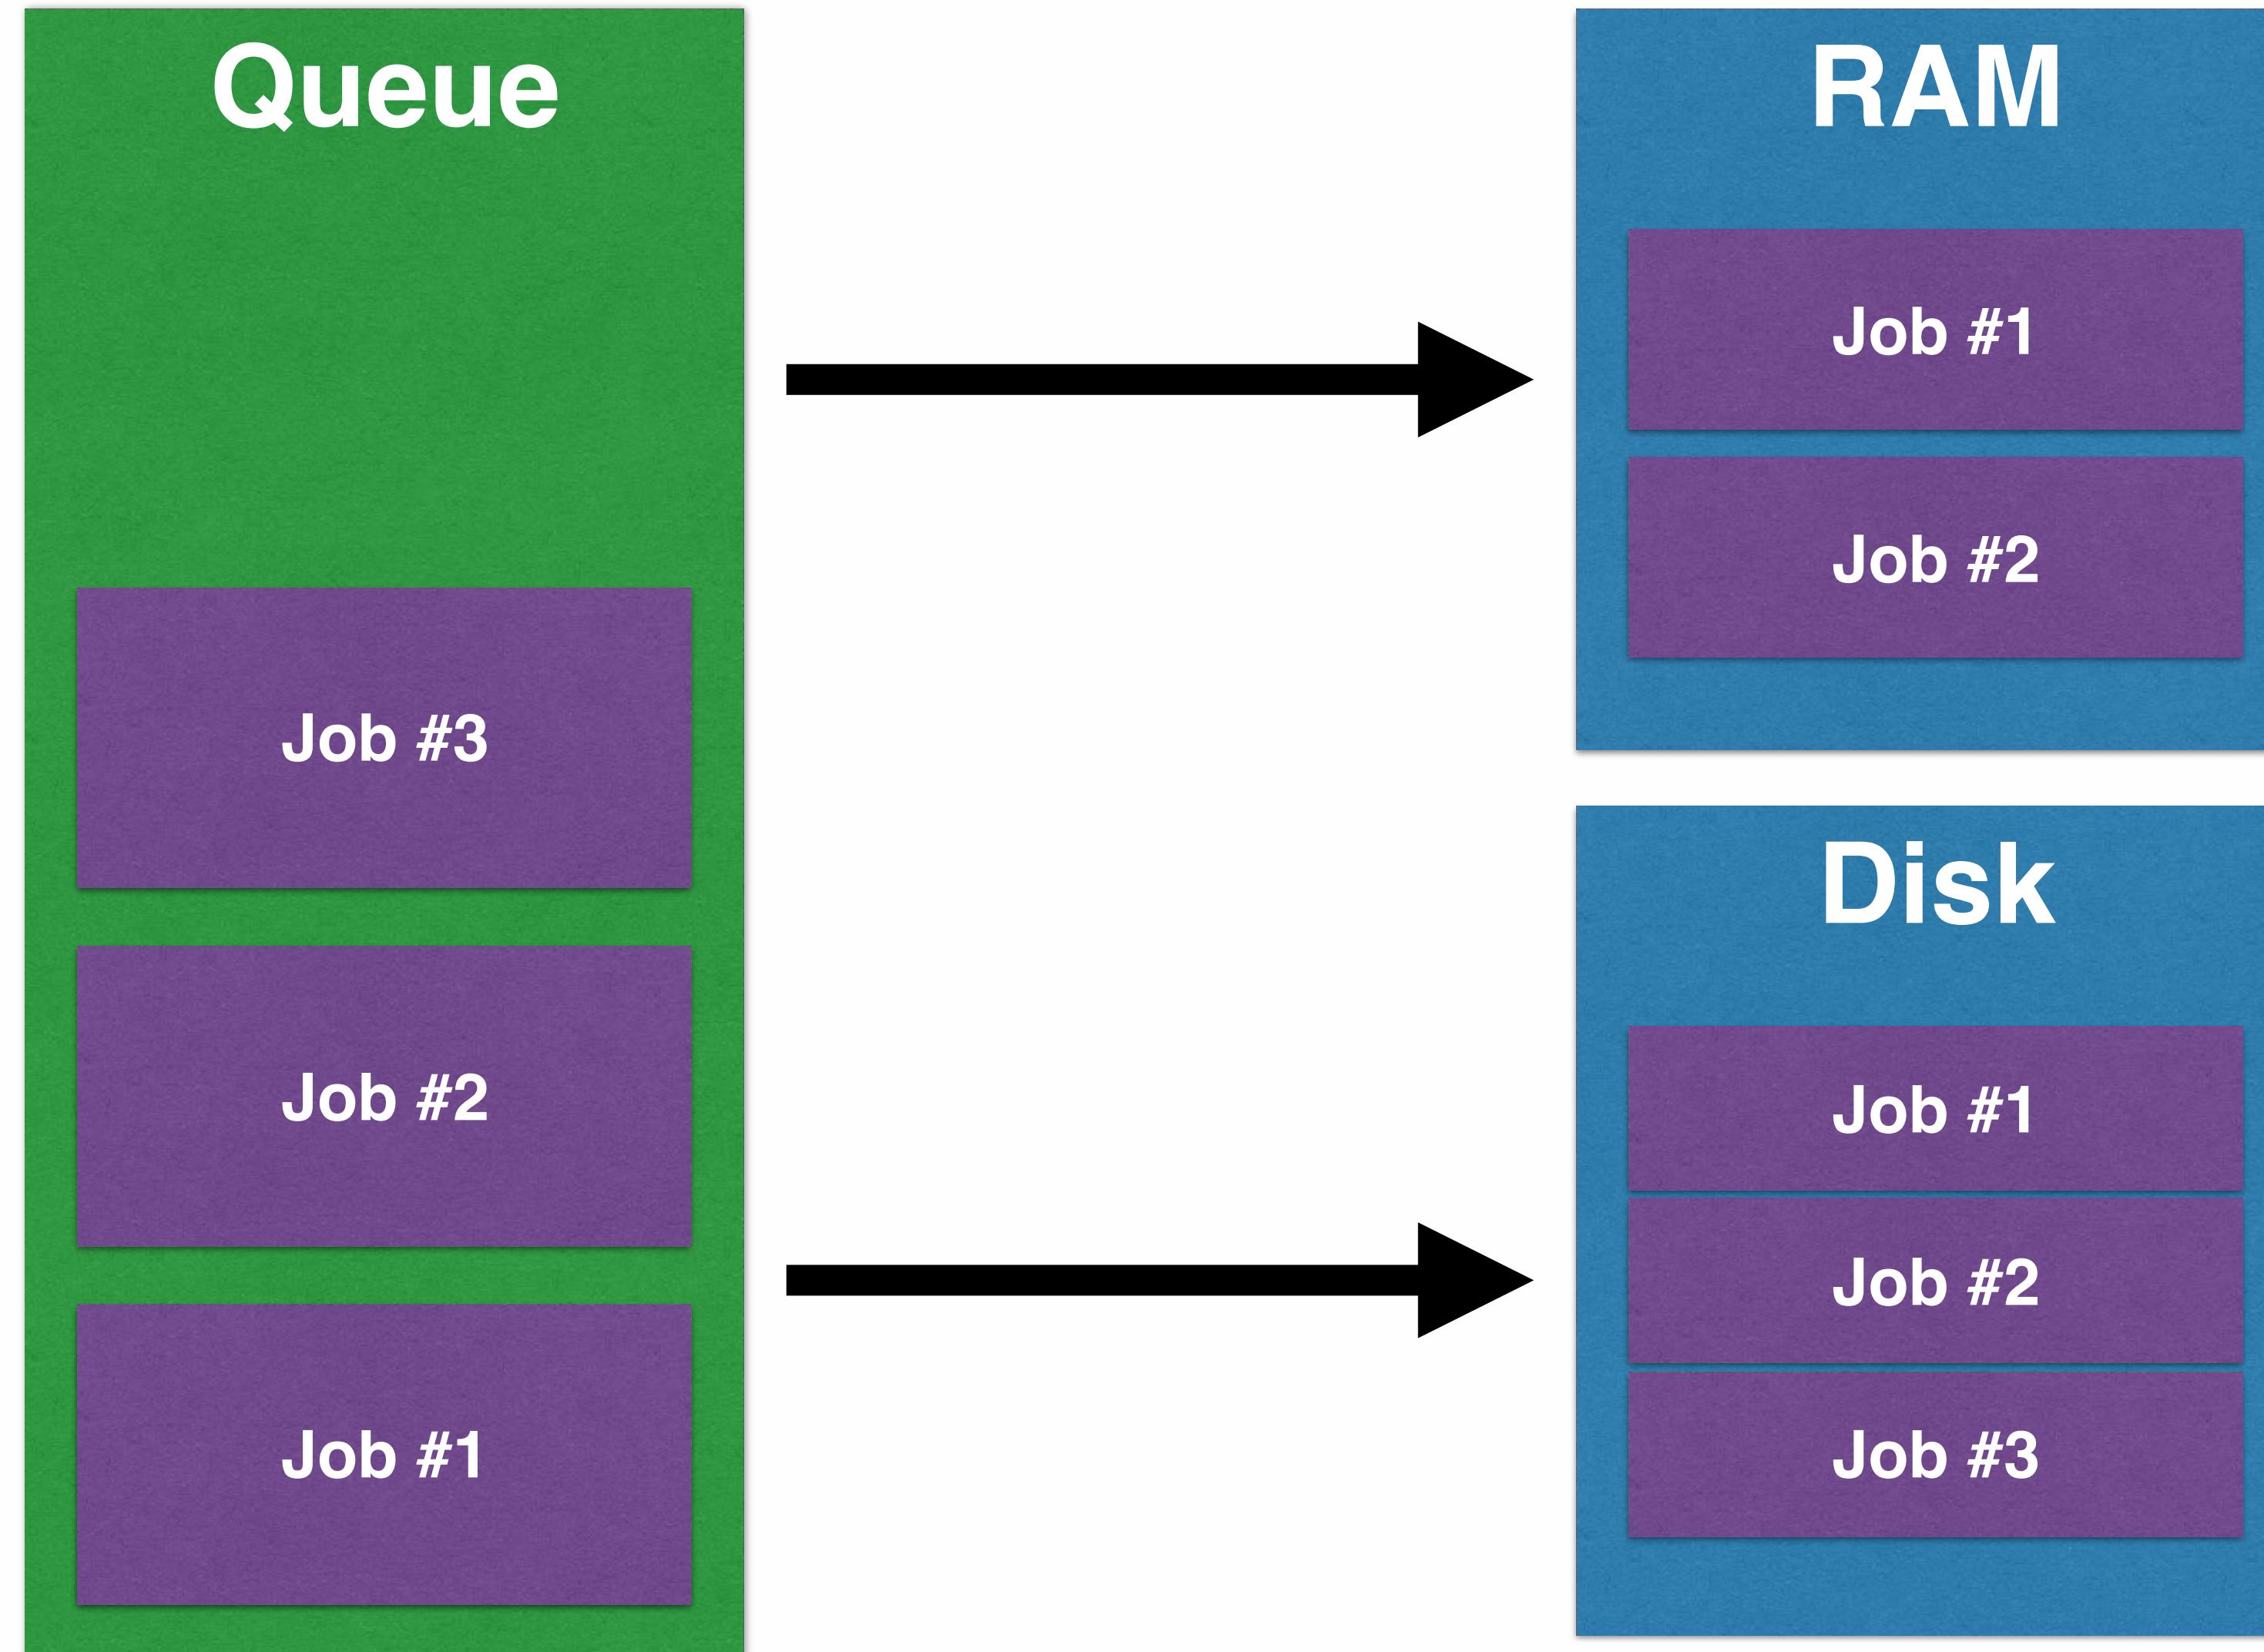
Job #2

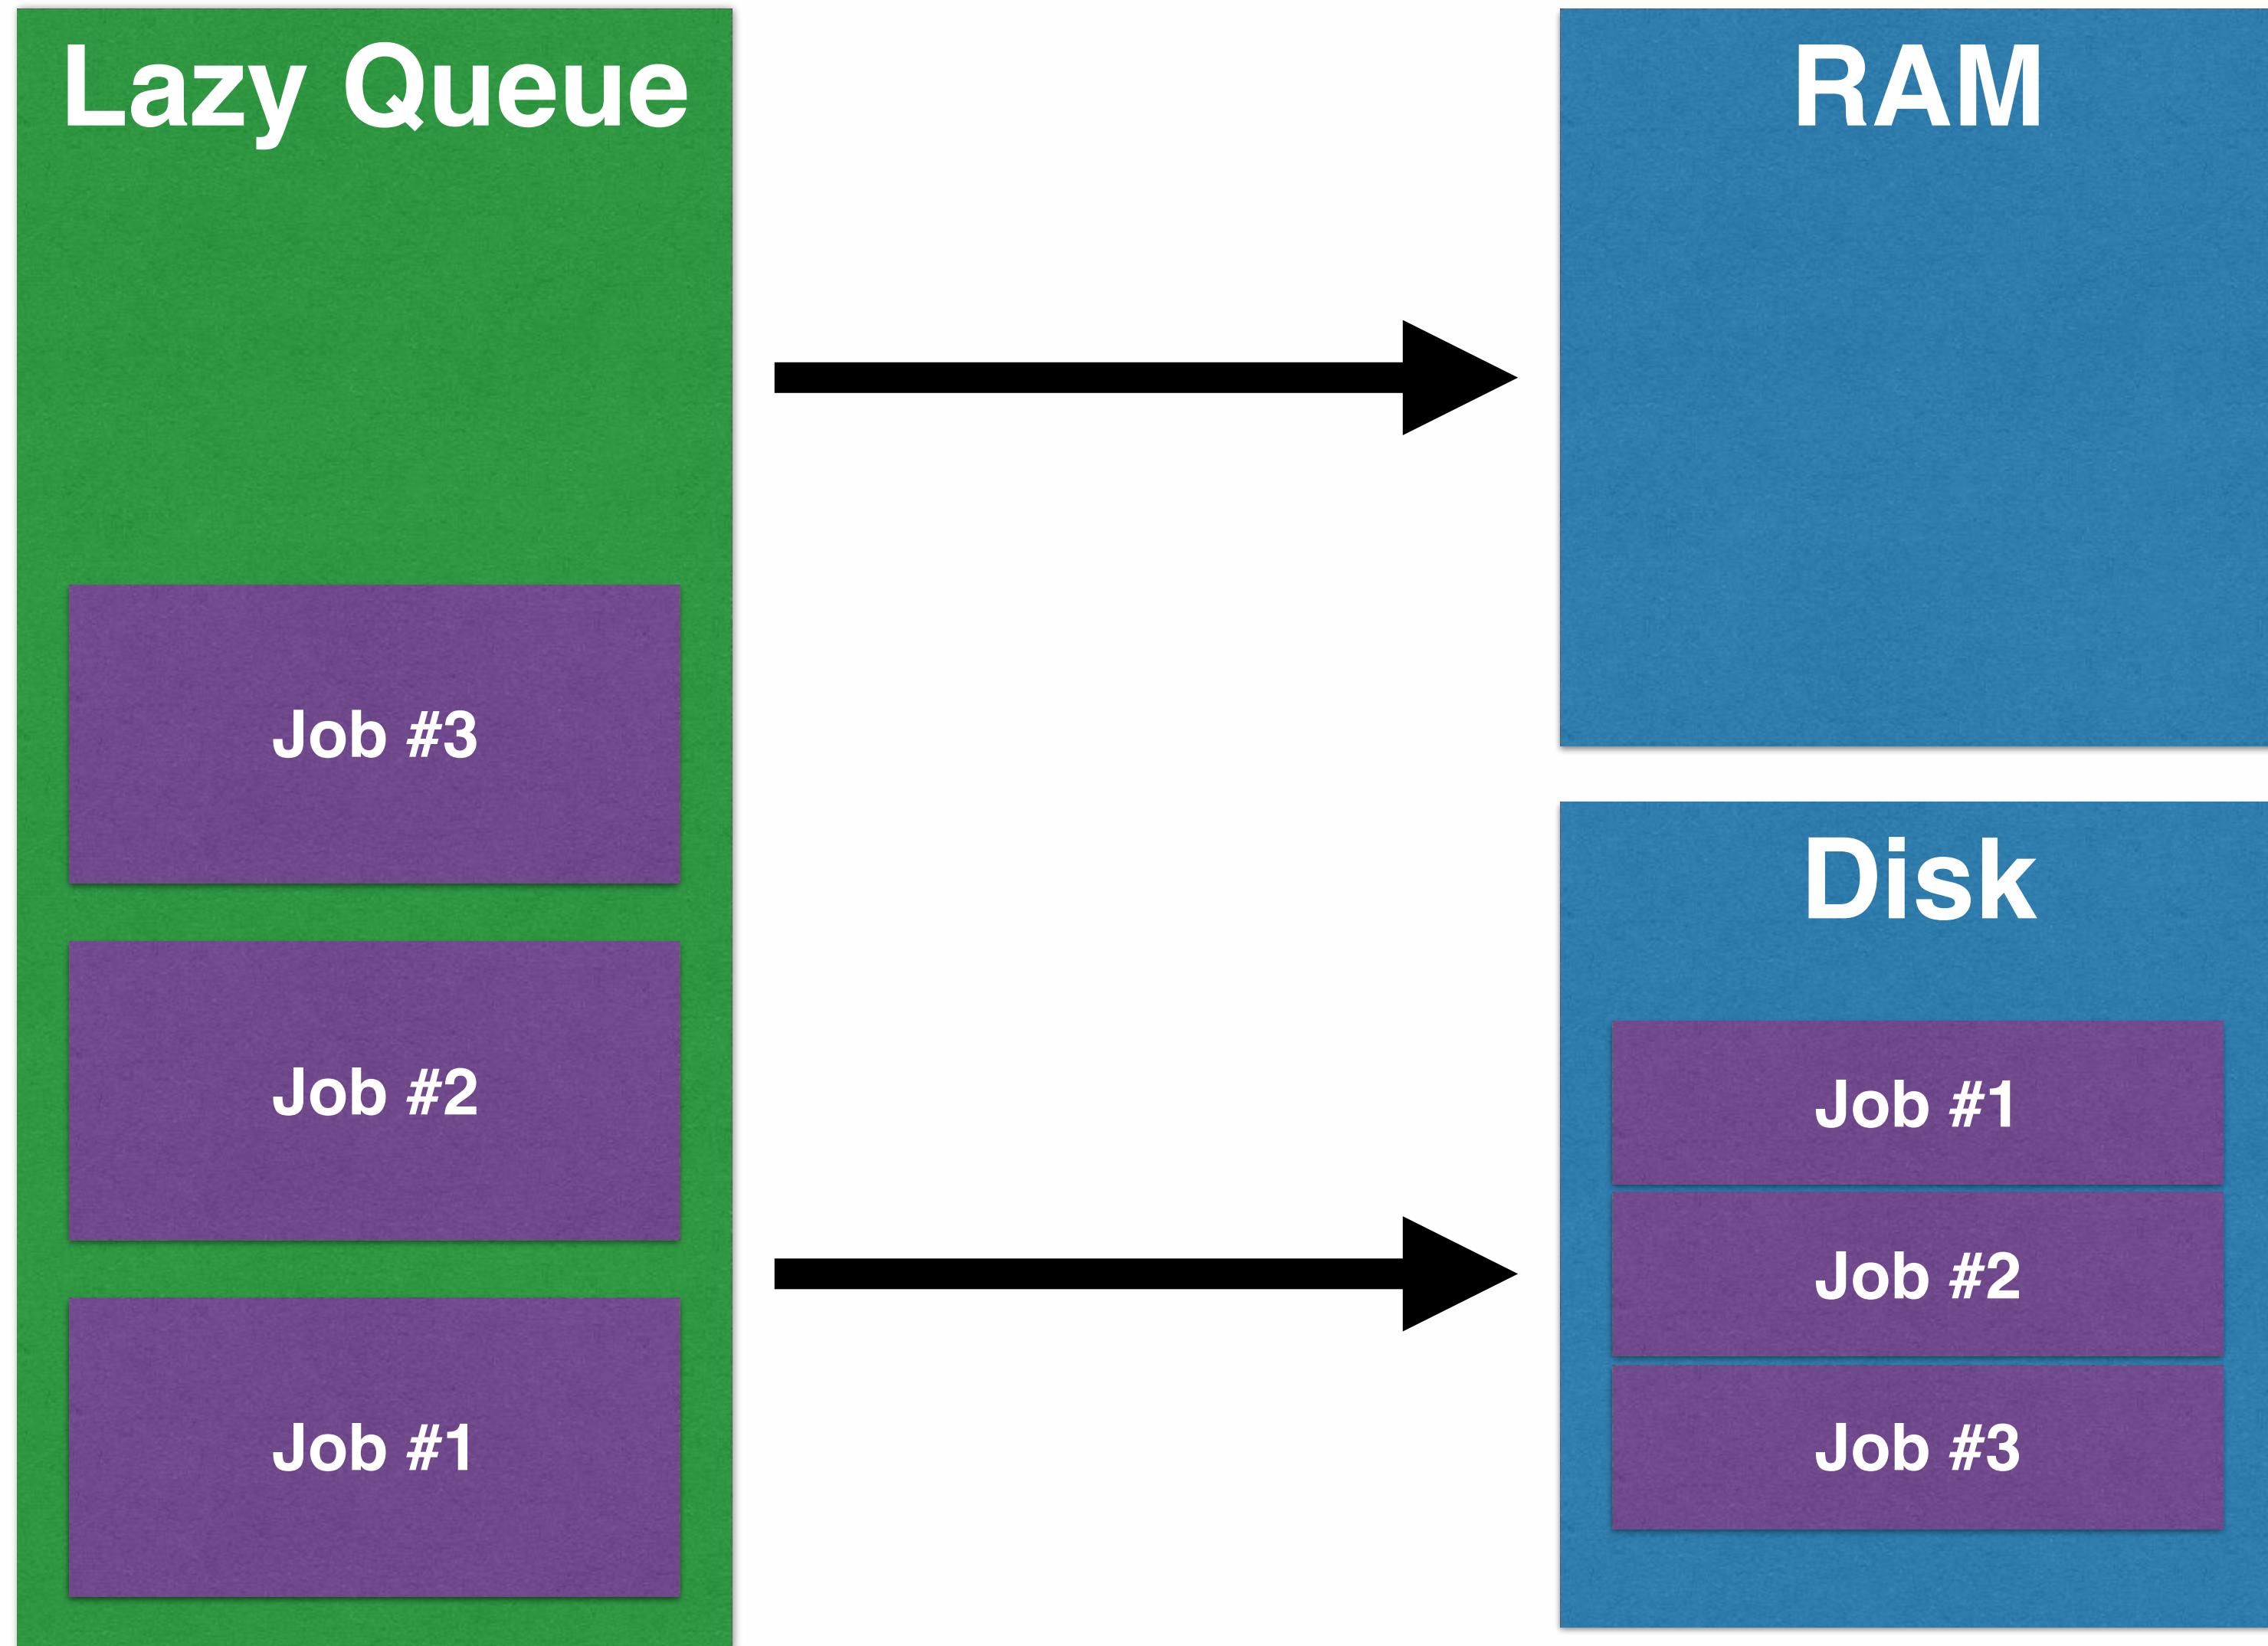
Worker

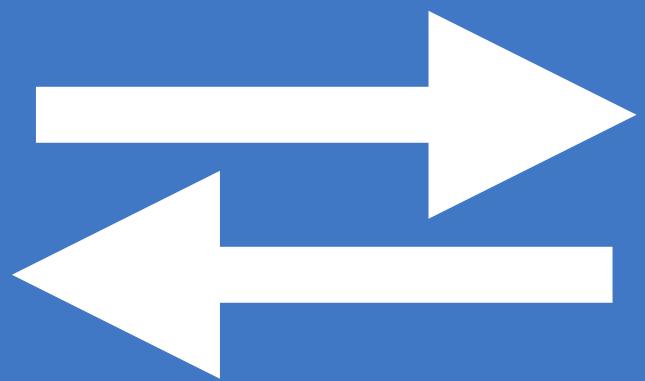
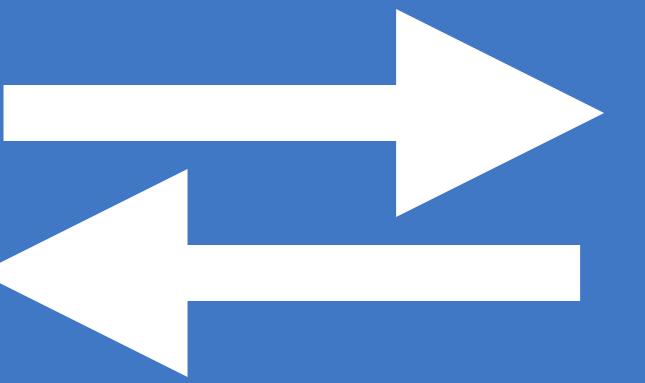


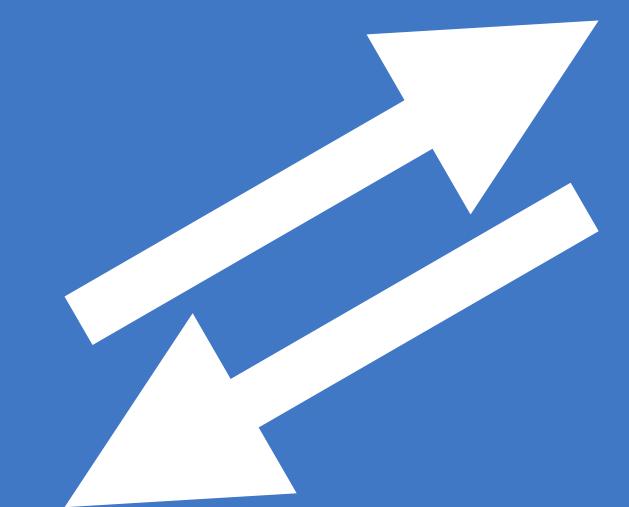
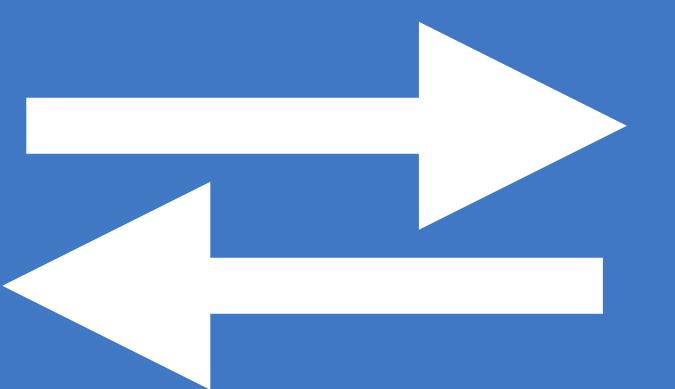
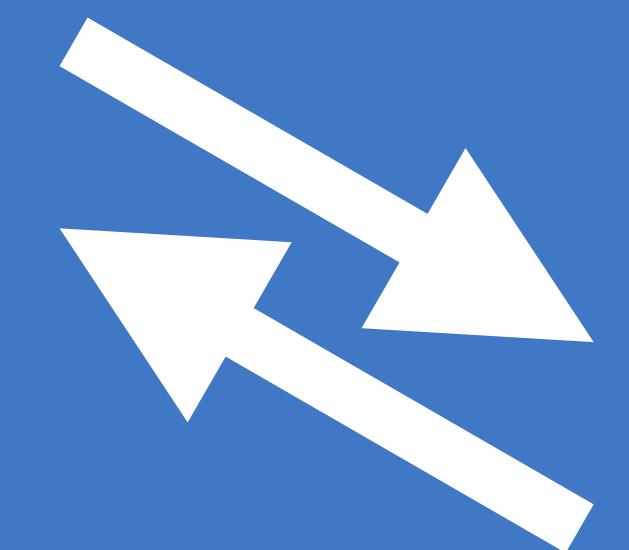












18
Processed**0**
Failed**0**
Busy**0**
Enqueued**0**
Retries**0**
Scheduled**0**
DeadPolling interval: **5 sec**Dashboard History [1 week](#) **1 month** [3 months](#) [6 months](#)

Redis

4.0.8

Version

10

Uptime (days)

30

Connections

1.69M

Memory Usage

3.15M

Peak Memory Usage



Username: *

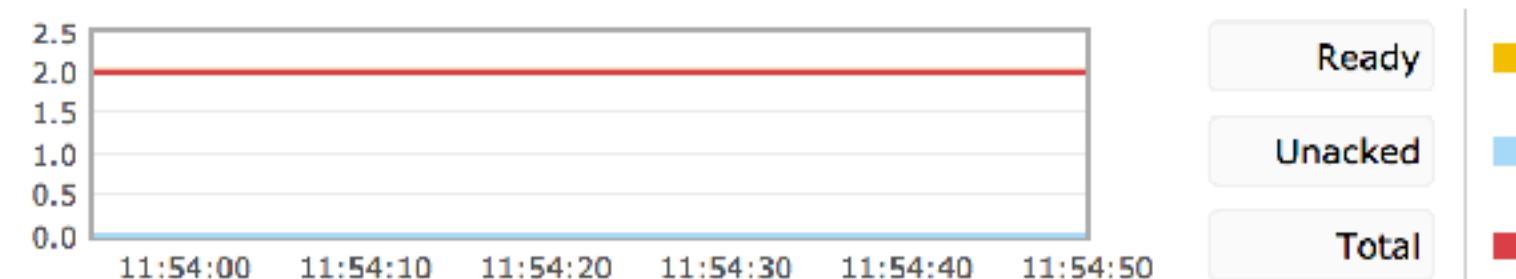
Password: *

WEB UI AUTHORIZATION	-	-	✓
DEDICATED SUPPORT	None	Email	Email
LICENSE	LGPL	Commercial No Custom Terms	Commercial with Custom Terms
PRICING	Free	\$950 / year Unlimited Usage	\$1950 / year per 100 Workers
PURCHASING	n/a	Credit Card	Credit Card, Invoice

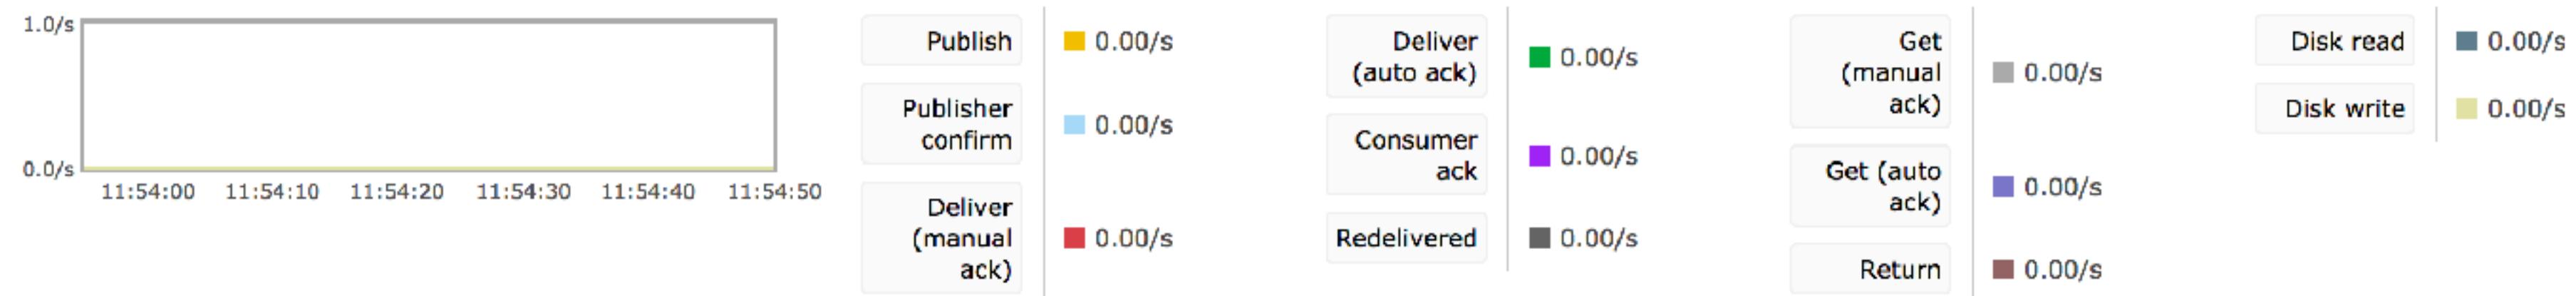
Overview

Totals

Queued messages (chart: last minute) (?)



Message rates (chart: last minute) (?)



Global counts (?)

Connections: 12

Channels: 12

Exchanges: 17

Queues: 8

Consumers: 12

Node

Node: rabbit@46e1483dee20 ([More about this node](#))

File descriptors (?)	Socket descriptors (?)	Erlang processes	Memory	Disk space	Rates mode	Info	Reset stats DB	+-
39 1048576 available	12 943626 available	474 1048576 available	77MB 800MB high watermark	25GB 48MB low watermark	basic	Disc 1	Reset	

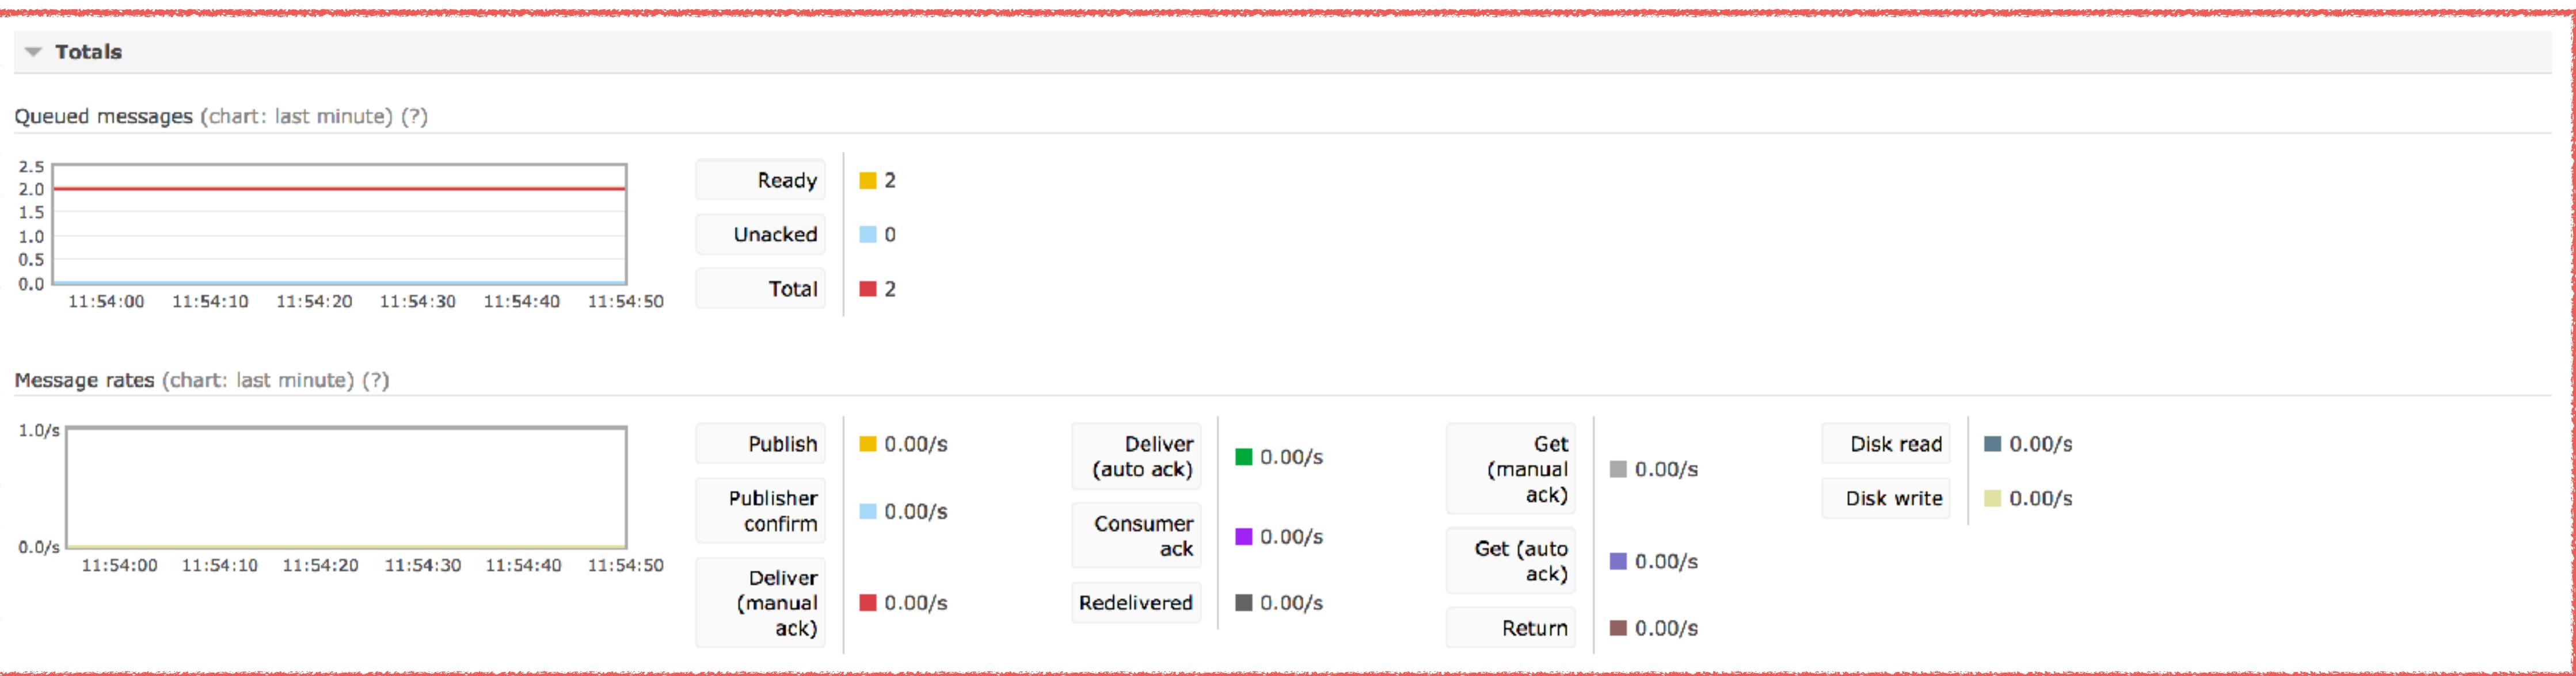
[Reset stats on all nodes](#)

Paths

Config file	/etc/rabbitmq/rabbitmq.config
Database directory	/var/lib/rabbitmq/mnesia/rabbit@46e1483dee20
Log file	tty
SASL log file	tty

[Ports and contexts](#)

Overview



Global counts (?)

Connections: 12

Channels: 12

Exchanges: 17

Queues: 8

Consumers: 12

Node

Node: rabbit@46e1483dee20 ([More about this node](#))

File descriptors (?)	Socket descriptors (?)	Erlang processes	Memory	Disk space	Rates mode	Info	Reset stats DB	+-
39 1048576 available	12 943626 available	474 1048576 available	77MB 800MB high watermark	25GB 48MB low watermark	basic	Disc 1	Reset	

[Reset stats on all nodes](#)

Paths

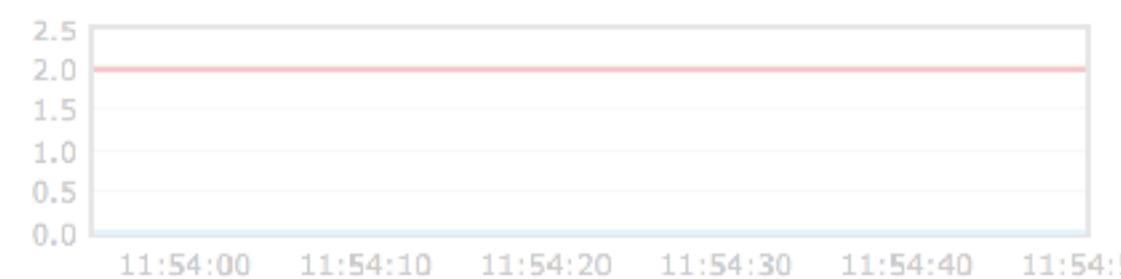
Config file	/etc/rabbitmq/rabbitmq.config
Database directory	/var/lib/rabbitmq/mnesia/rabbit@46e1483dee20
Log file	tty
SASL log file	tty

[Ports and contexts](#)

Overview

▼ Totals

Queued messages (chart: last minute) (?)



Message rates (chart: last minute) (?)



Global counts (?)

Connections: 12

Channels: 12

Exchanges: 17

Queues: 8

Consumers: 12

▼ Node

Node: rabbit@46e1483dee20 ([More about this node](#))

File descriptors (?)	Socket descriptors (?)	Erlang processes	Memory	Disk space	Rates mode	Info	Reset stats DB	+-
39 1048576 available	12 943626 available	474 1048576 available	77MB 800MB high watermark	25GB 48MB low watermark	basic	Disc 1	Reset	

[Reset stats on all nodes](#)

Paths

Config file	/etc/rabbitmq/rabbitmq.config
Database directory	/var/lib/rabbitmq/mnesia/rabbit@46e1483dee20
Log file	tty
SASL log file	tty

▶ Ports and contexts

Queues

All queues (8)

Pagination

Page 1 of 1 - Filter: Regex (?)(?)

Displaying 8 items , page size up to:

Overview			Messages			Message bytes		Message rates		
Name	Features	State	Ready	Unacked	Total	In Memory	Persistent	Incoming	deliver / get	ack
chatbot_messages	D DLX DLK	idle	0	0	0	0B	0B			
chatbot_messages.error	D	idle	0	0	0	0B	0B			
default	D DLX	running	0	0	0	0B	0B	0.00/s	0.00/s	0.00/s
default-error	D	idle	2	0	2	2.0kB	2.0kB	0.00/s		
default-retry	D TTL DLX	idle	0	0	0	0B	0B			
mailers	D DLX	idle	0	0	0	0B	0B	0.00/s	0.00/s	0.00/s
mailers-error	D	idle	0	0	0	0B	0B			
mailers-retry	D TTL DLX	idle	0	0	0	0B	0B			

+/-

[Add a new queue](#)
[HTTP API](#) | [Command Line](#)

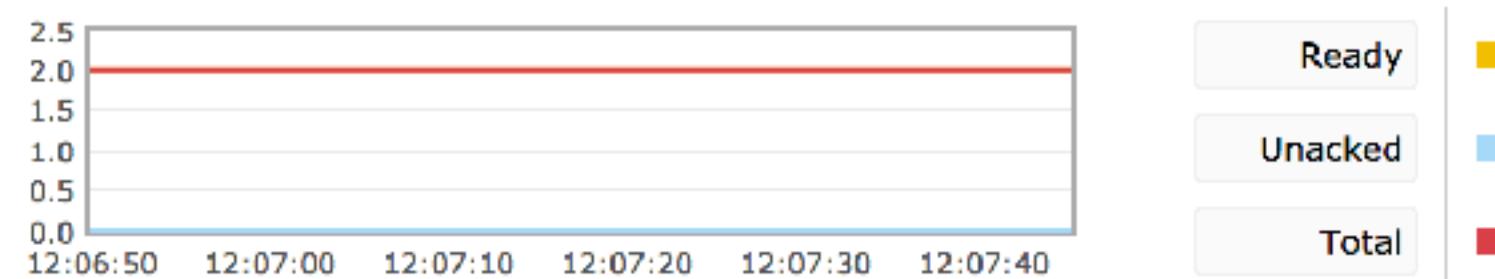
Update

Last update: 2018-02-27 12:04:46

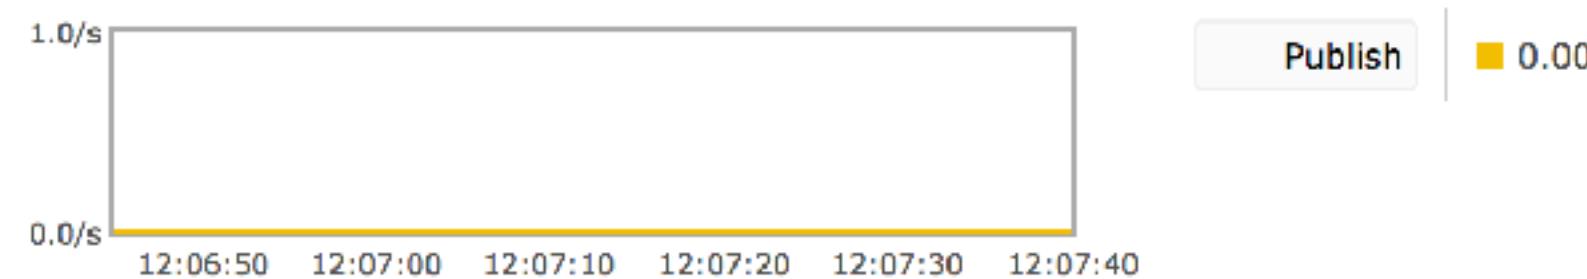
Queue default-error

Overview

Queued messages (chart: last minute) (?)



Message rates (chart: last minute) (?)



Details

Features	durable: <code>true</code>	State	idle	Messages (?)	Total	Ready	Unacked	In memory	Persistent	Transient, Paged Out
Policy		Consumers		0	2	2	0	2	2	0
Consumer utilisation (?)		0%		Message body bytes (?)	2.0kB	2.0kB	0B	2.0kB	2.0kB	0B
				Process memory (?)	15kB					

▶ Consumers

▶ Bindings

▶ Publish message

▶ Get messages

▶ Move messages

▶ Delete

▶ Purge

▼ Runtime Metrics (Advanced)

Reductions (per second) (chart: last minute) (?)



User:

[Log out](#)

Cluster: rabbit@46e1483dee20 ([change](#))

RabbitMQ 3.6.10, Erlang 19.3

[Overview](#)

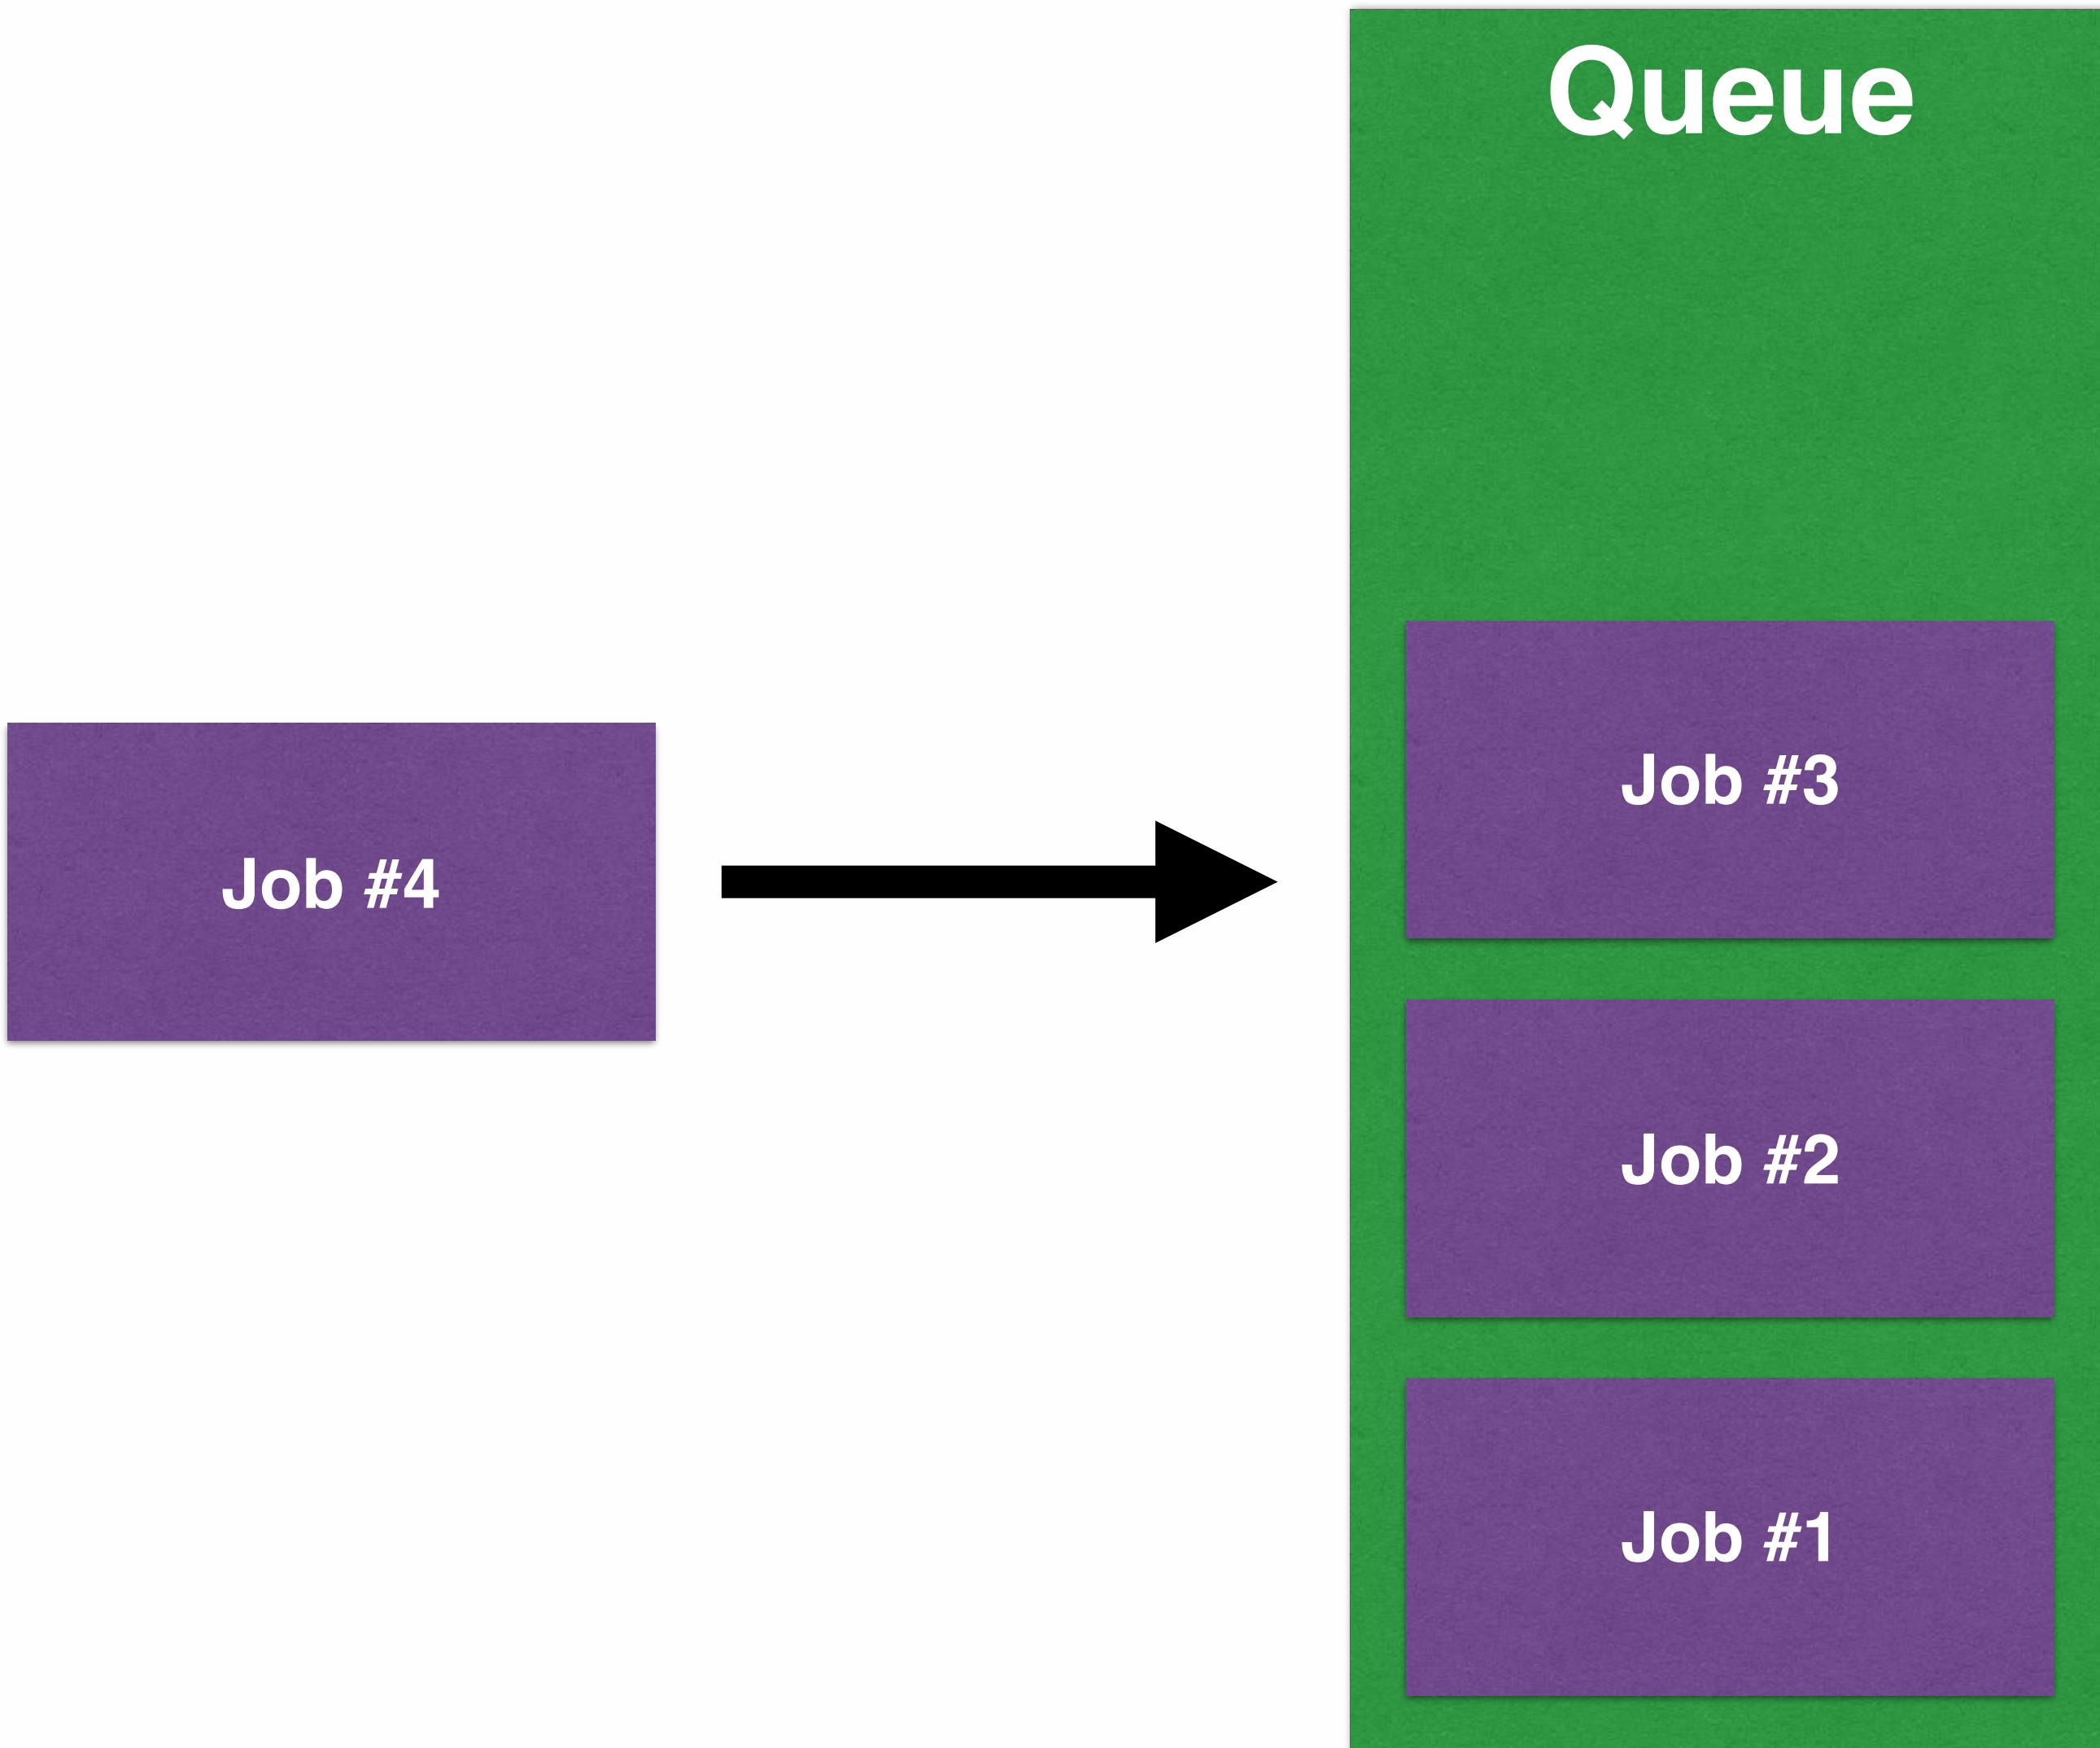
[Connections](#)

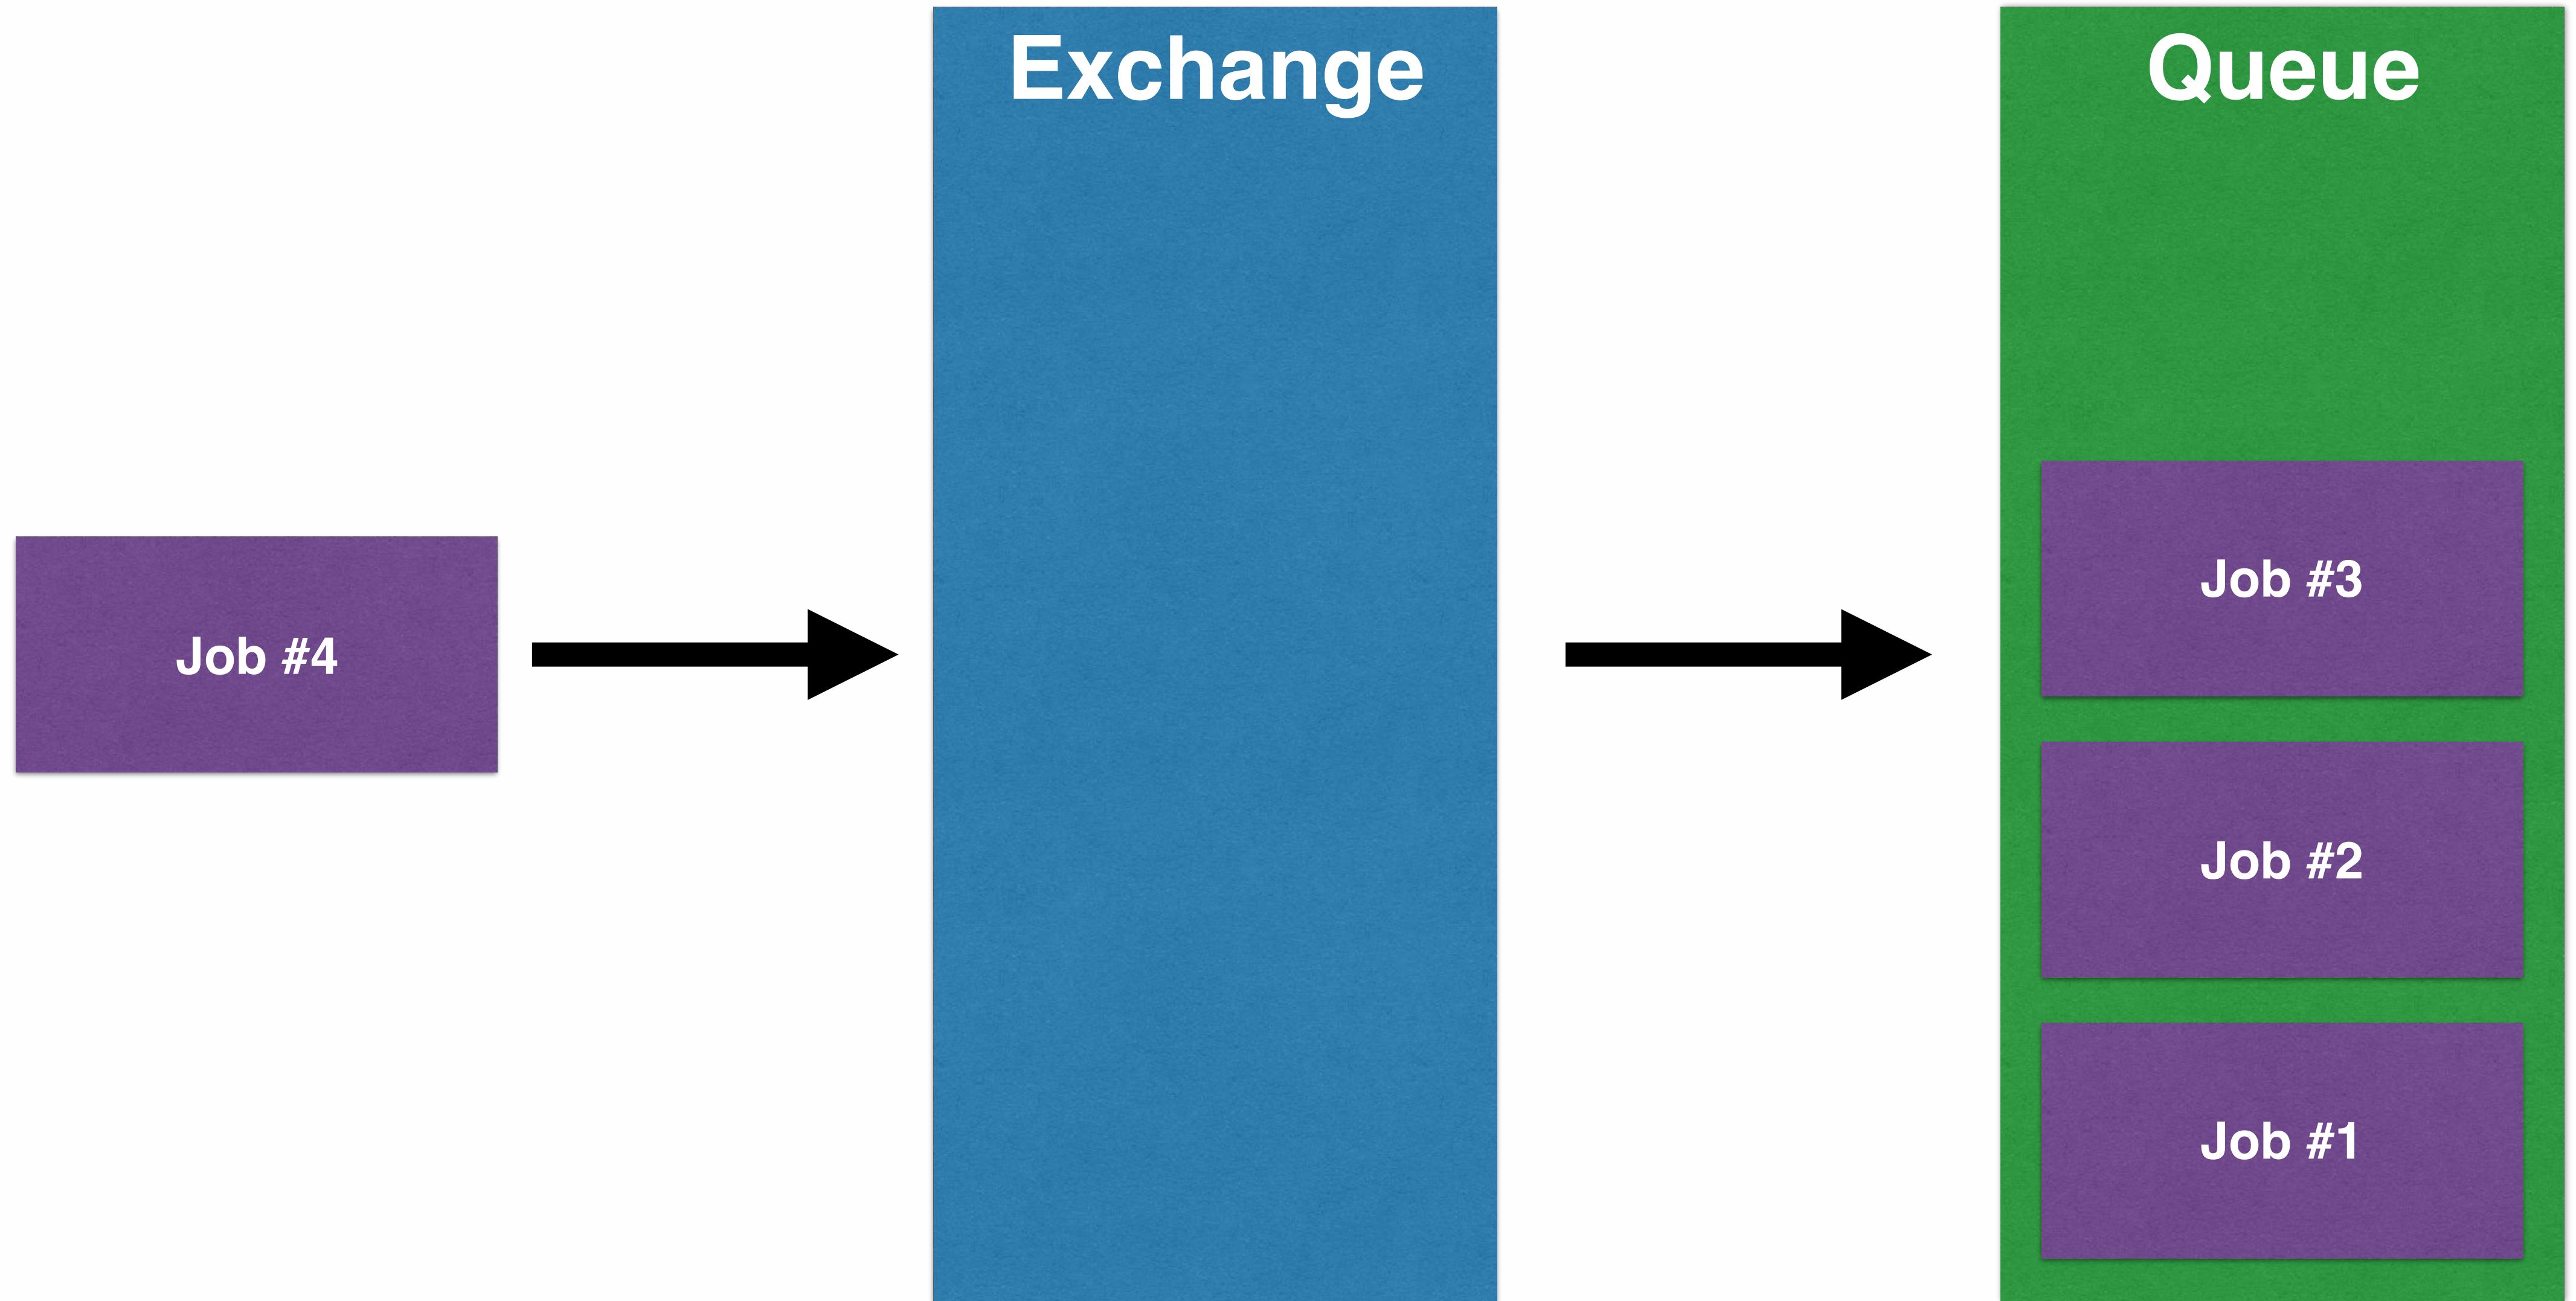
[Channels](#)

[Exchanges](#)

[Queues](#)

[Admin](#)





Direct Exchange

Job #4



Queue

Job #3

Job #2

Job #1

Ruby Chat Exchange

Message #4



Ruby Chat

Message #3

Message #2

Message #1

Fanout Exchange



Queue #1

Job #1



Queue #2

Job #2

Job #3

**Awesome
languages
Exchange**



Ruby Chat

Message #1

Message #3



Rust Chat

Message #2

Topic Exchange



Queue A

Job #1
Topic A

Job #3
Topic A



Queue B

Job #2
Topic B

Chat Exchange



Chat

Message #1

Text

Message #3

Video

Video Chat

Message #2

Video

Header Exchange



Queue A

Job #1

Job #3



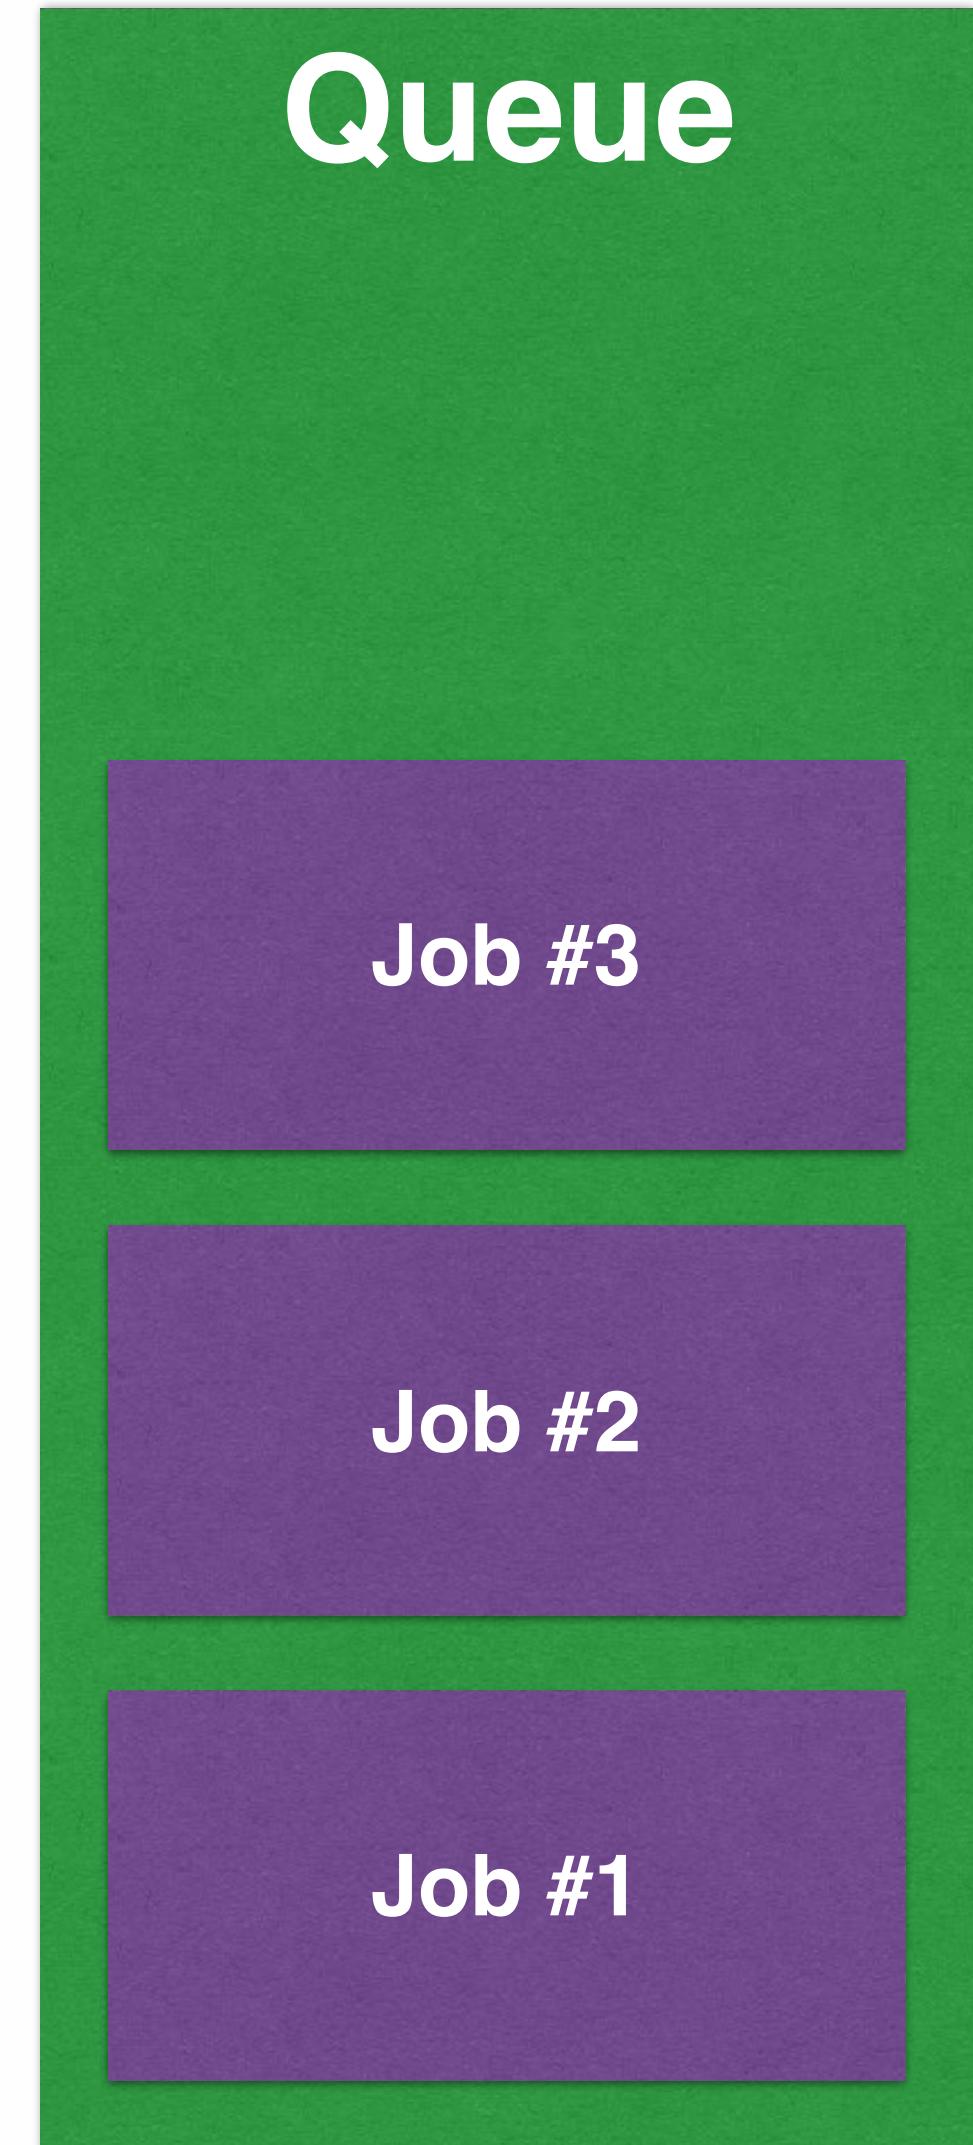
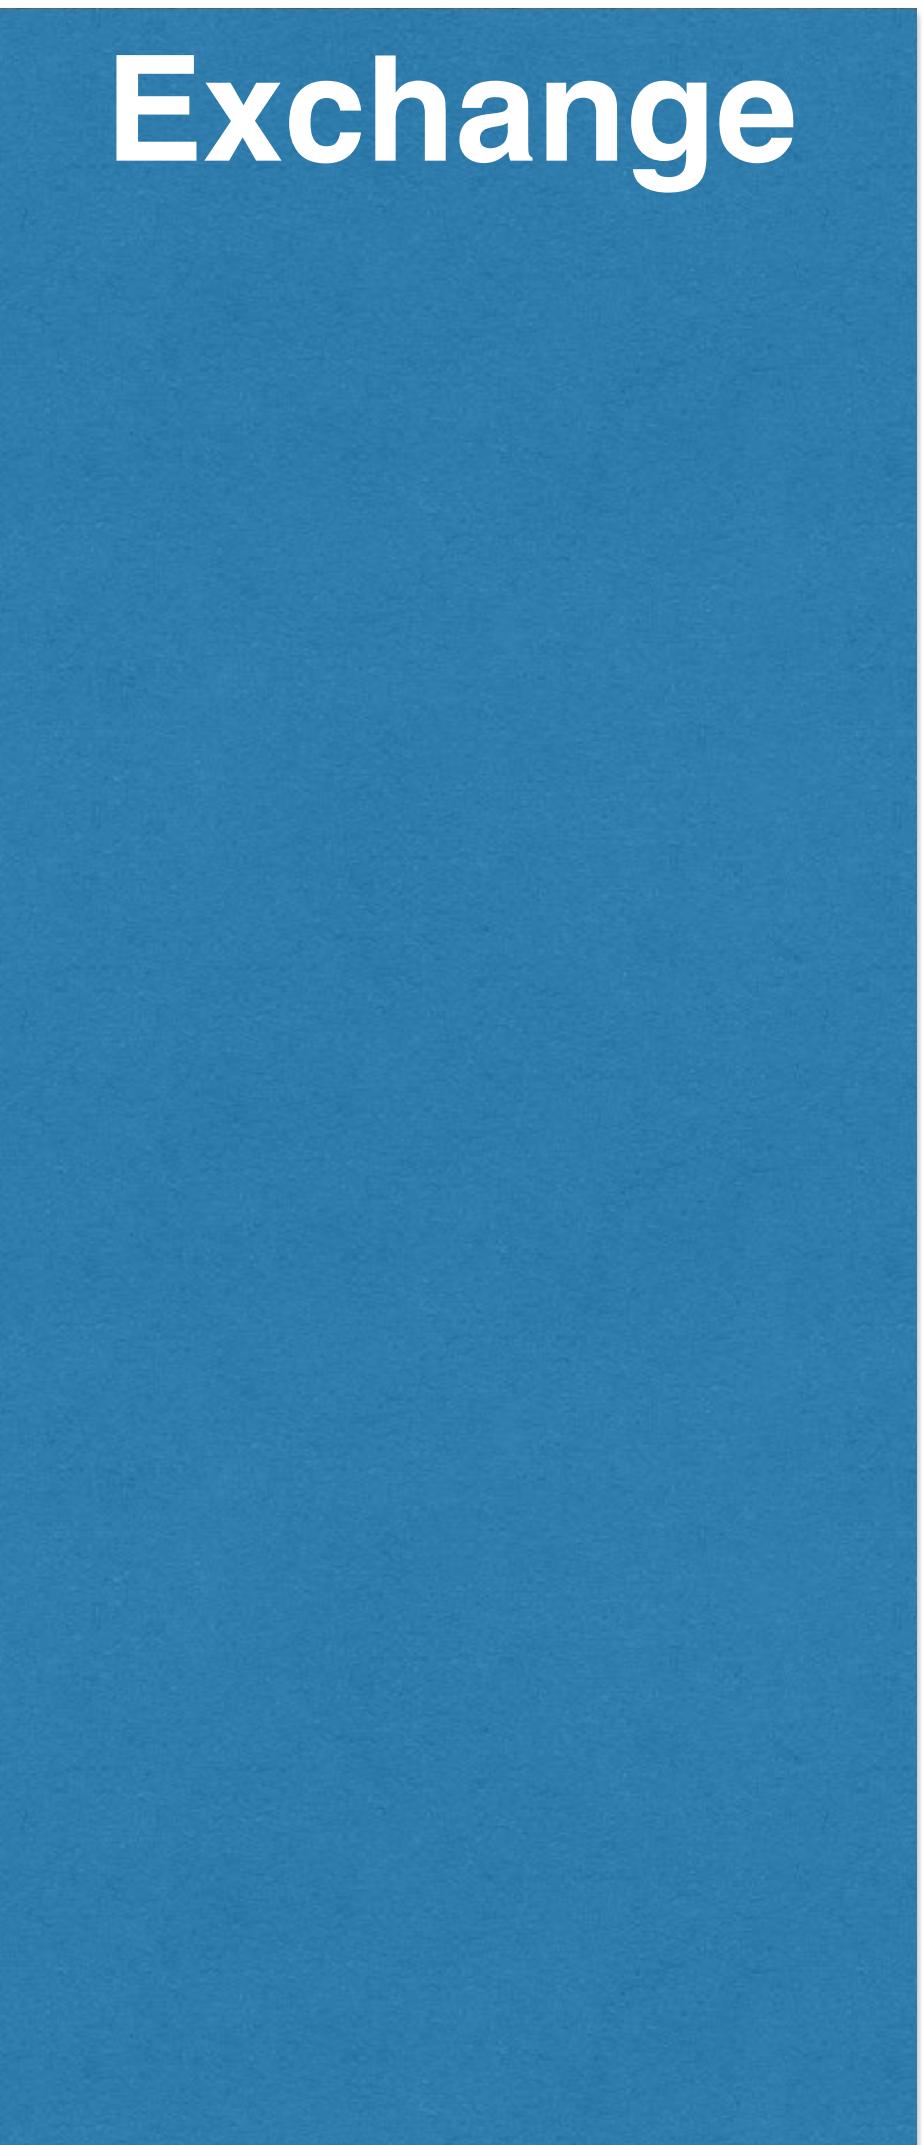
Queue B

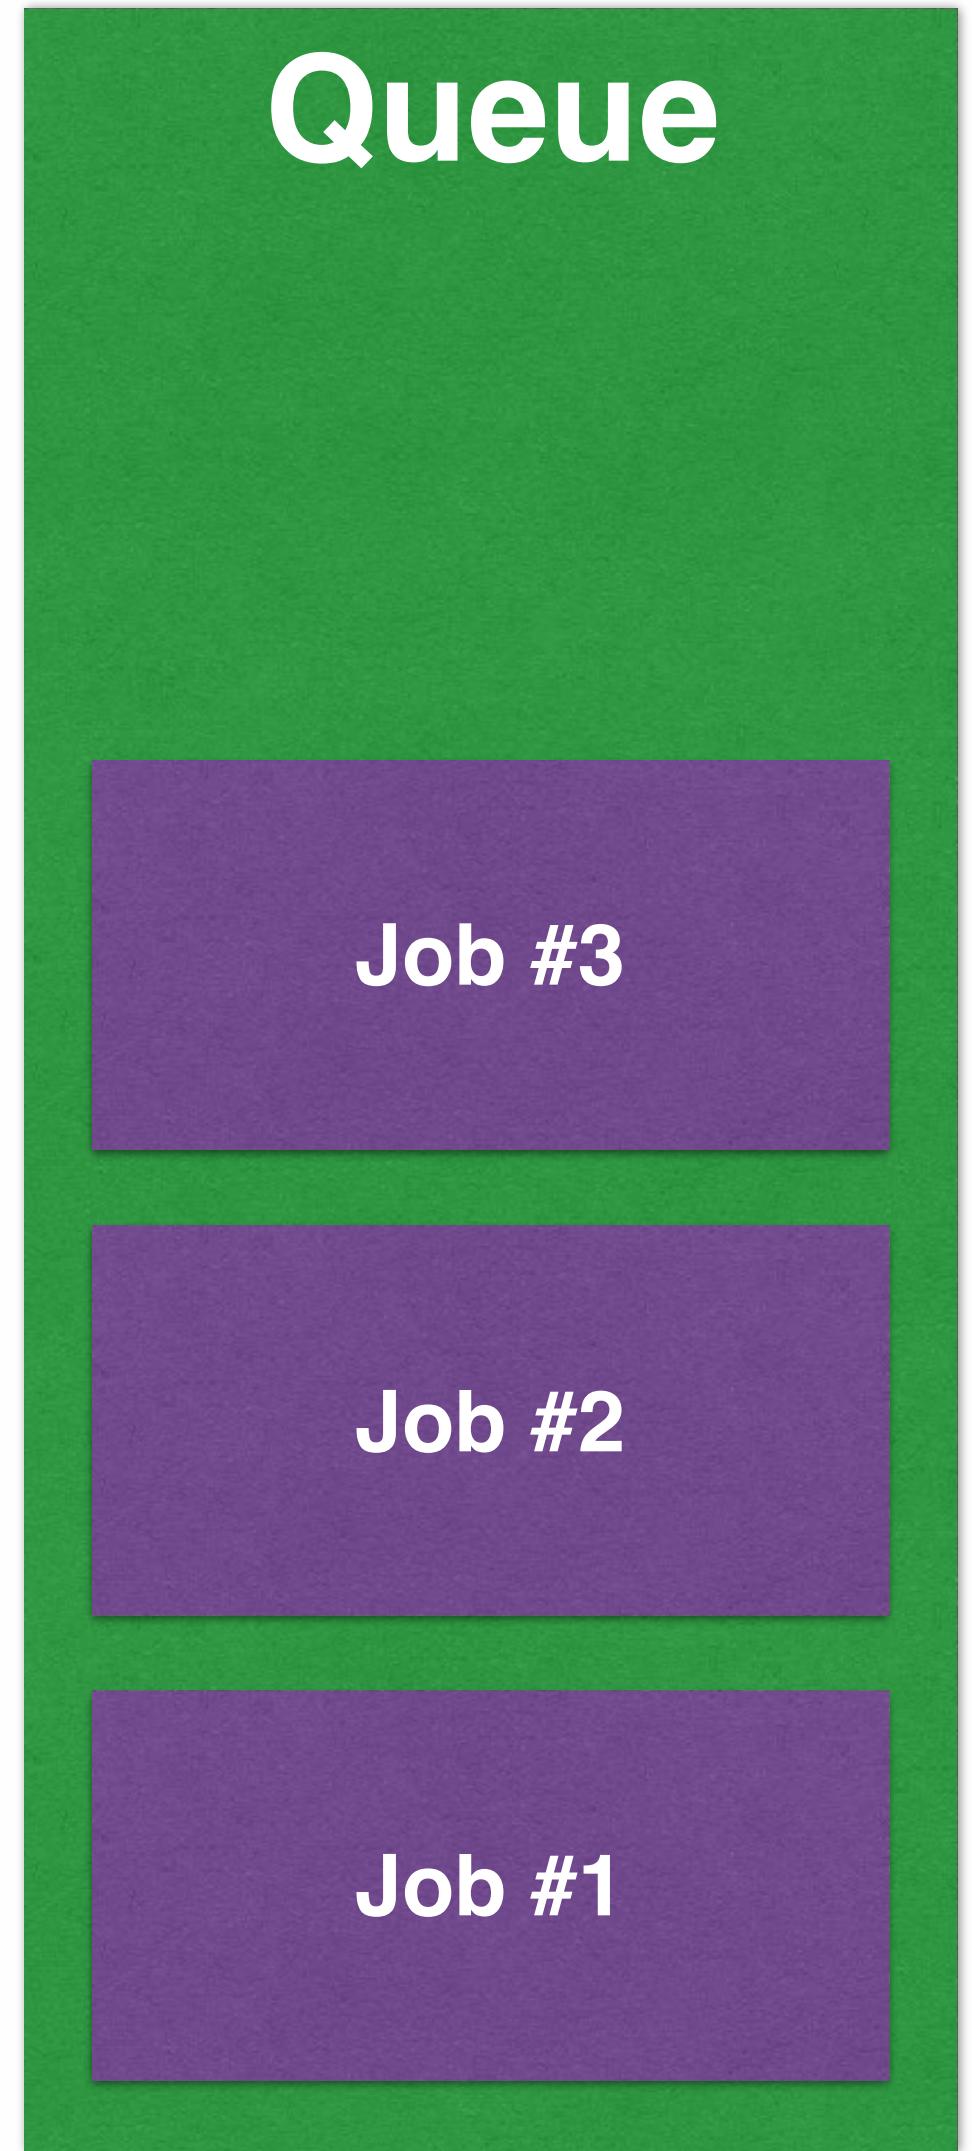
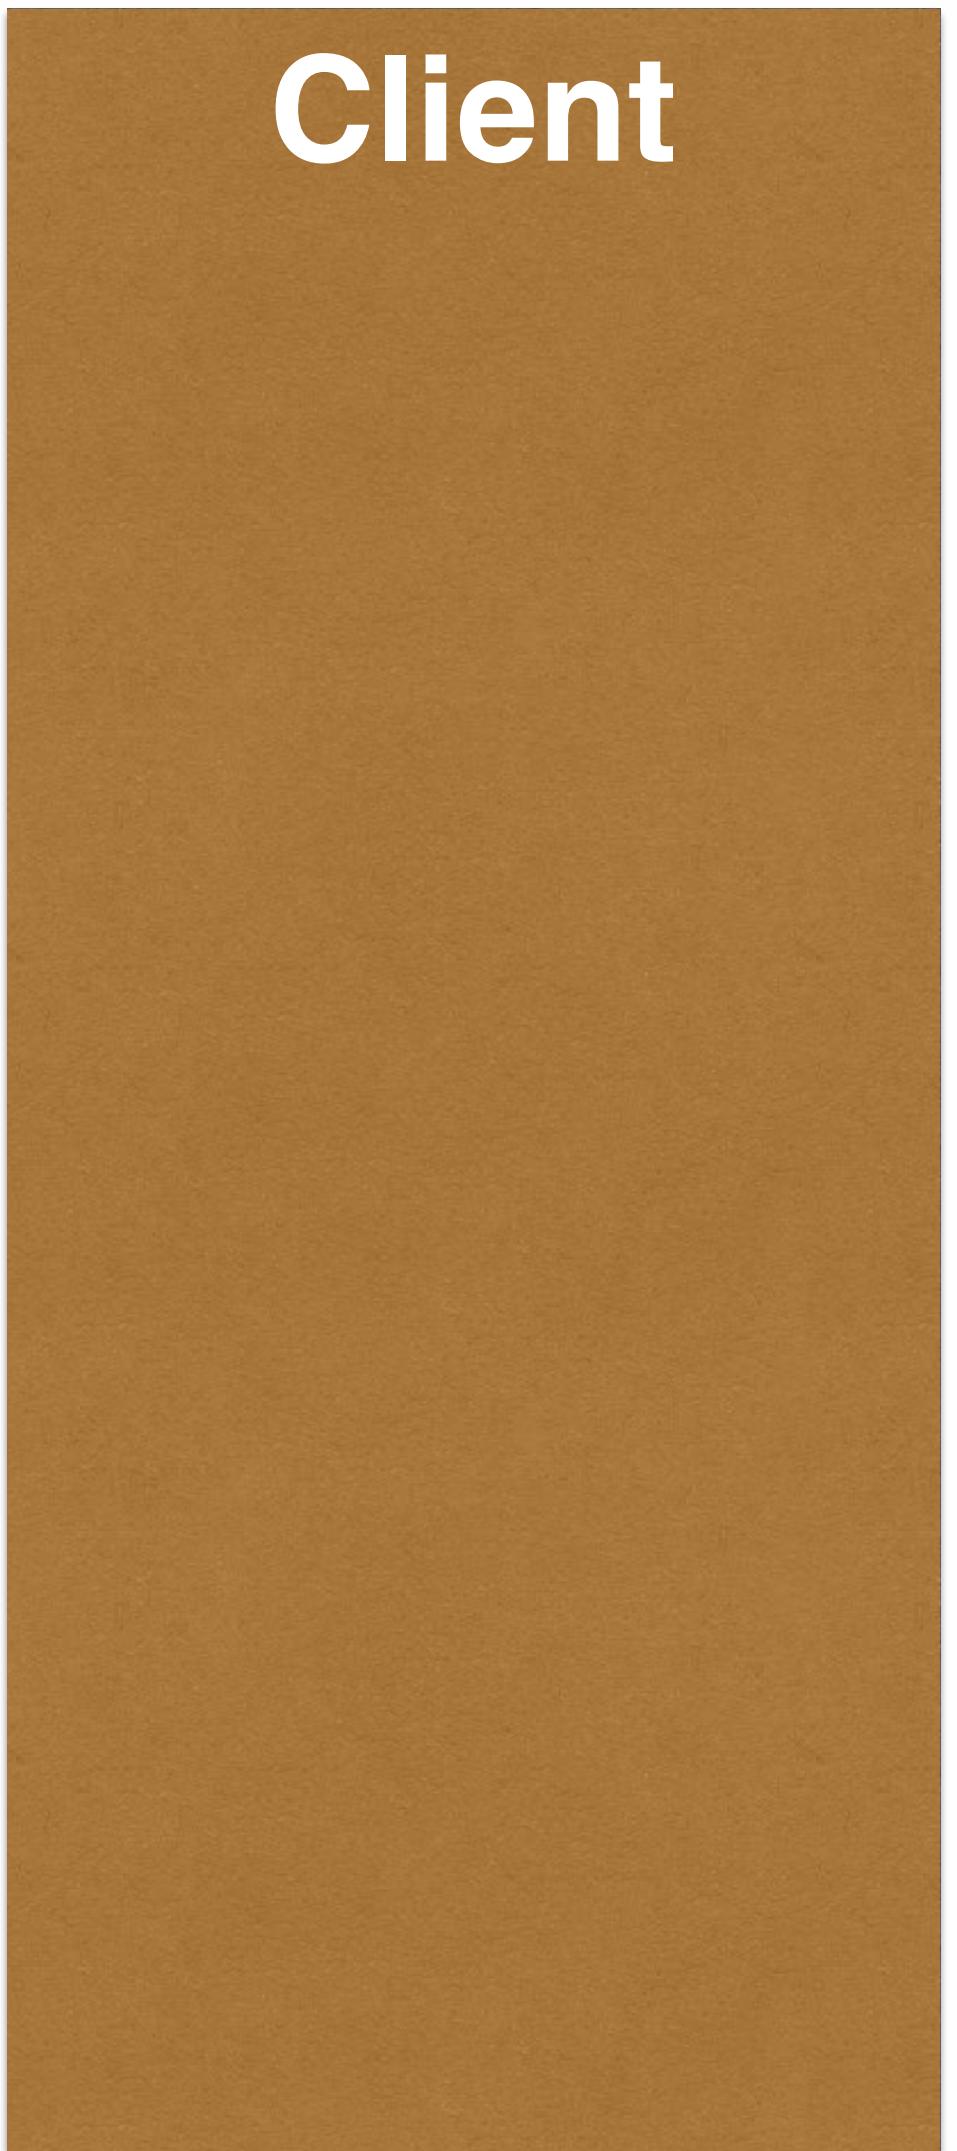
Job #2

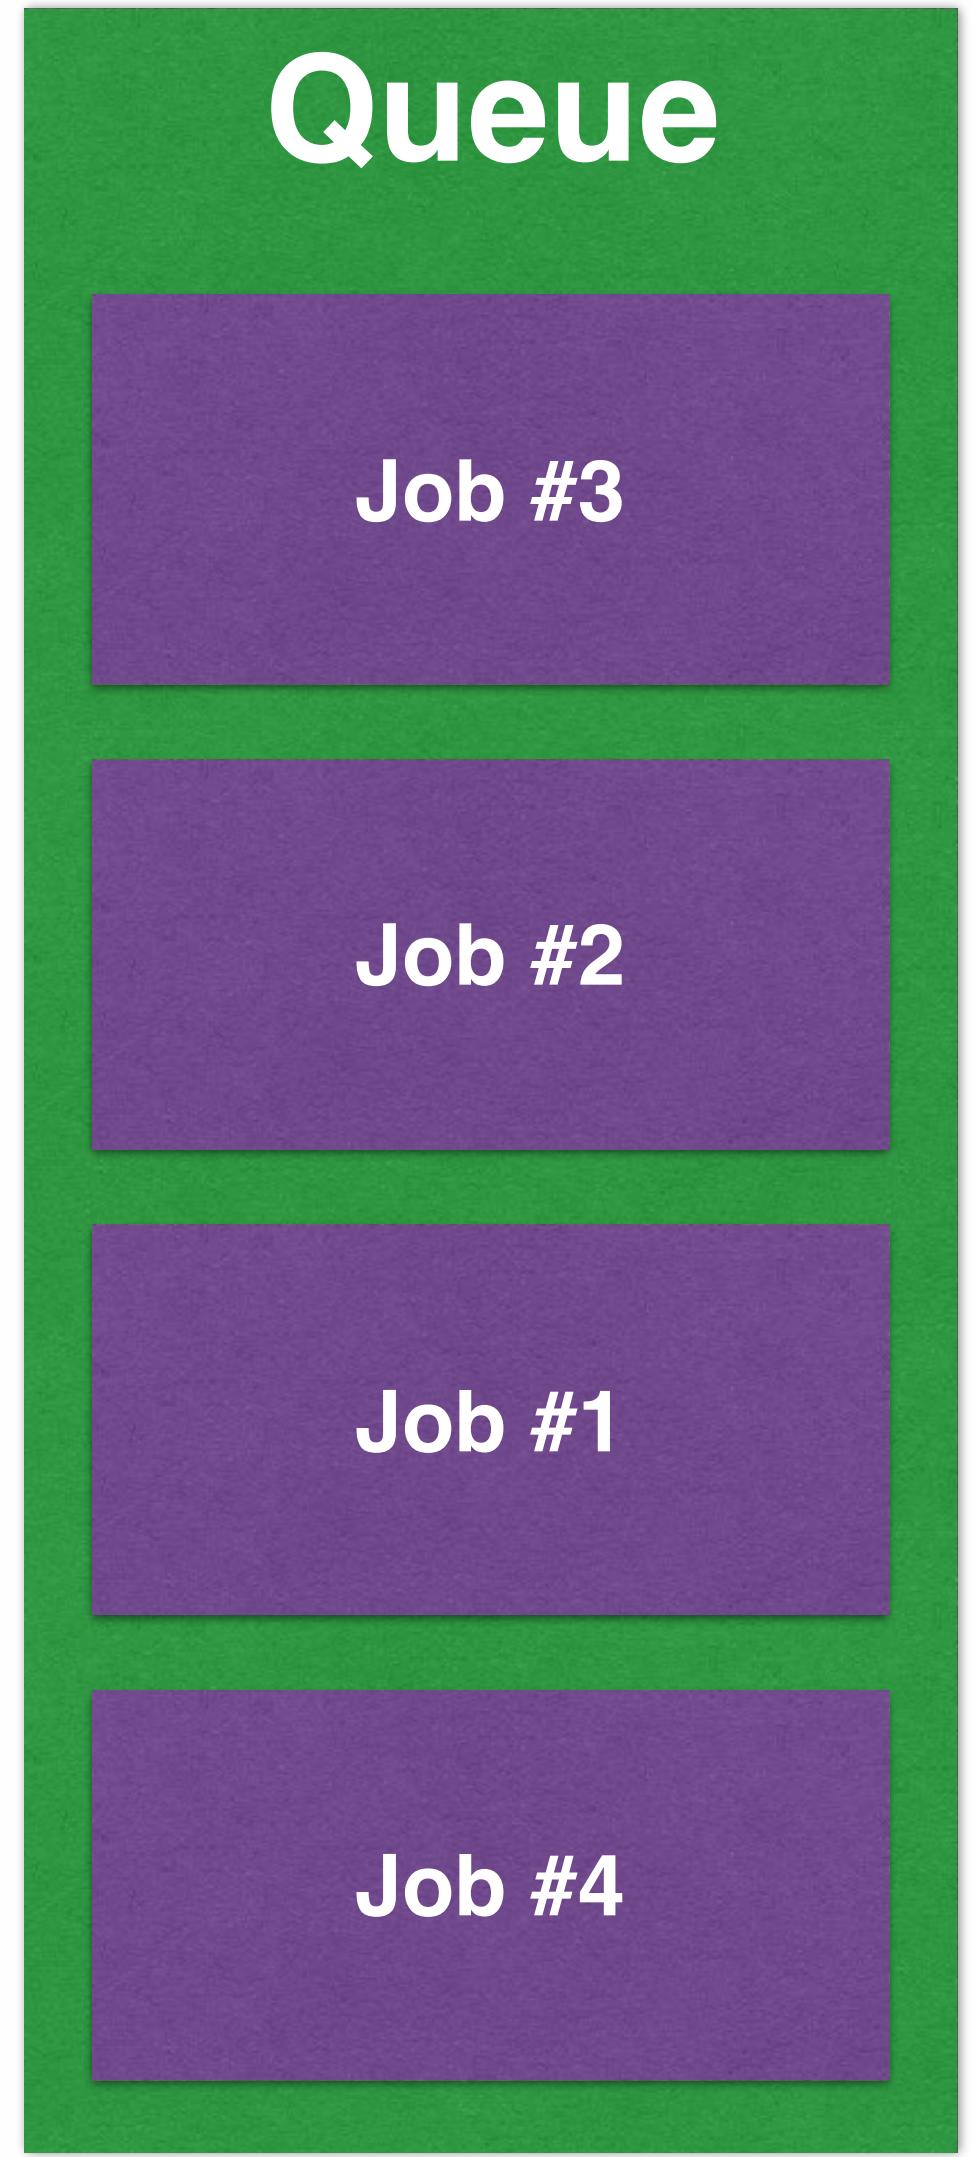
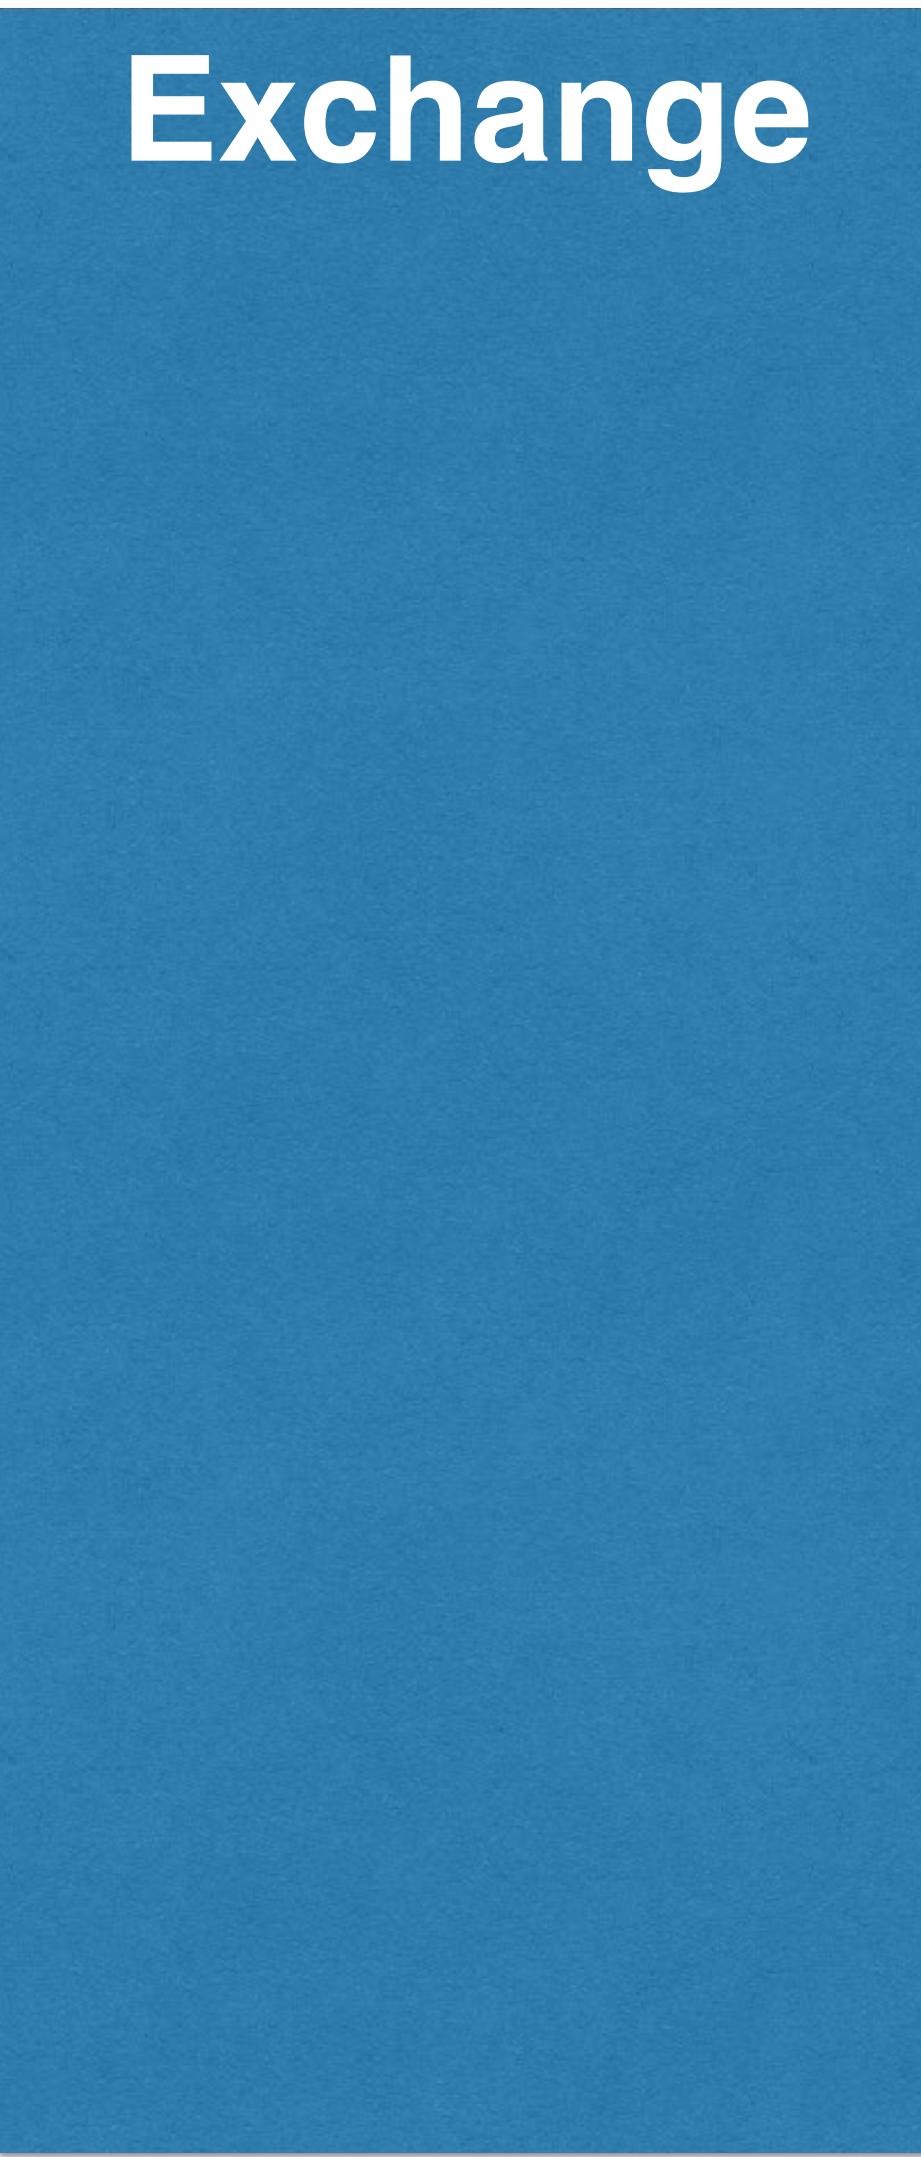
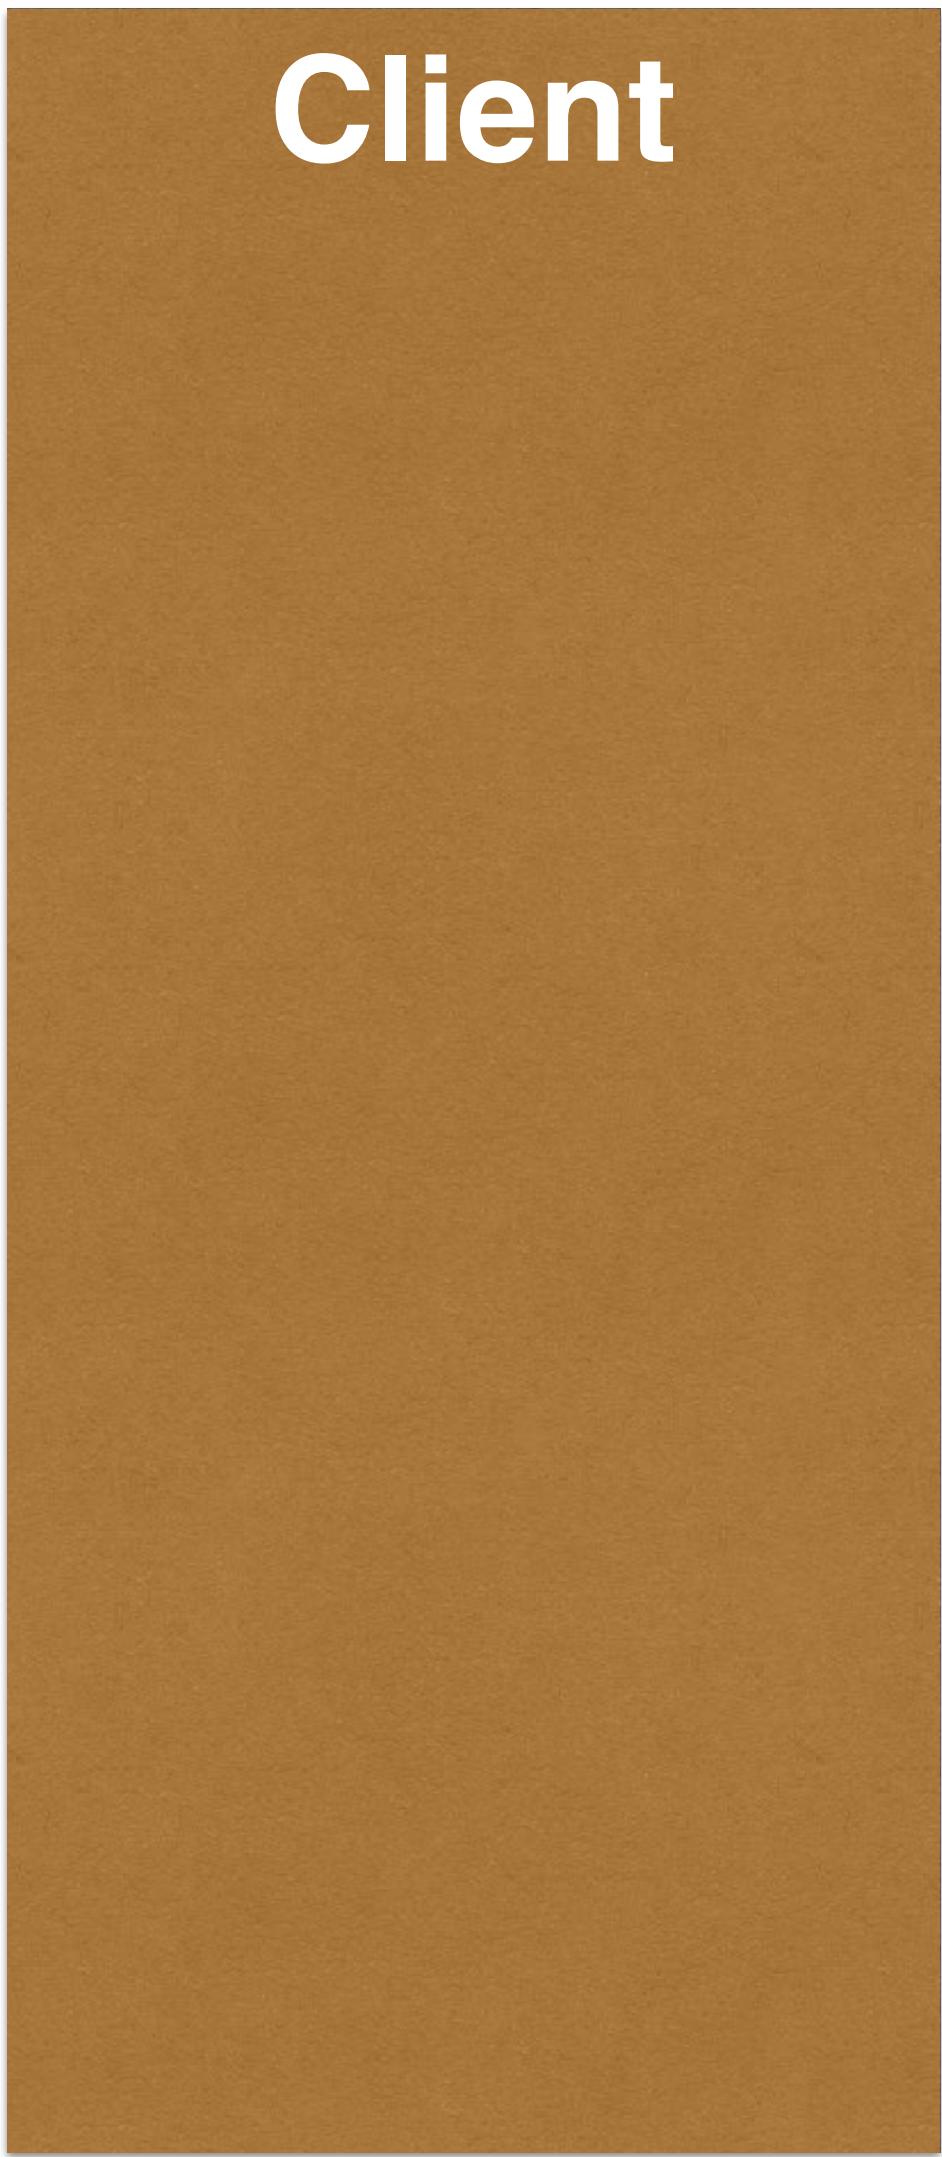
But what would you use the?

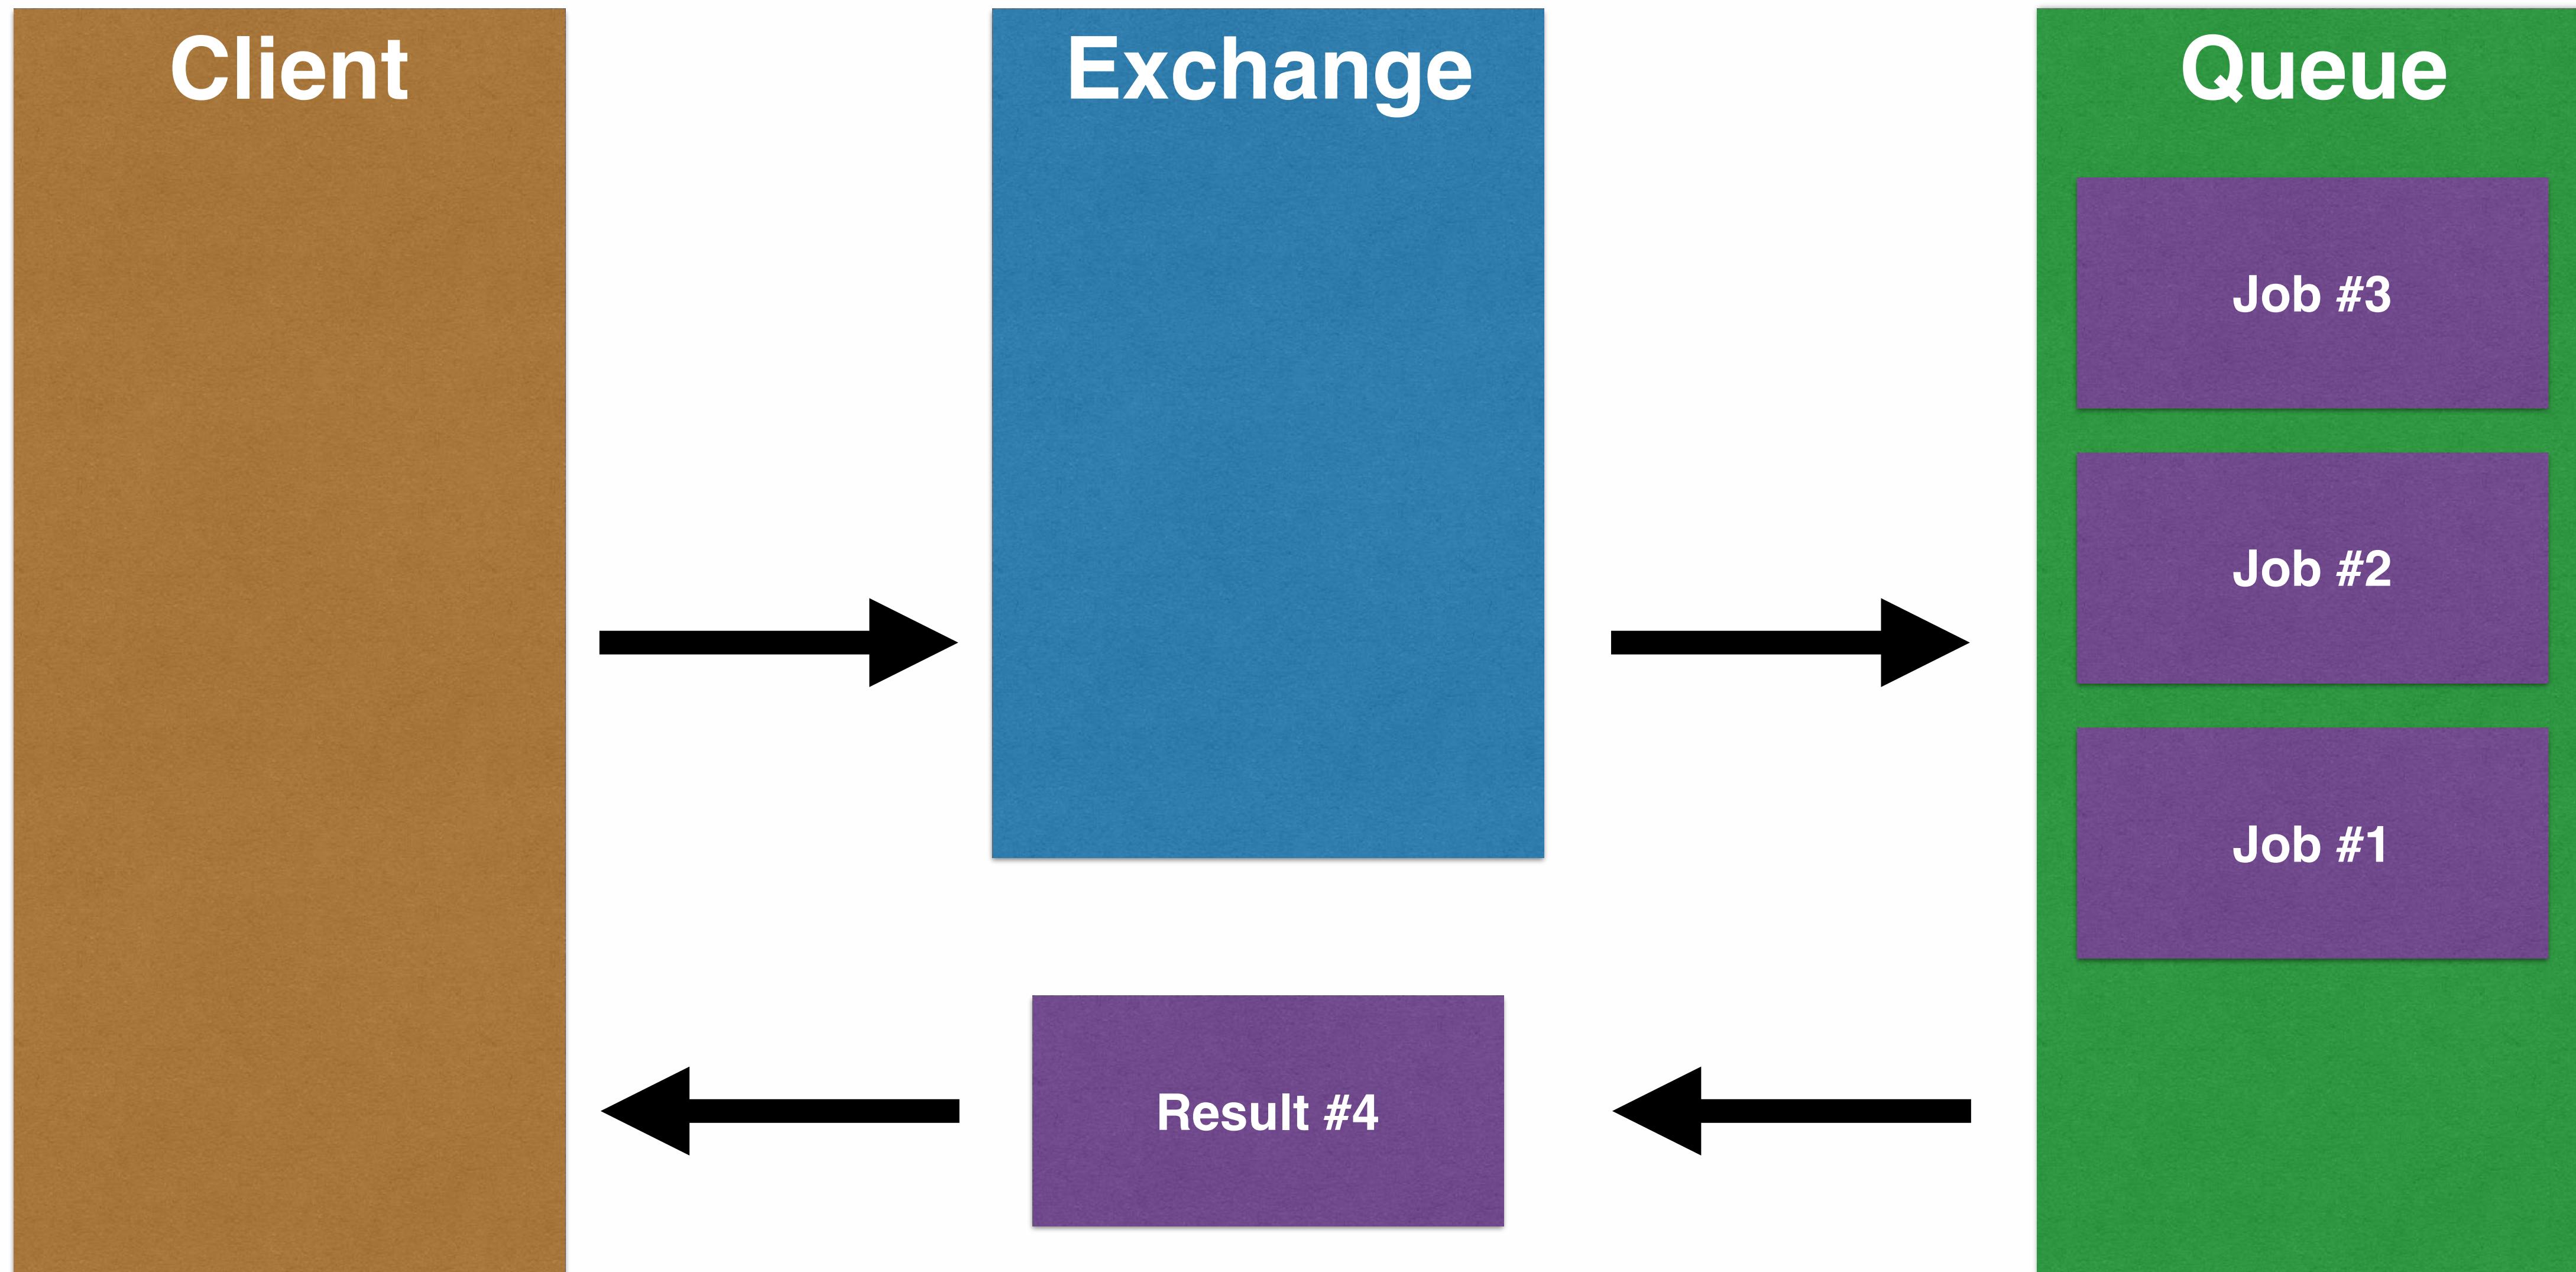
SPECIAL FEATURES

Direct reply-to









Client

Result #4

Exchange

Queue

Job #3

Job #2

Job #1

Smart Lock

Check user PIN

Heartbeat

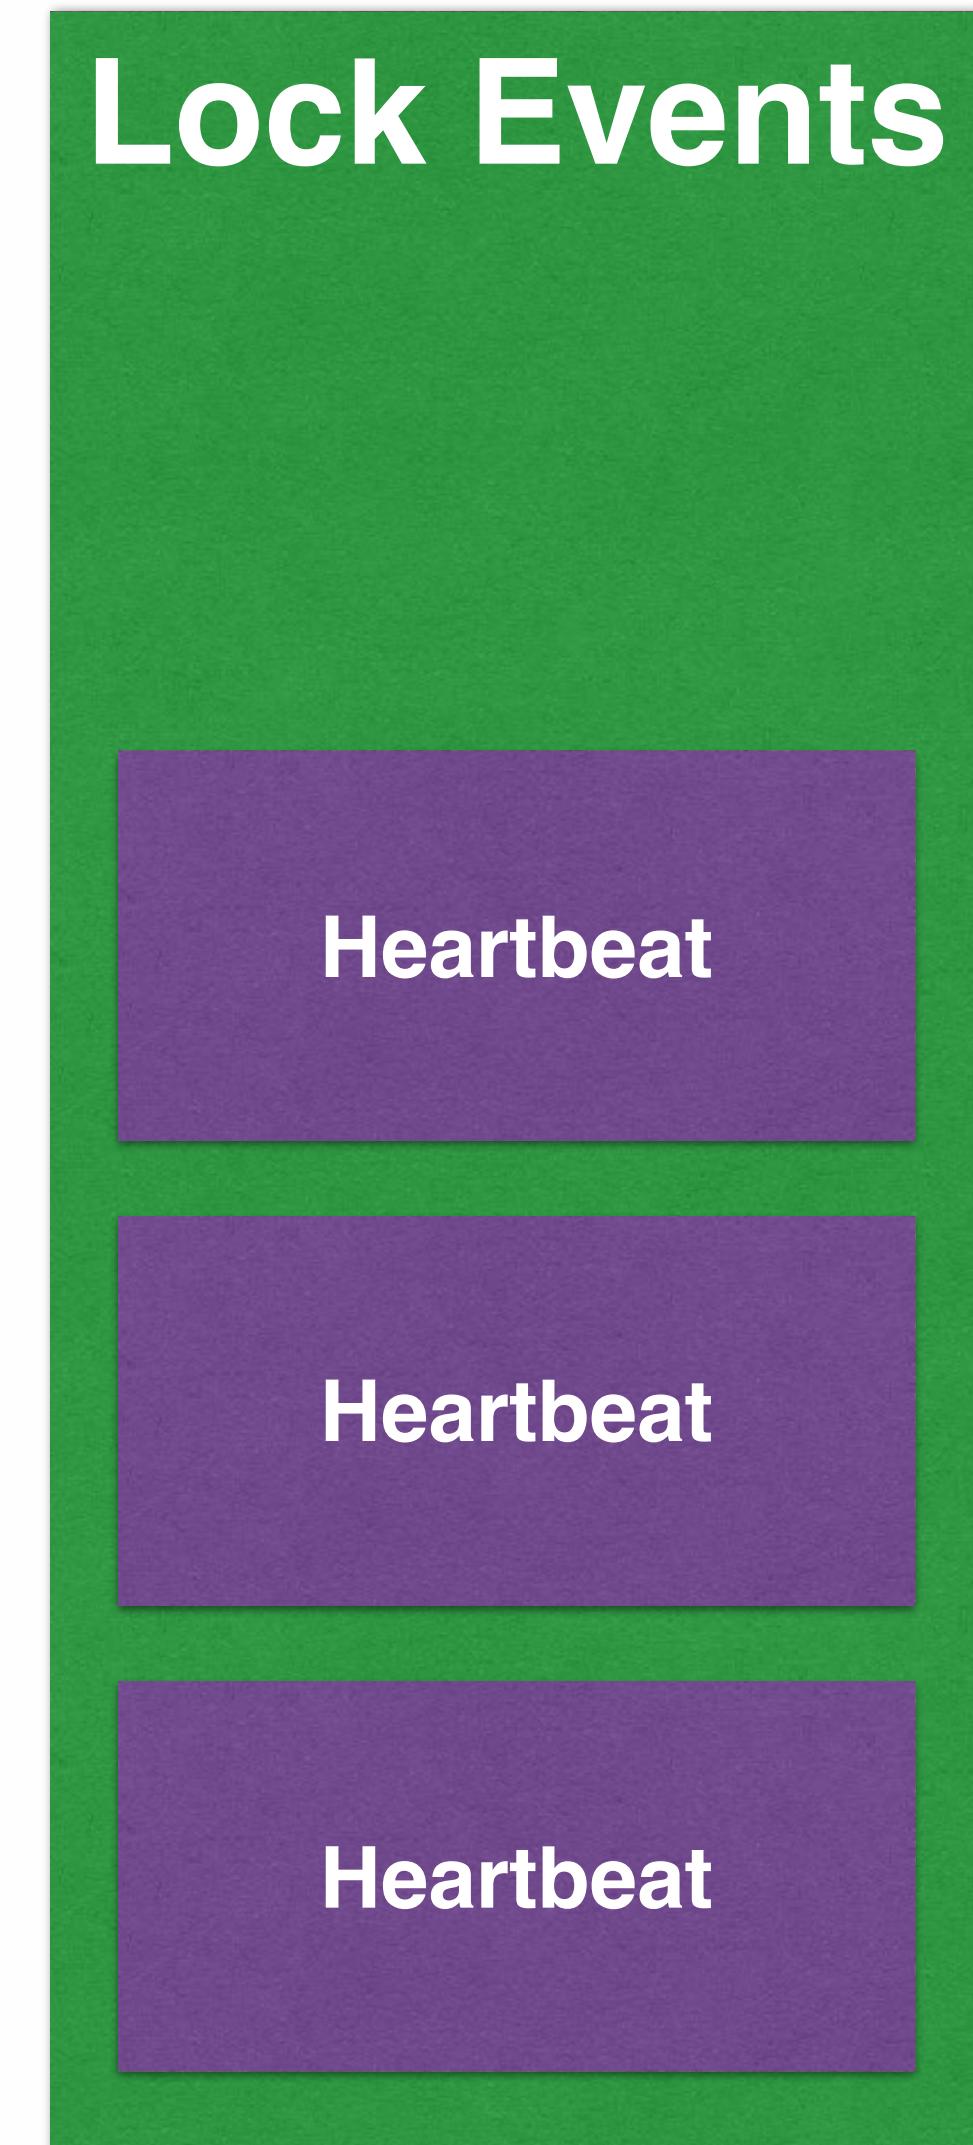
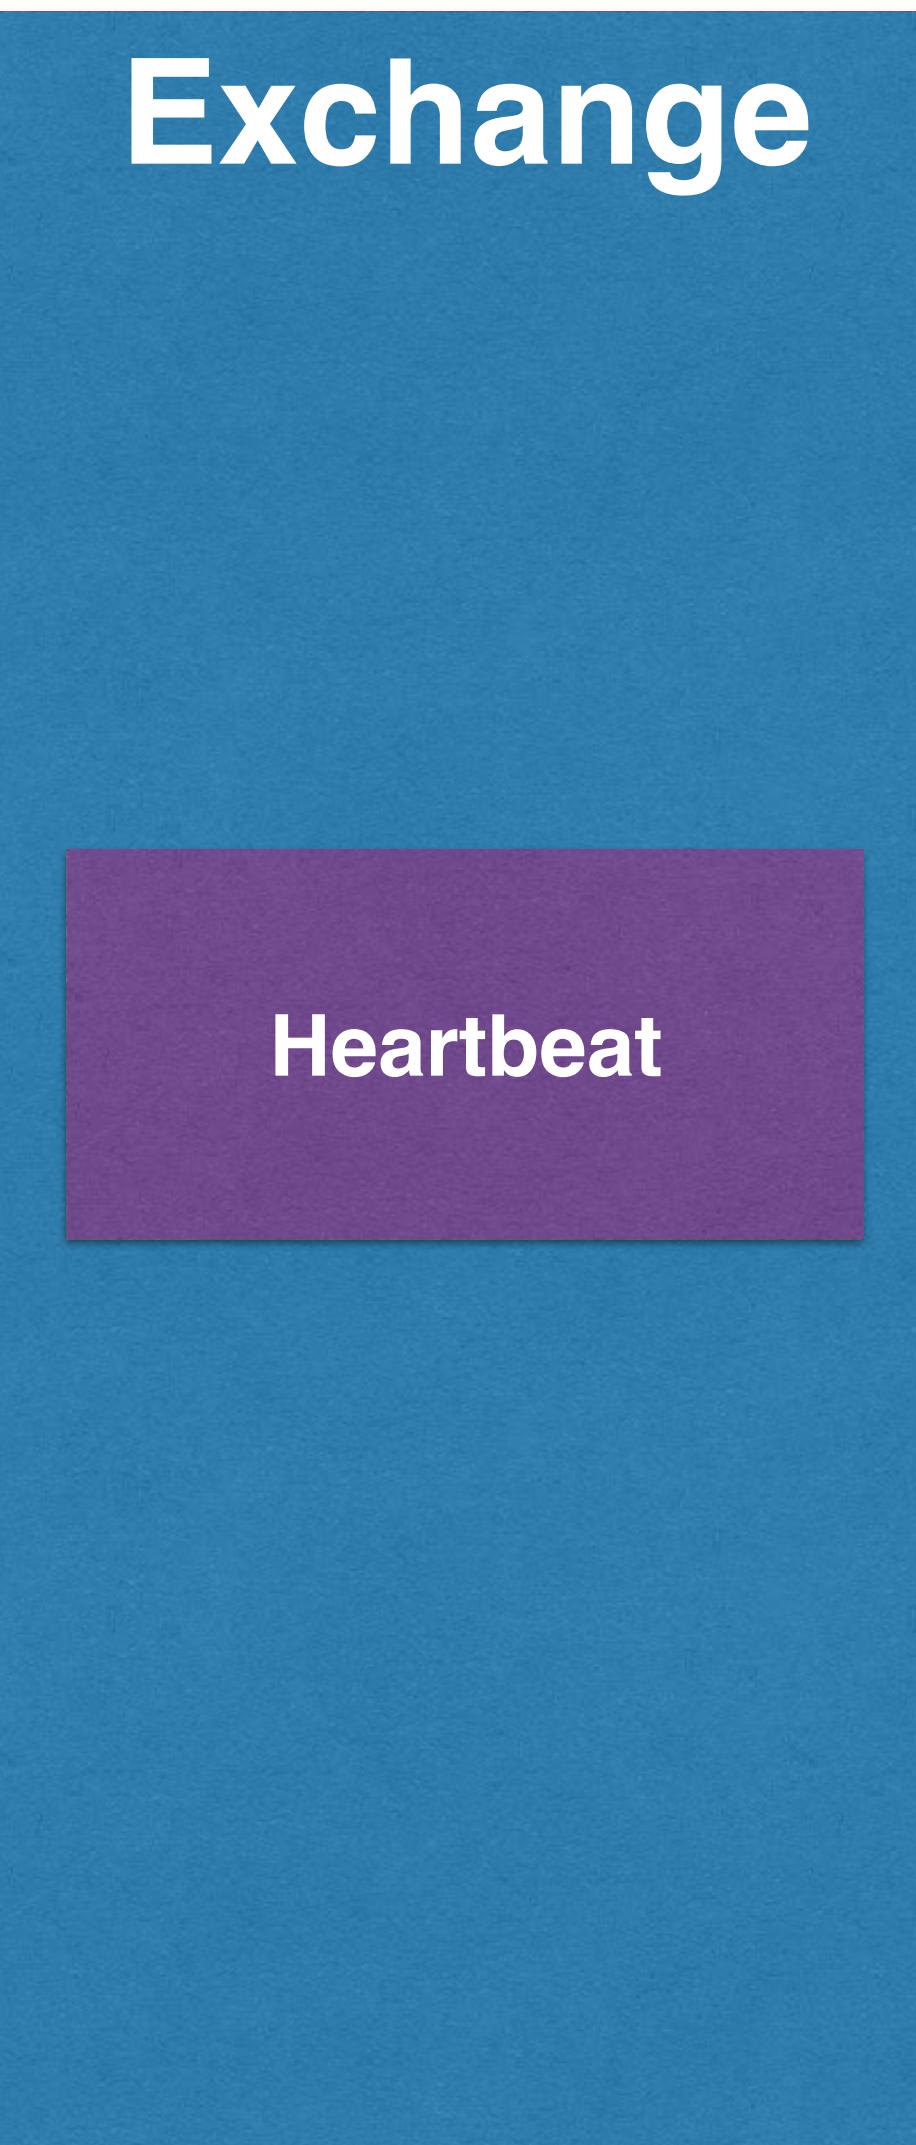
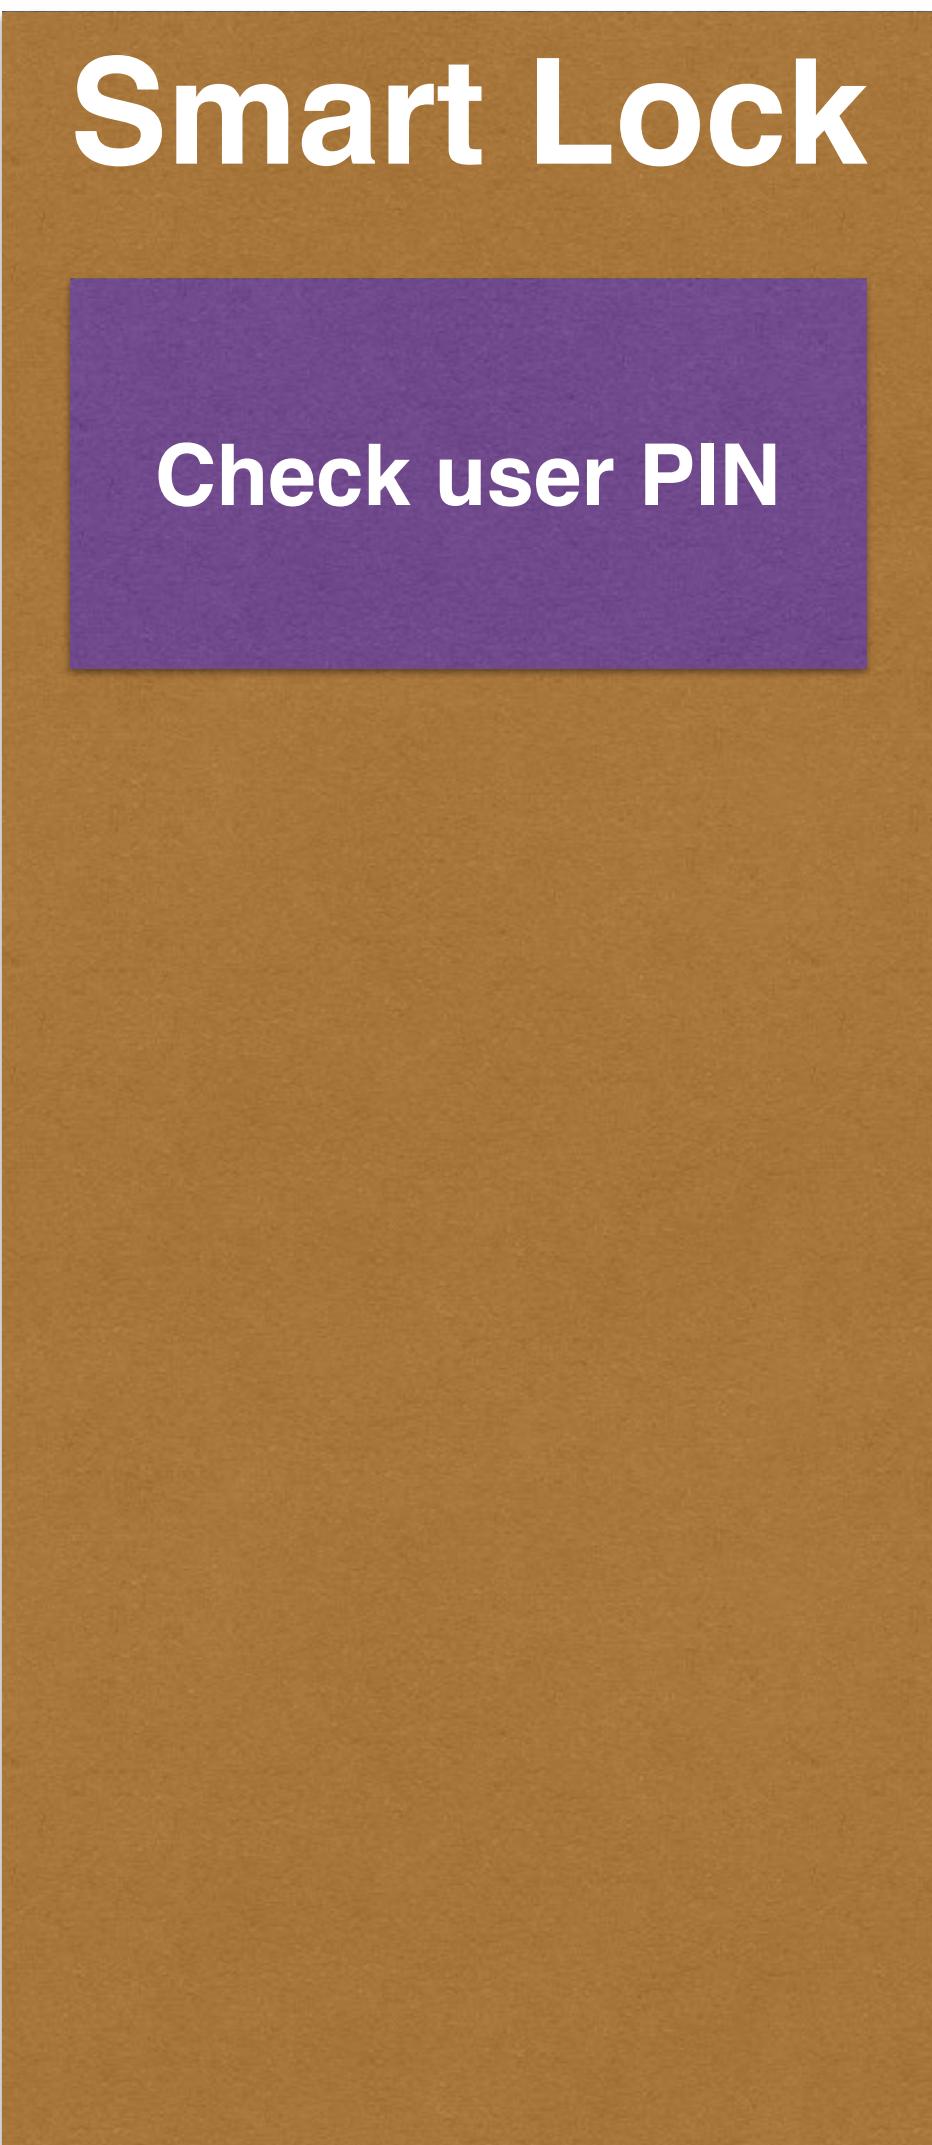
Exchange

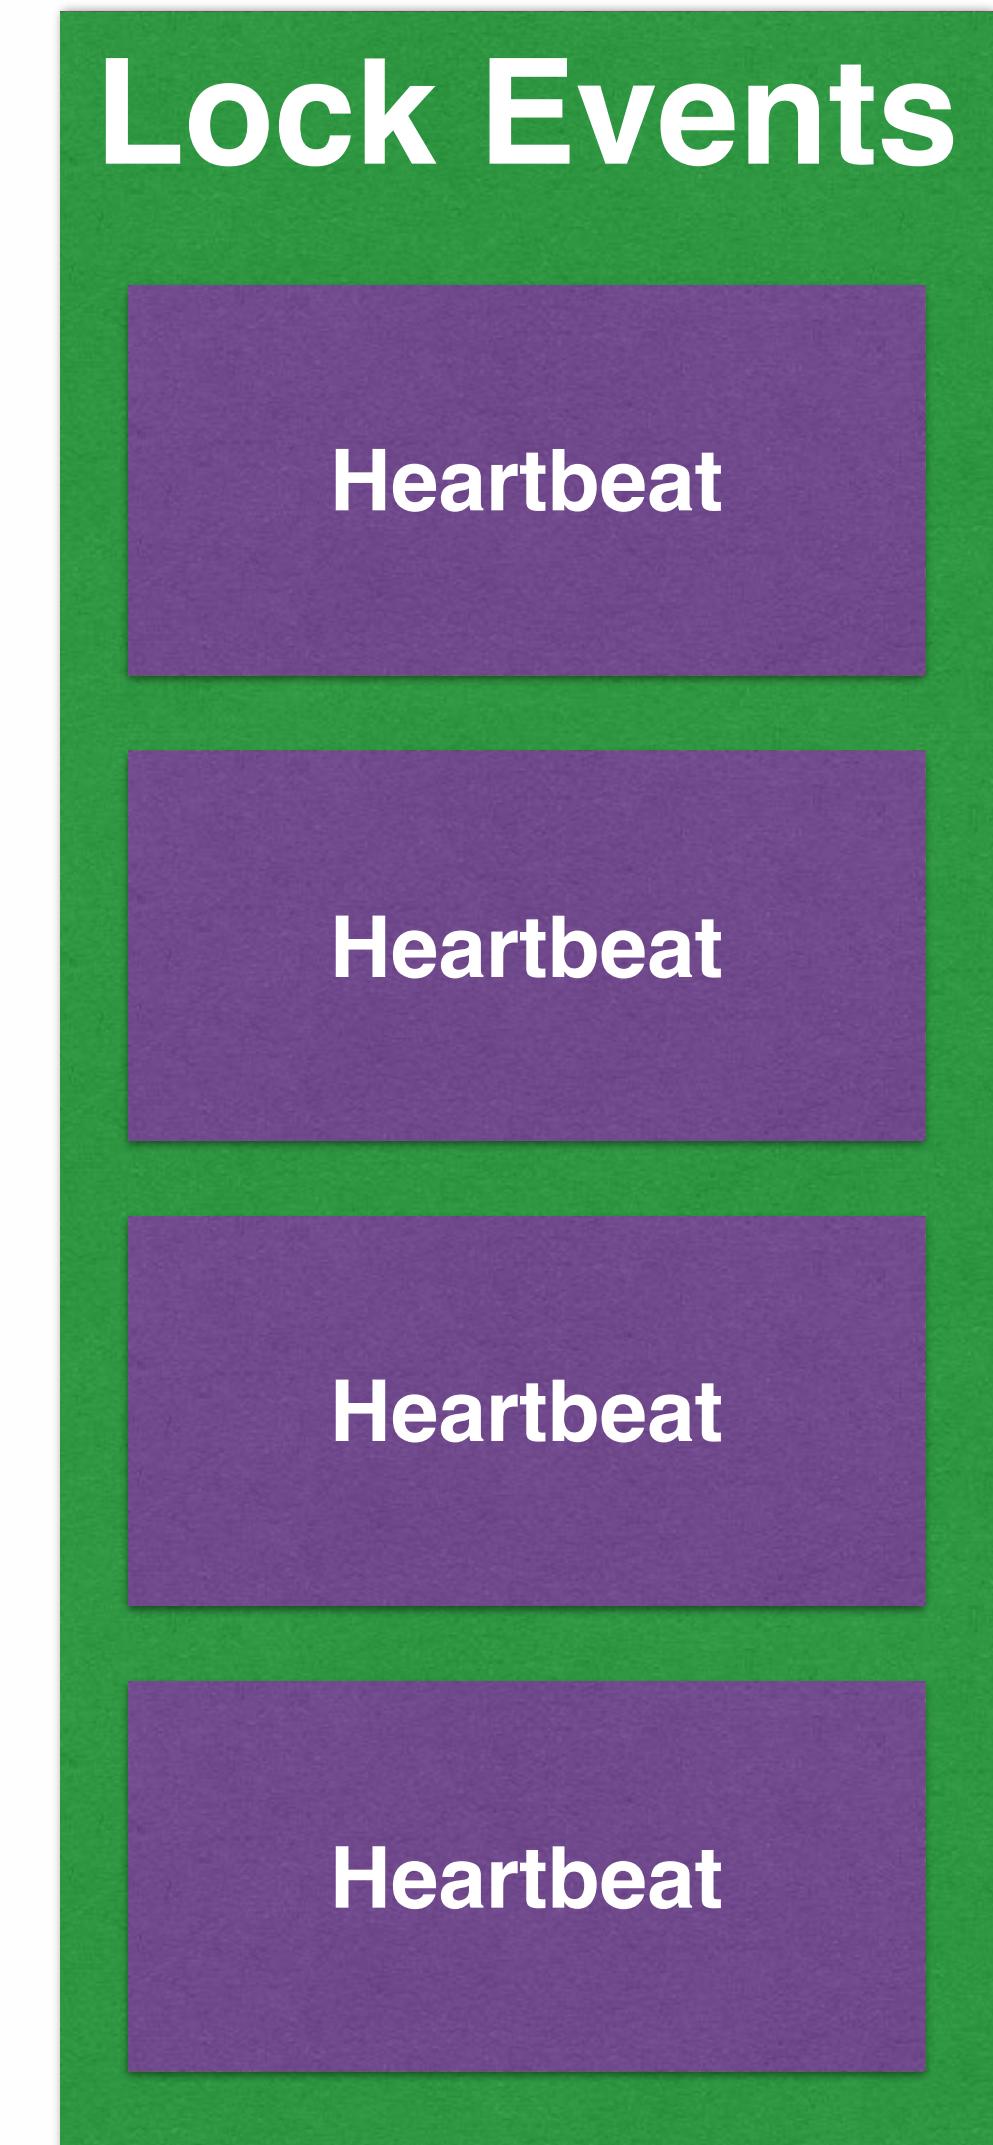
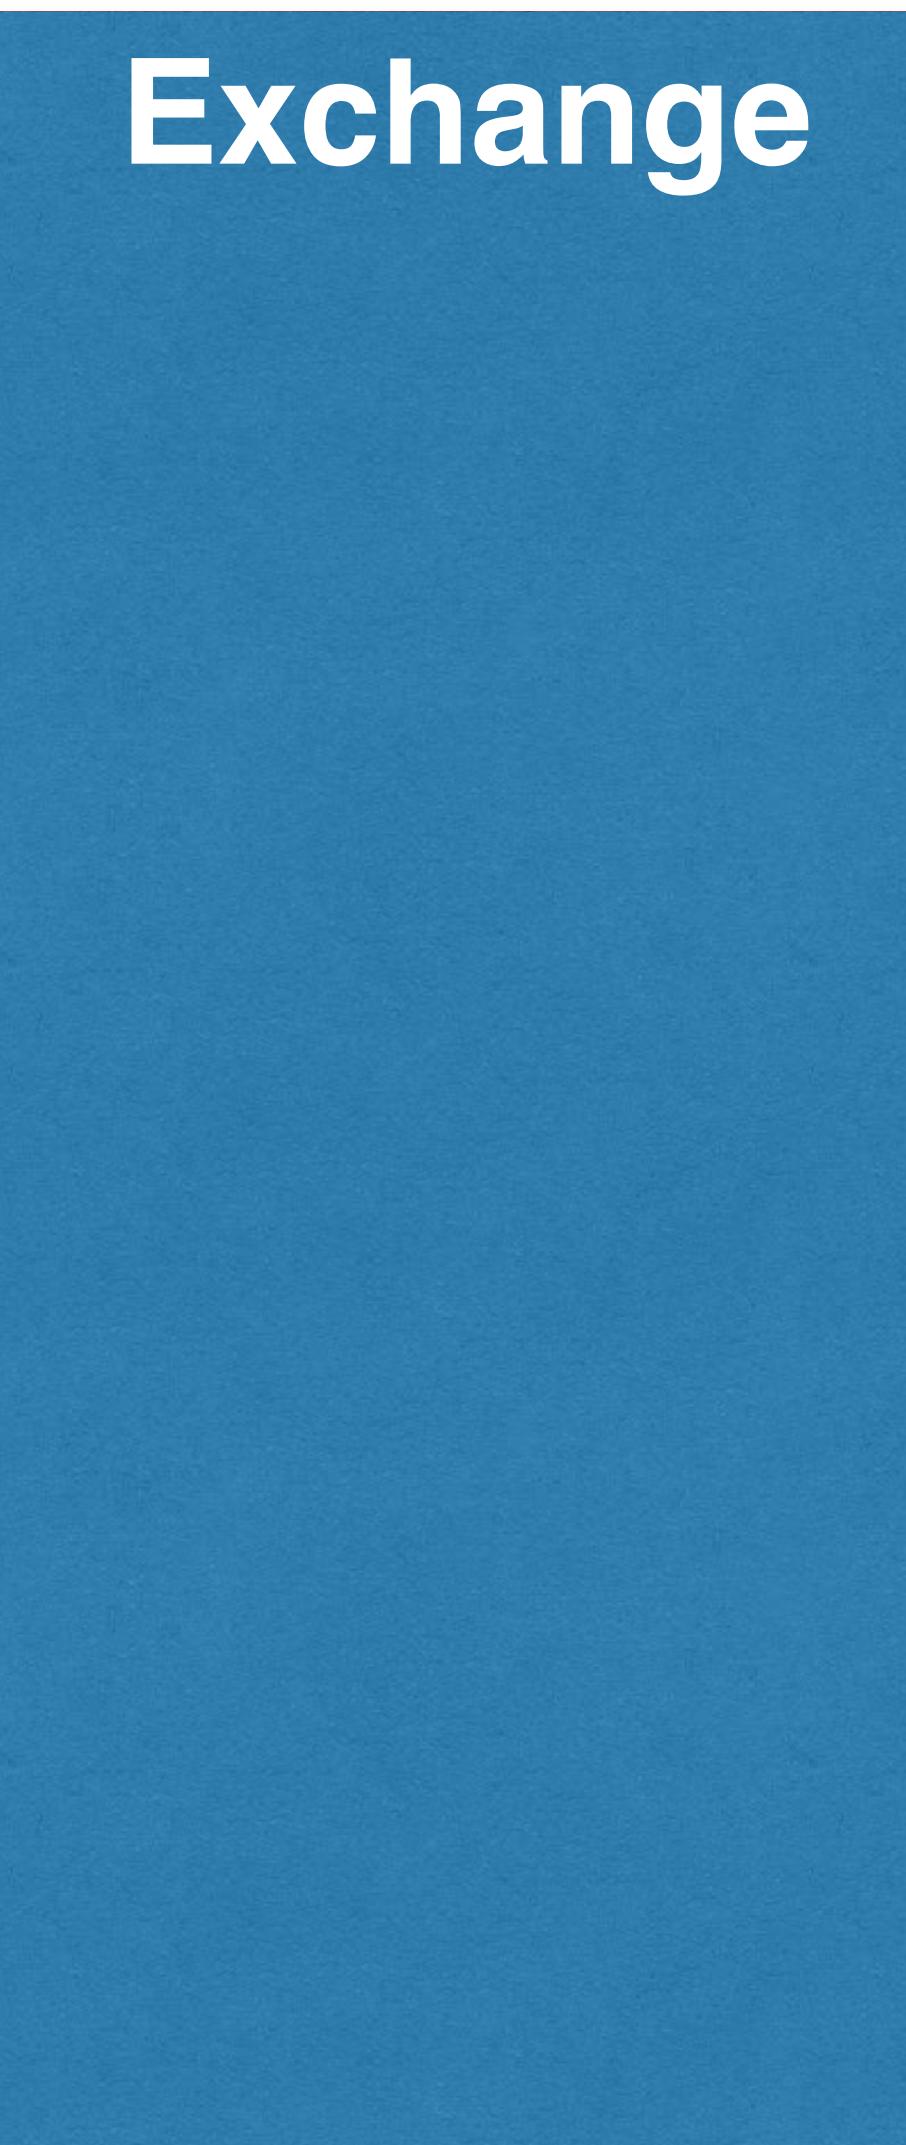
Lock Events

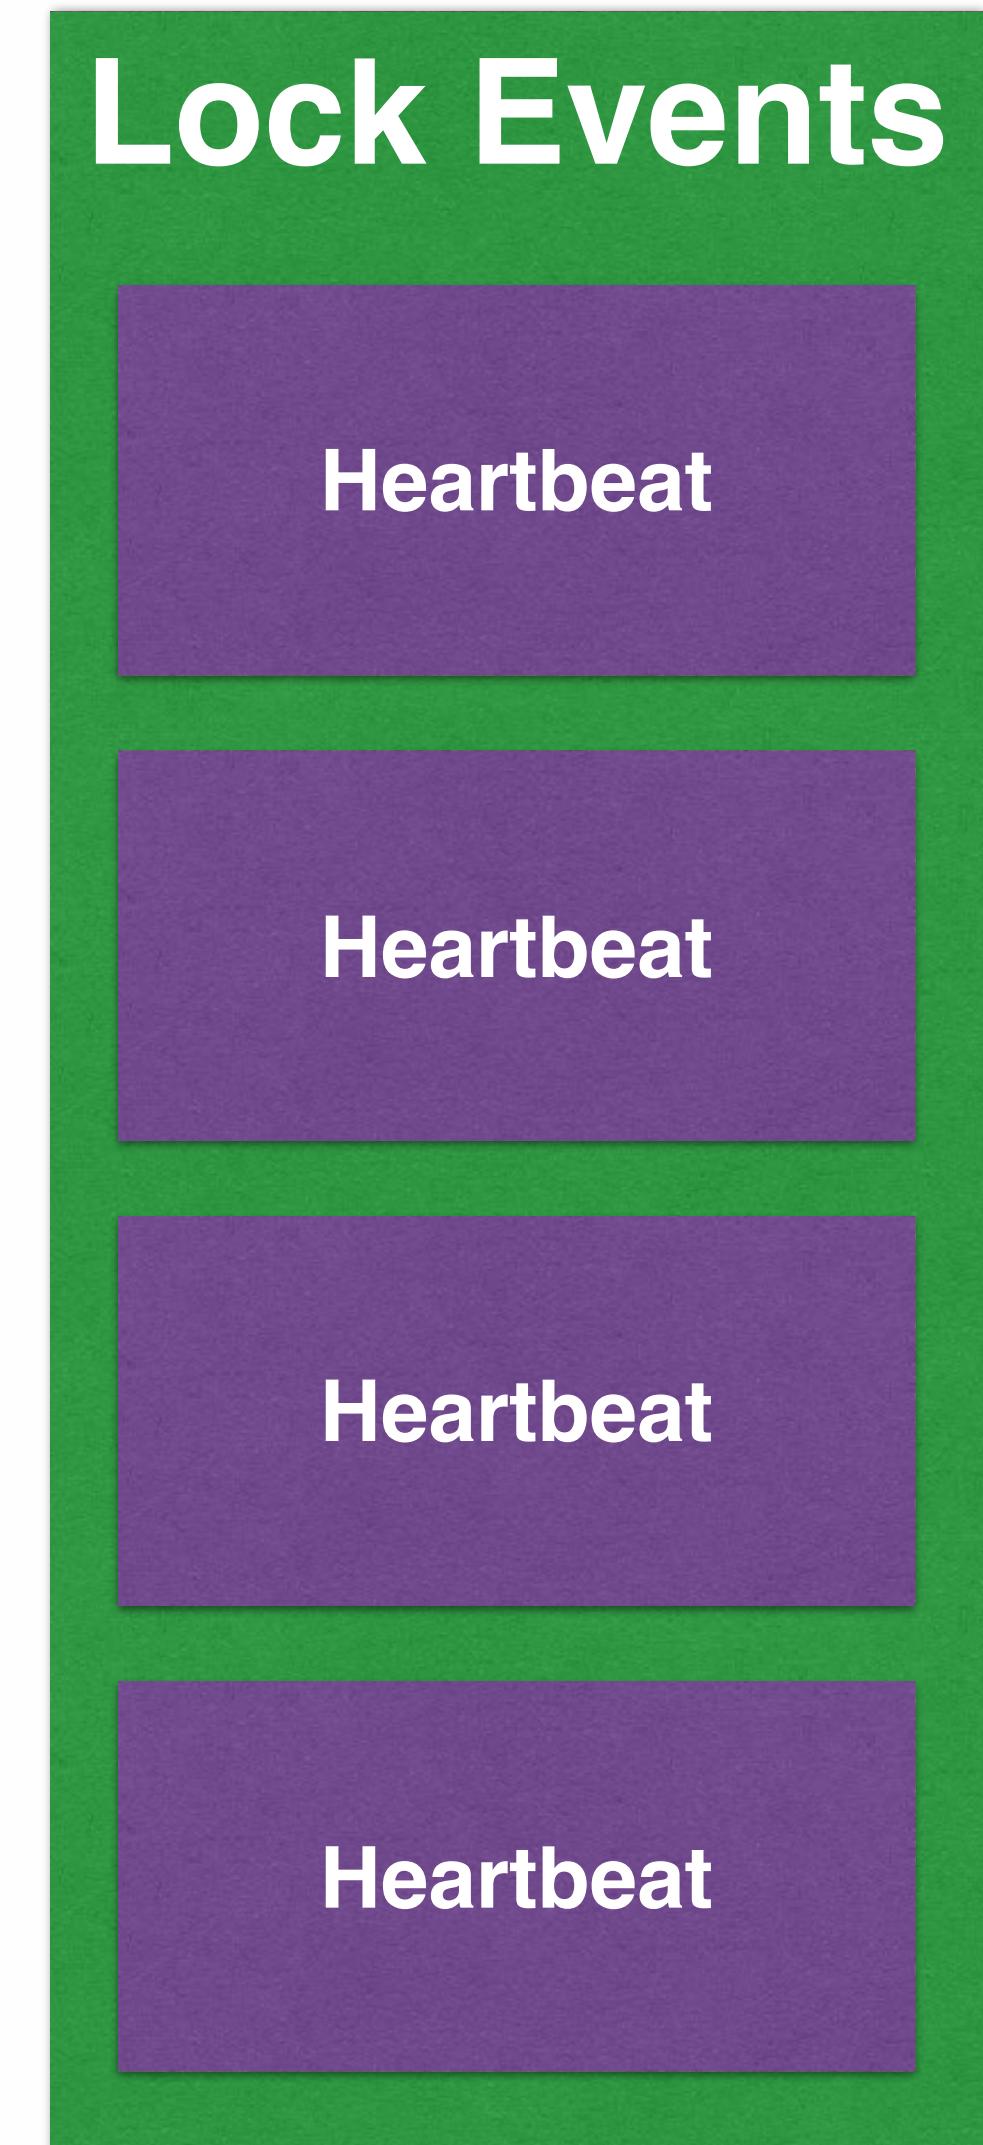
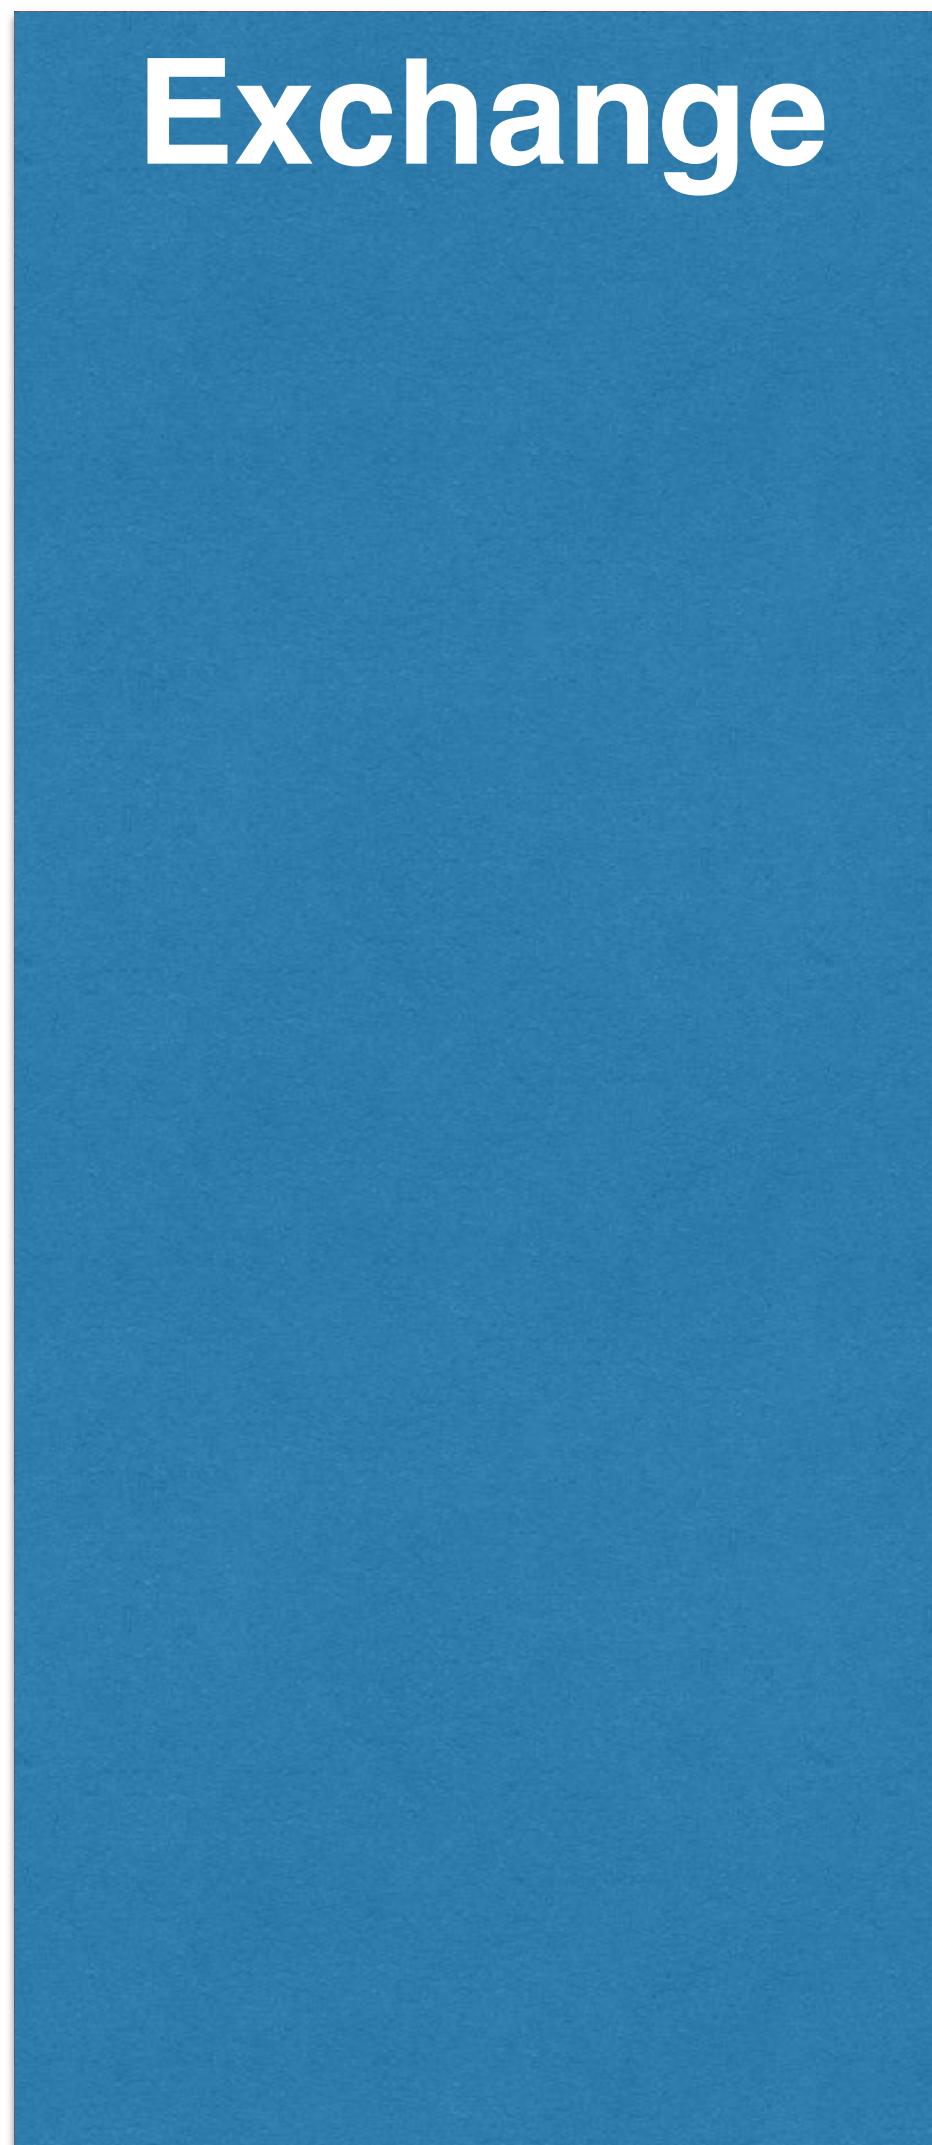
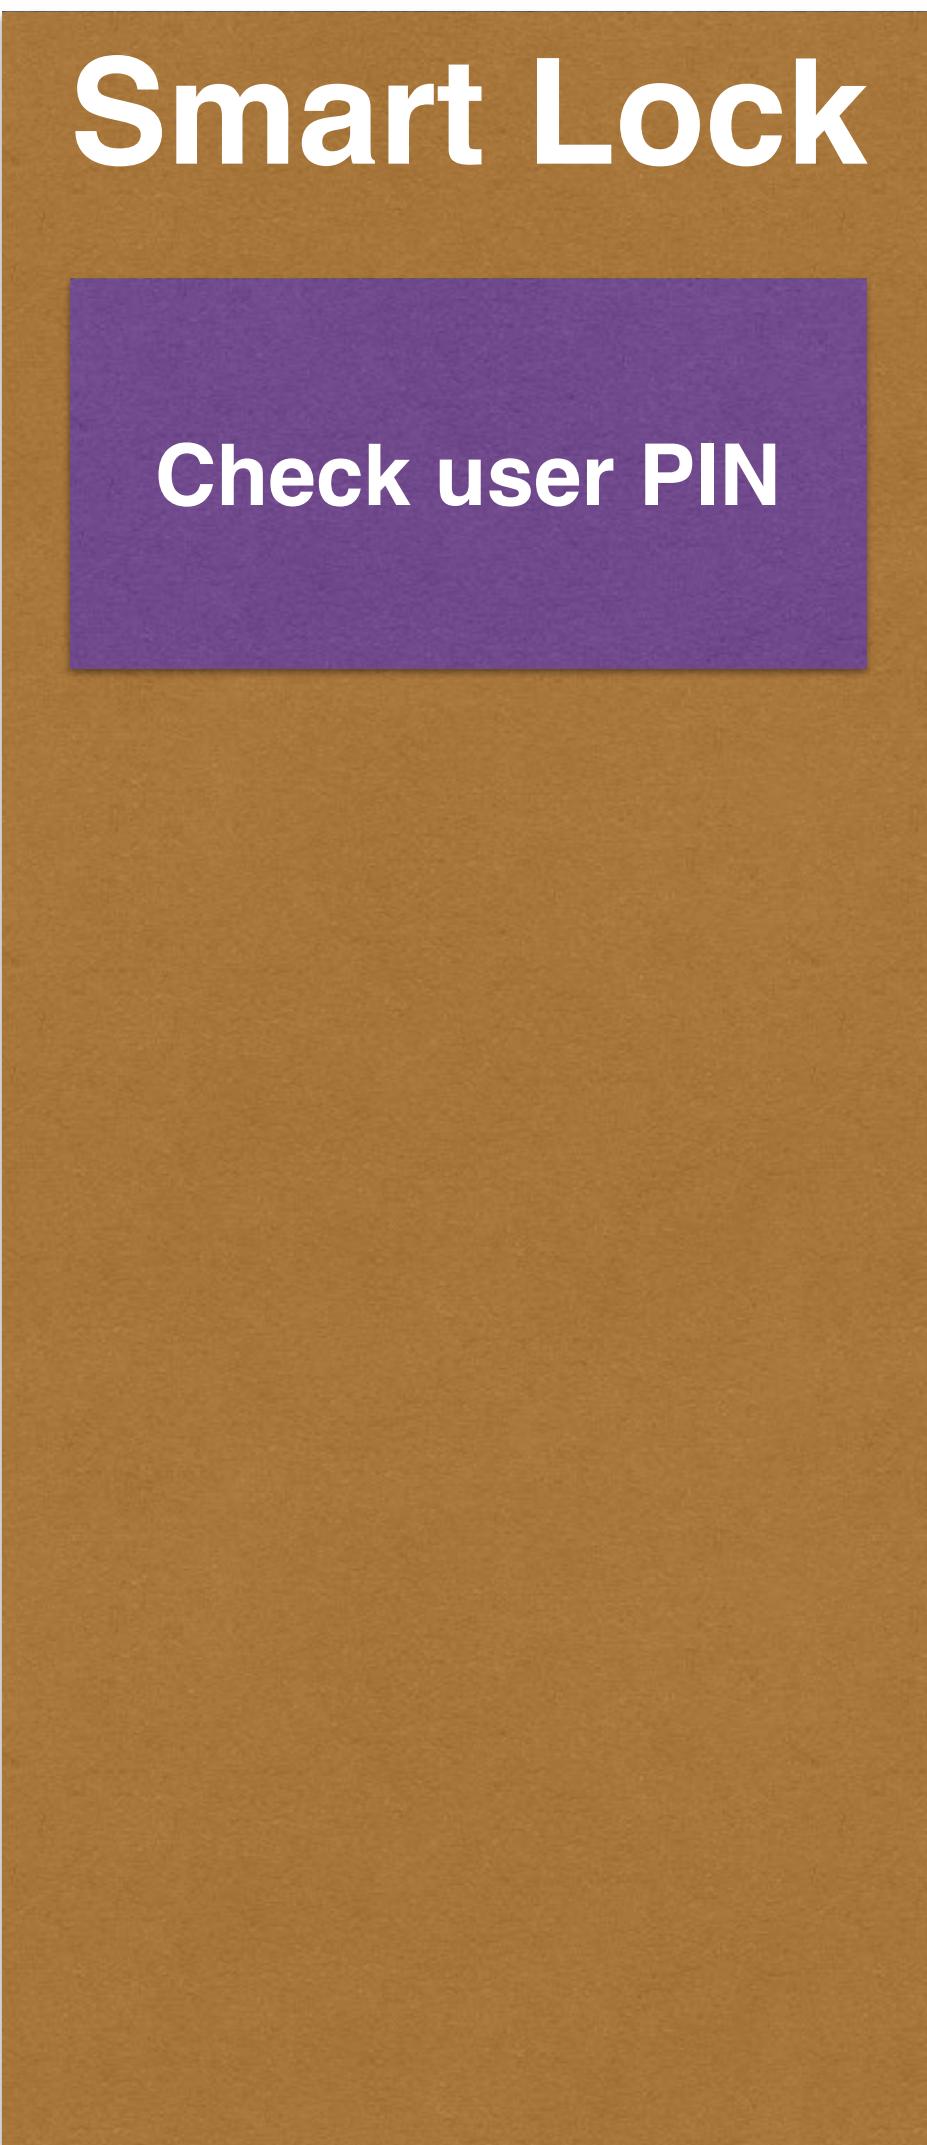
Heartbeat

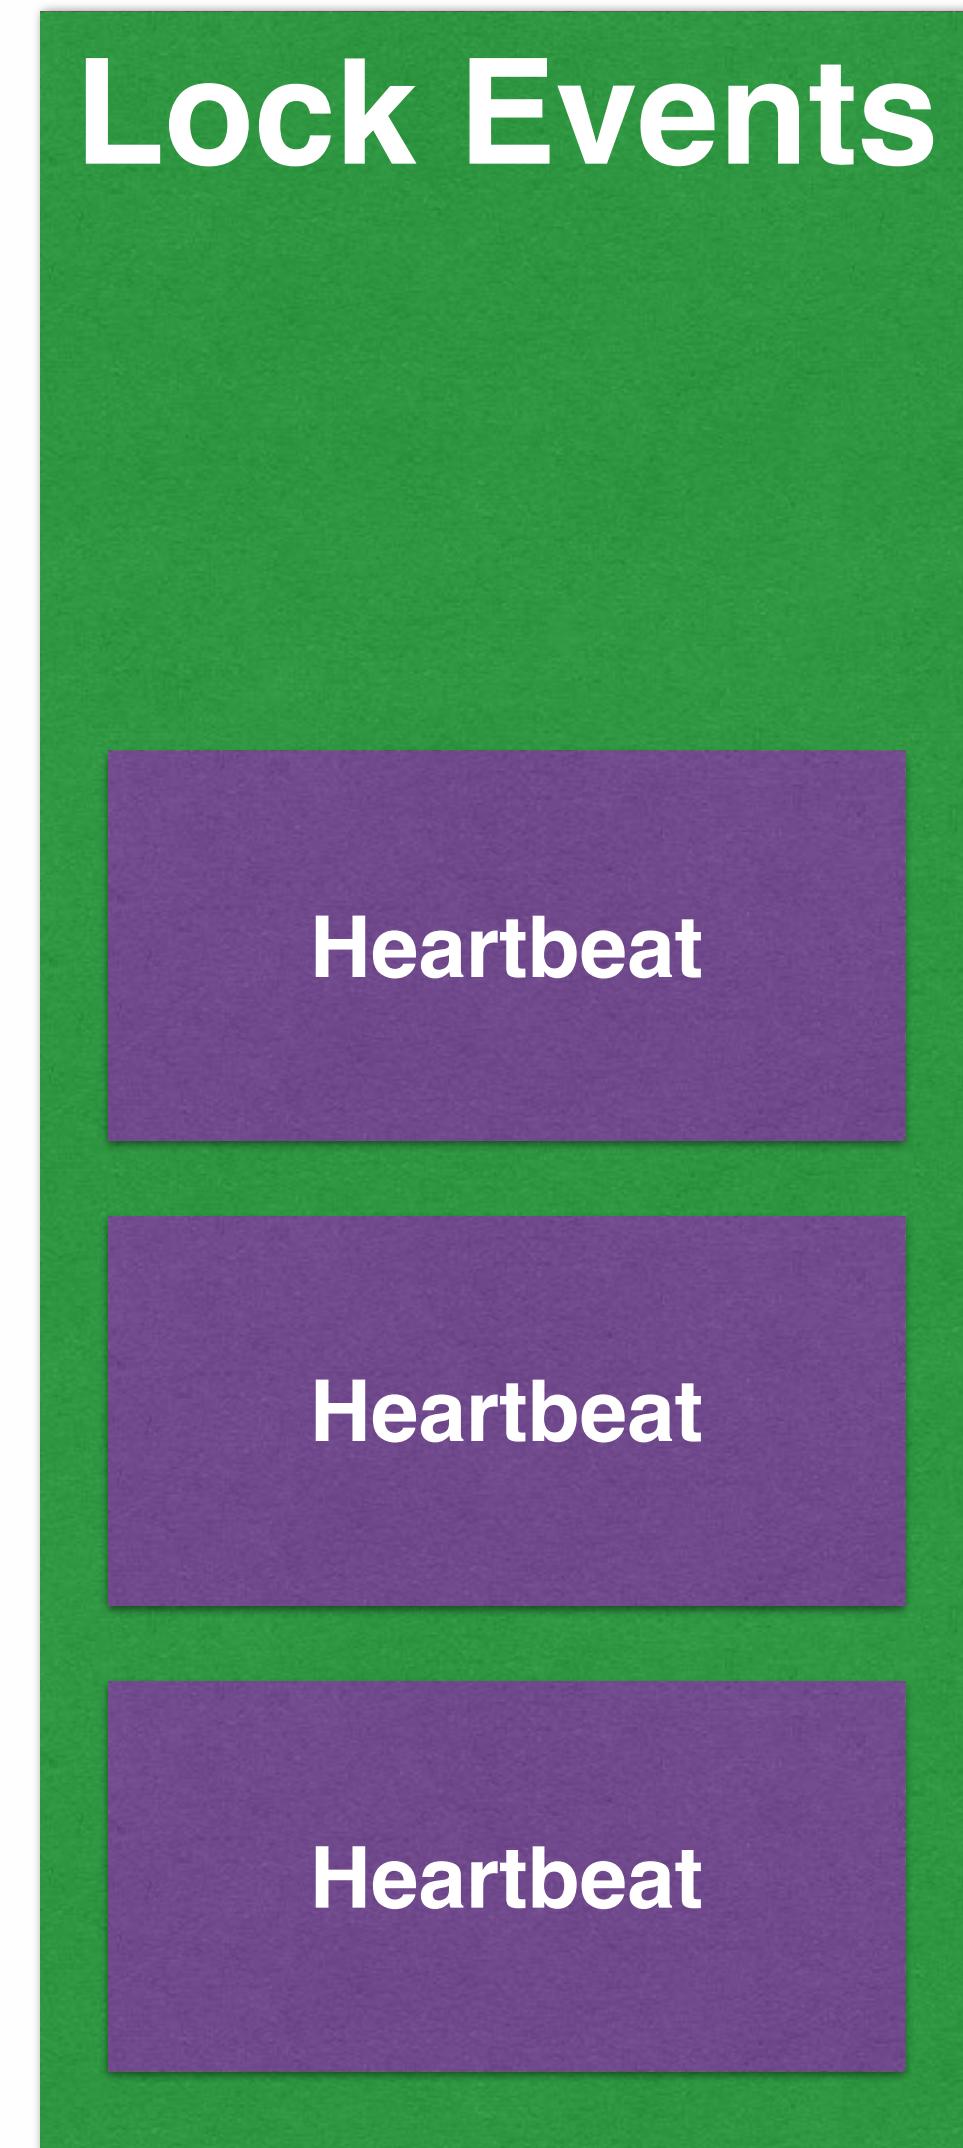
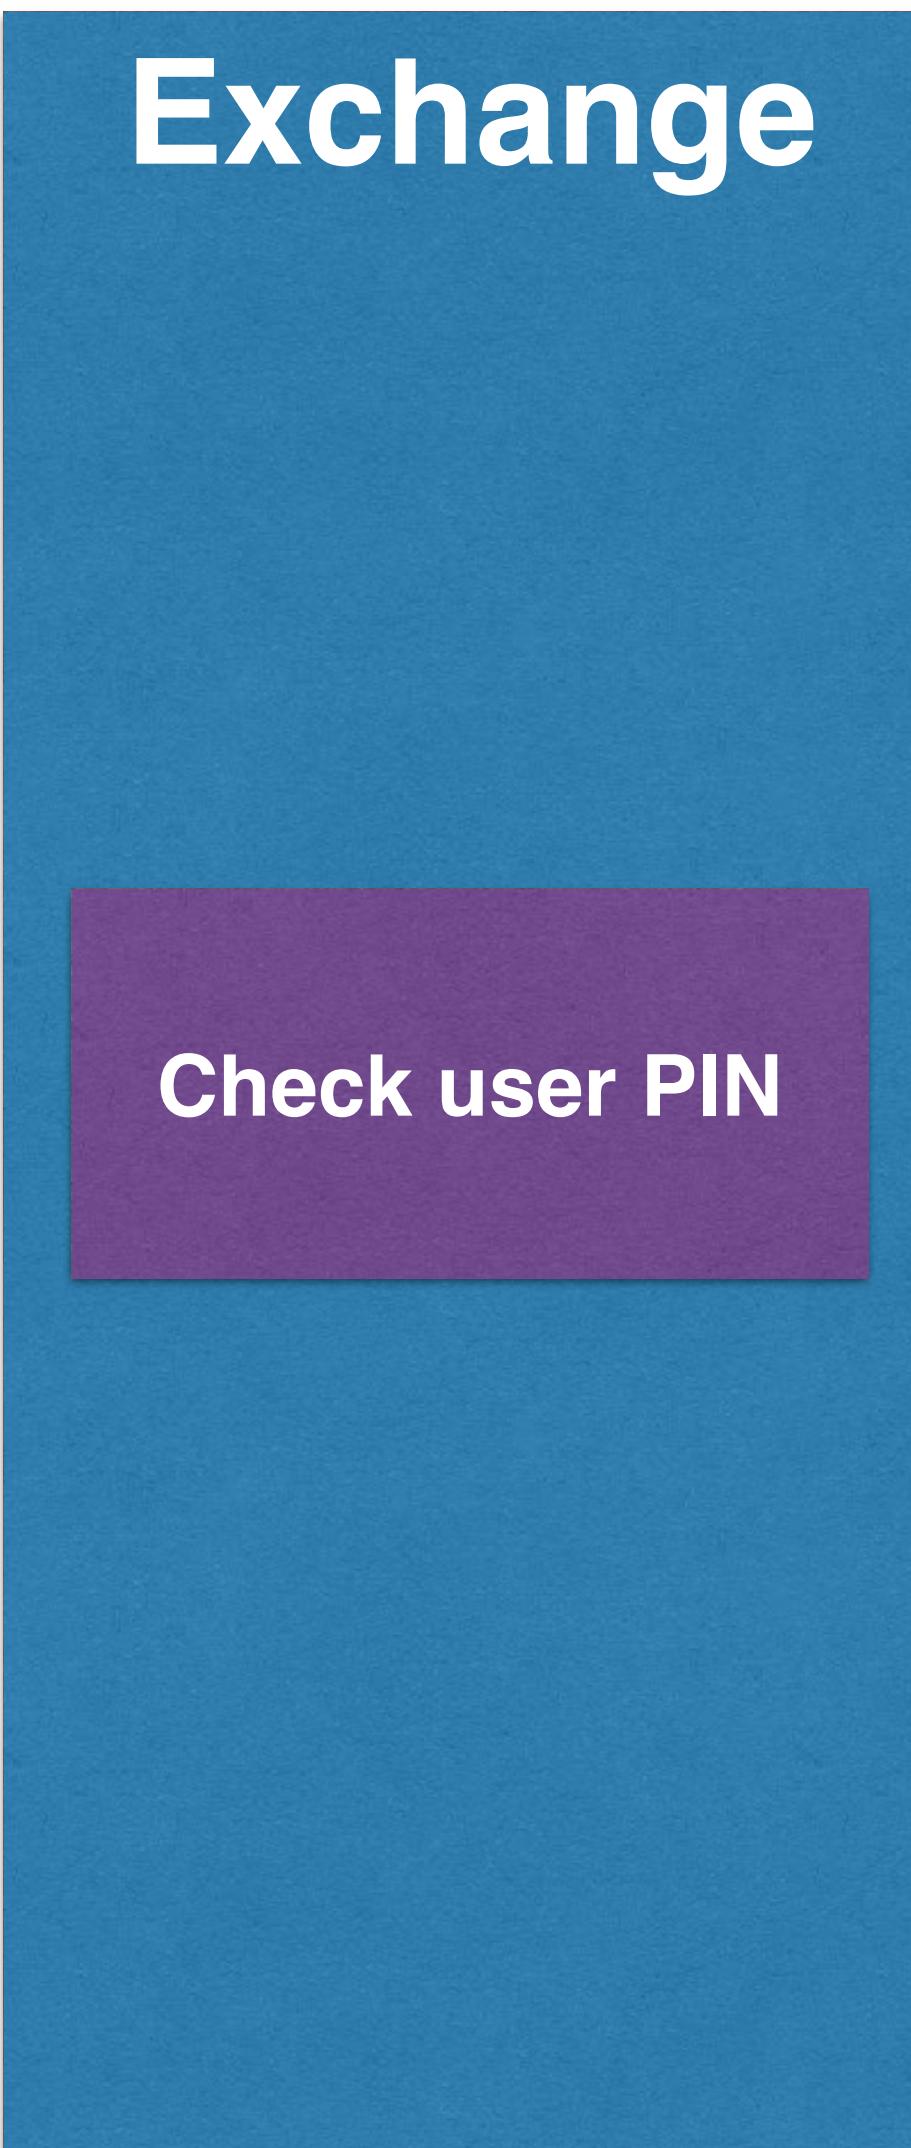
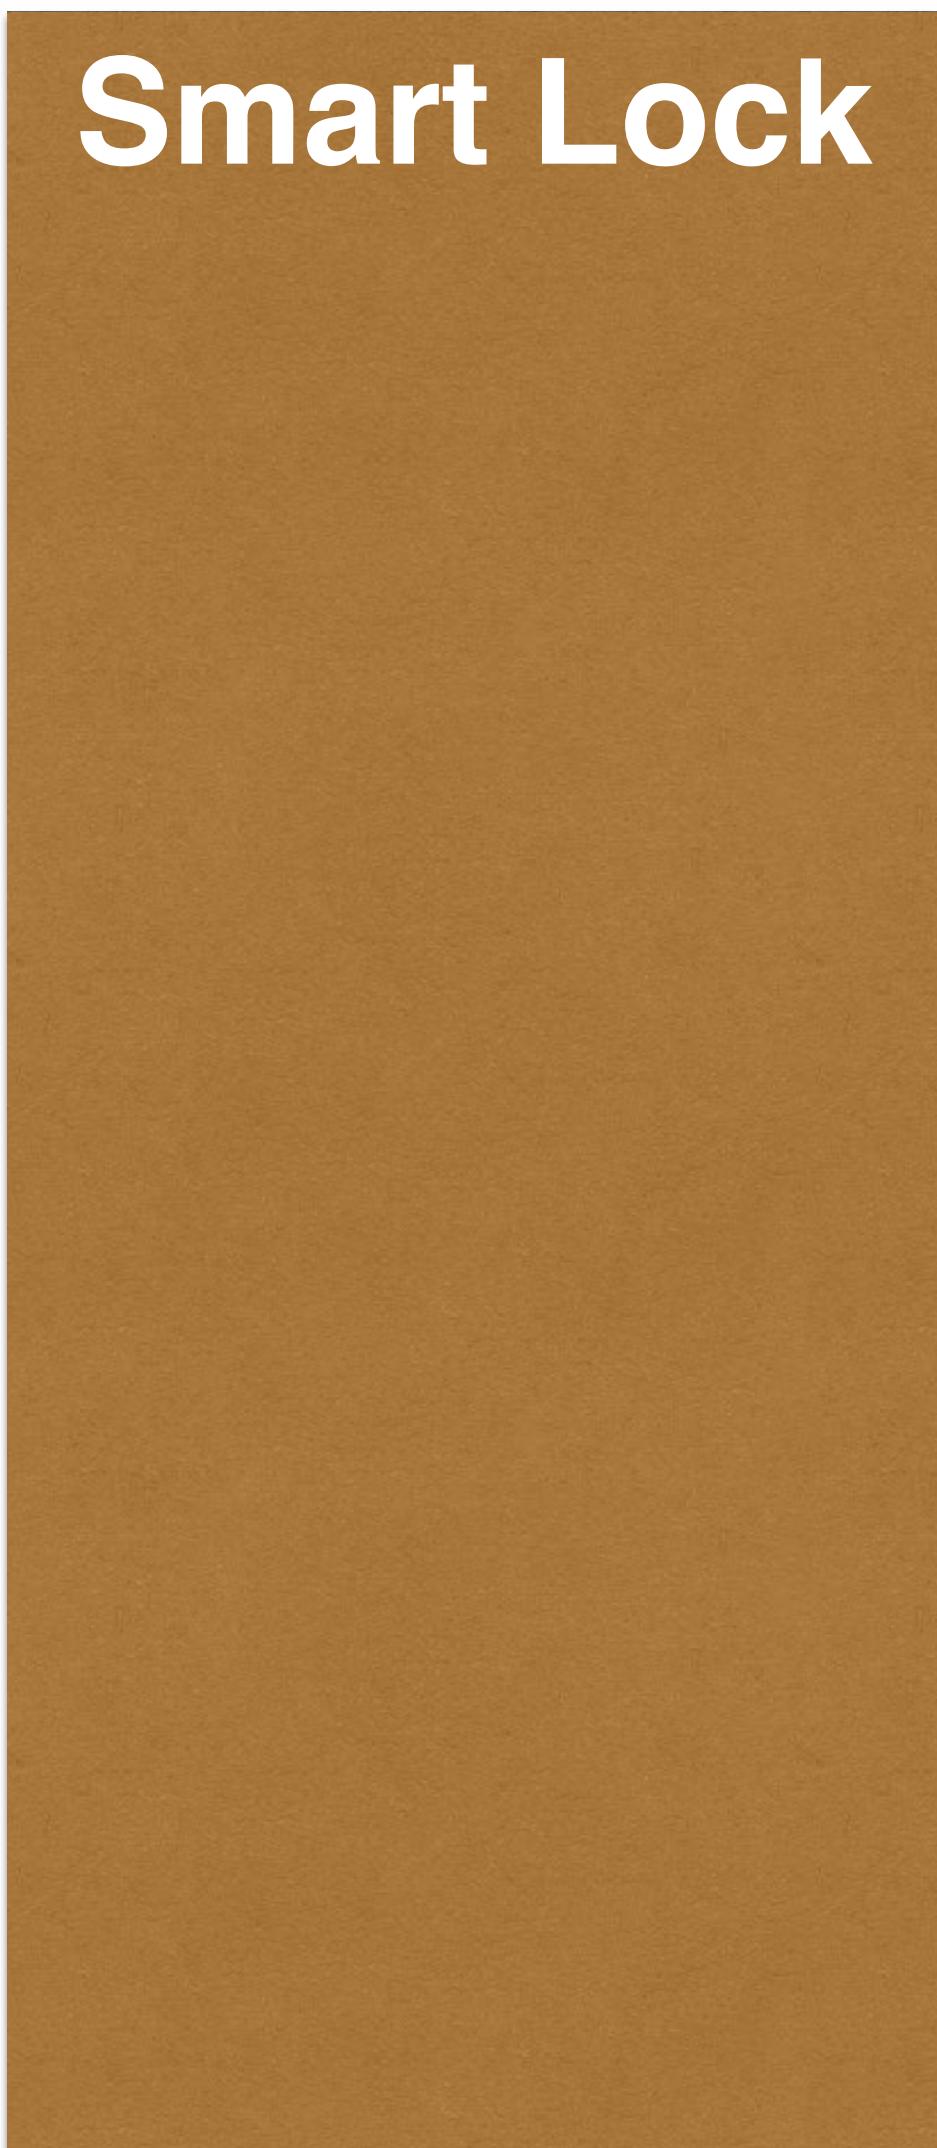
Heartbeat

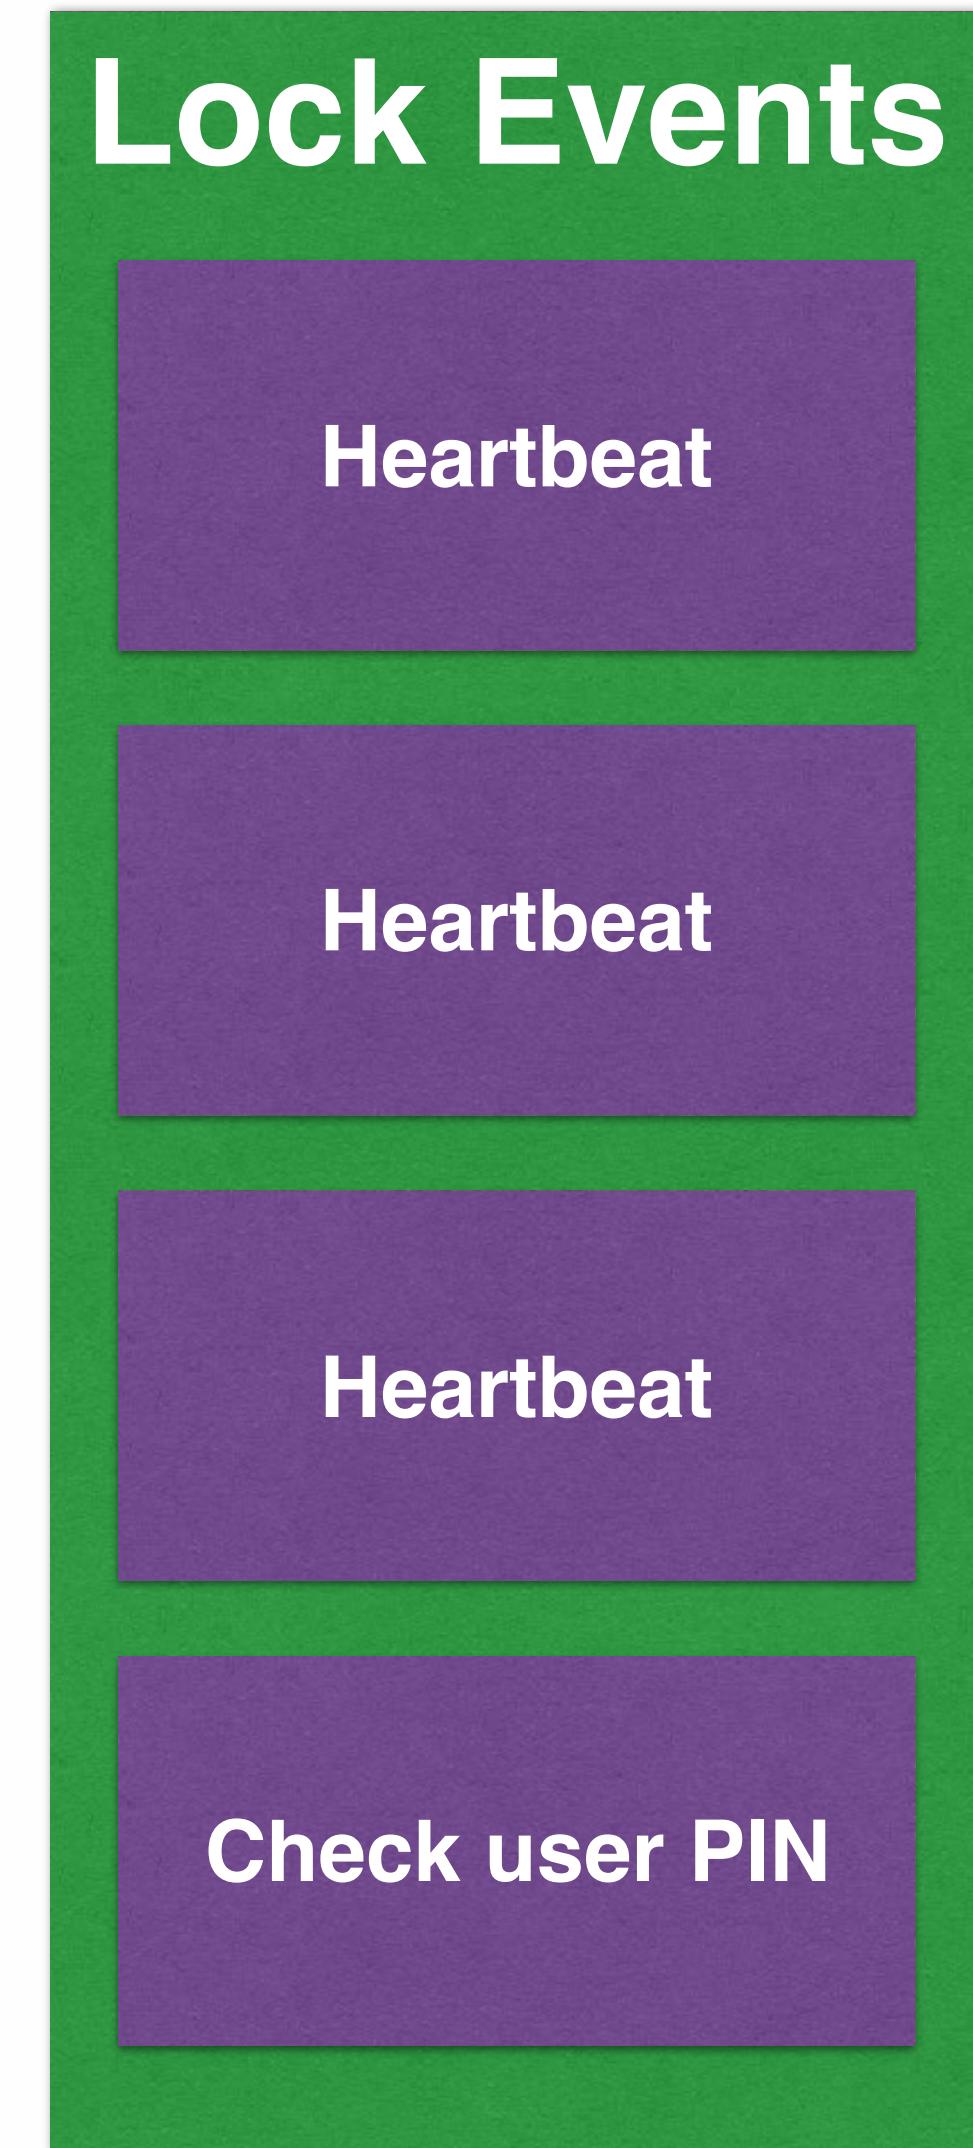
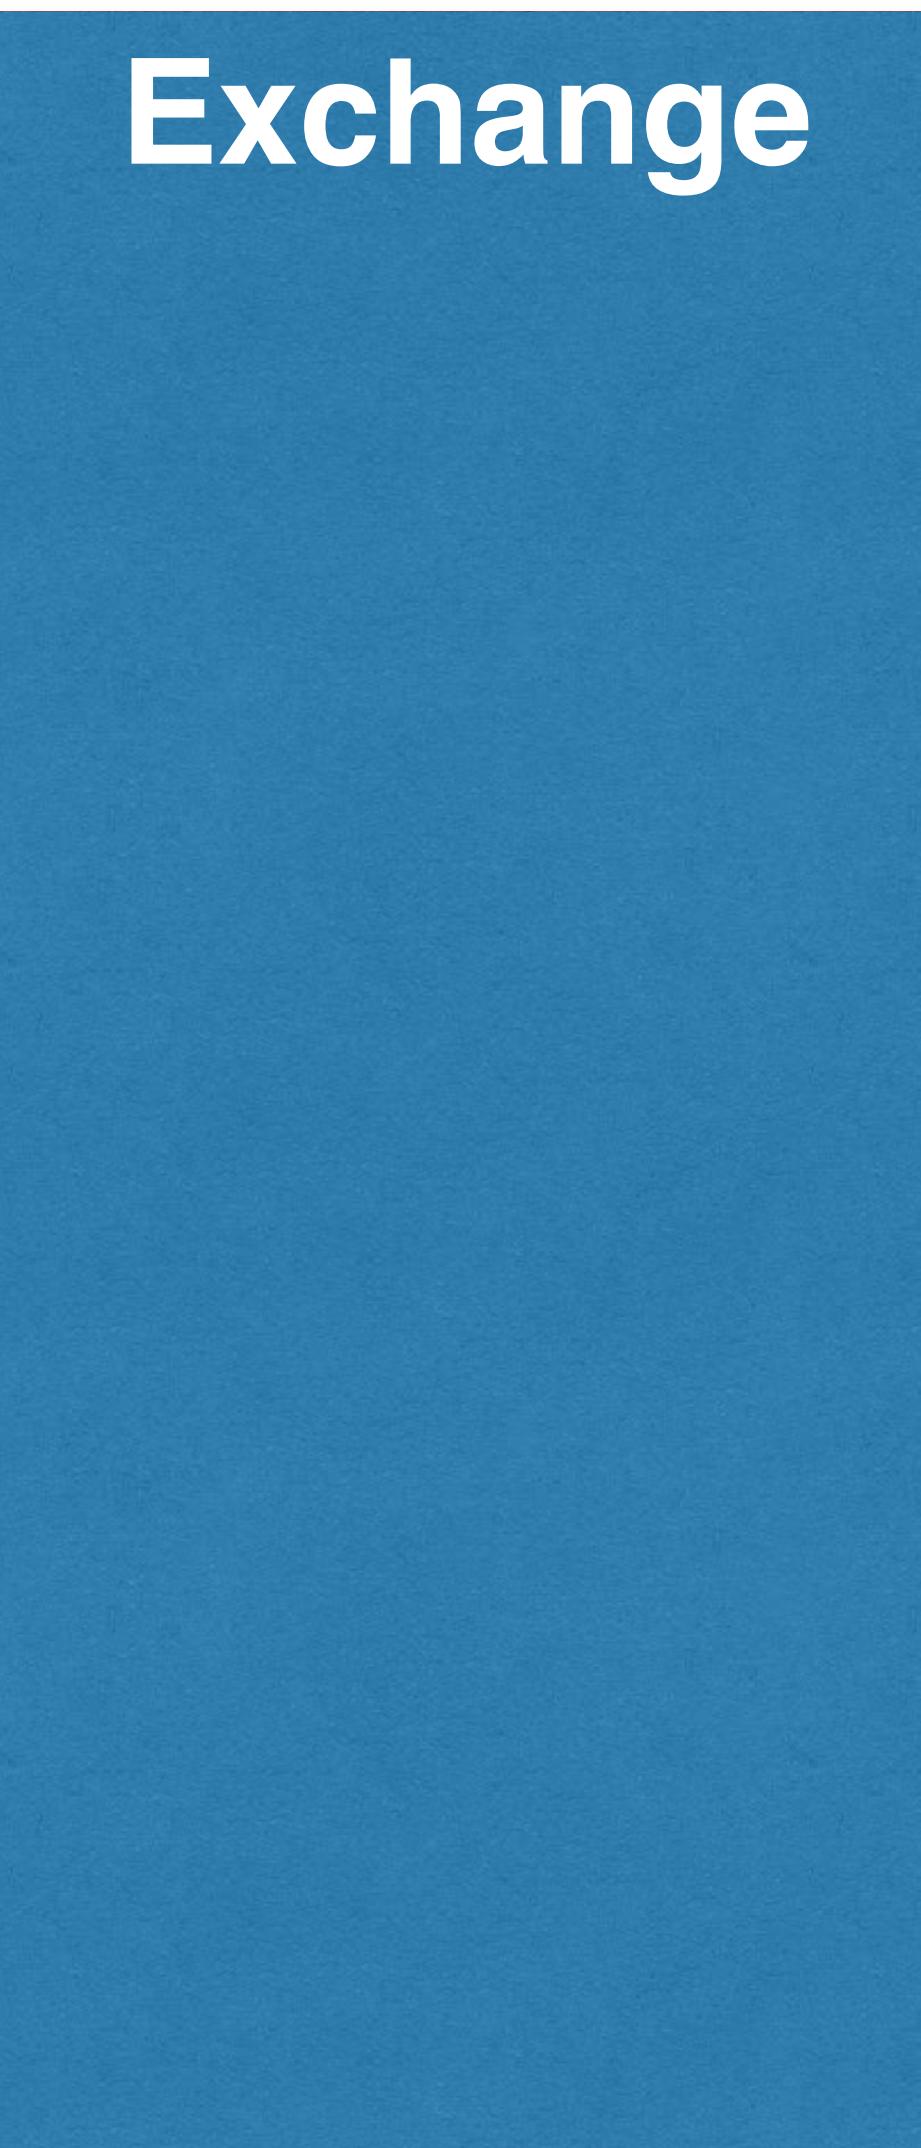
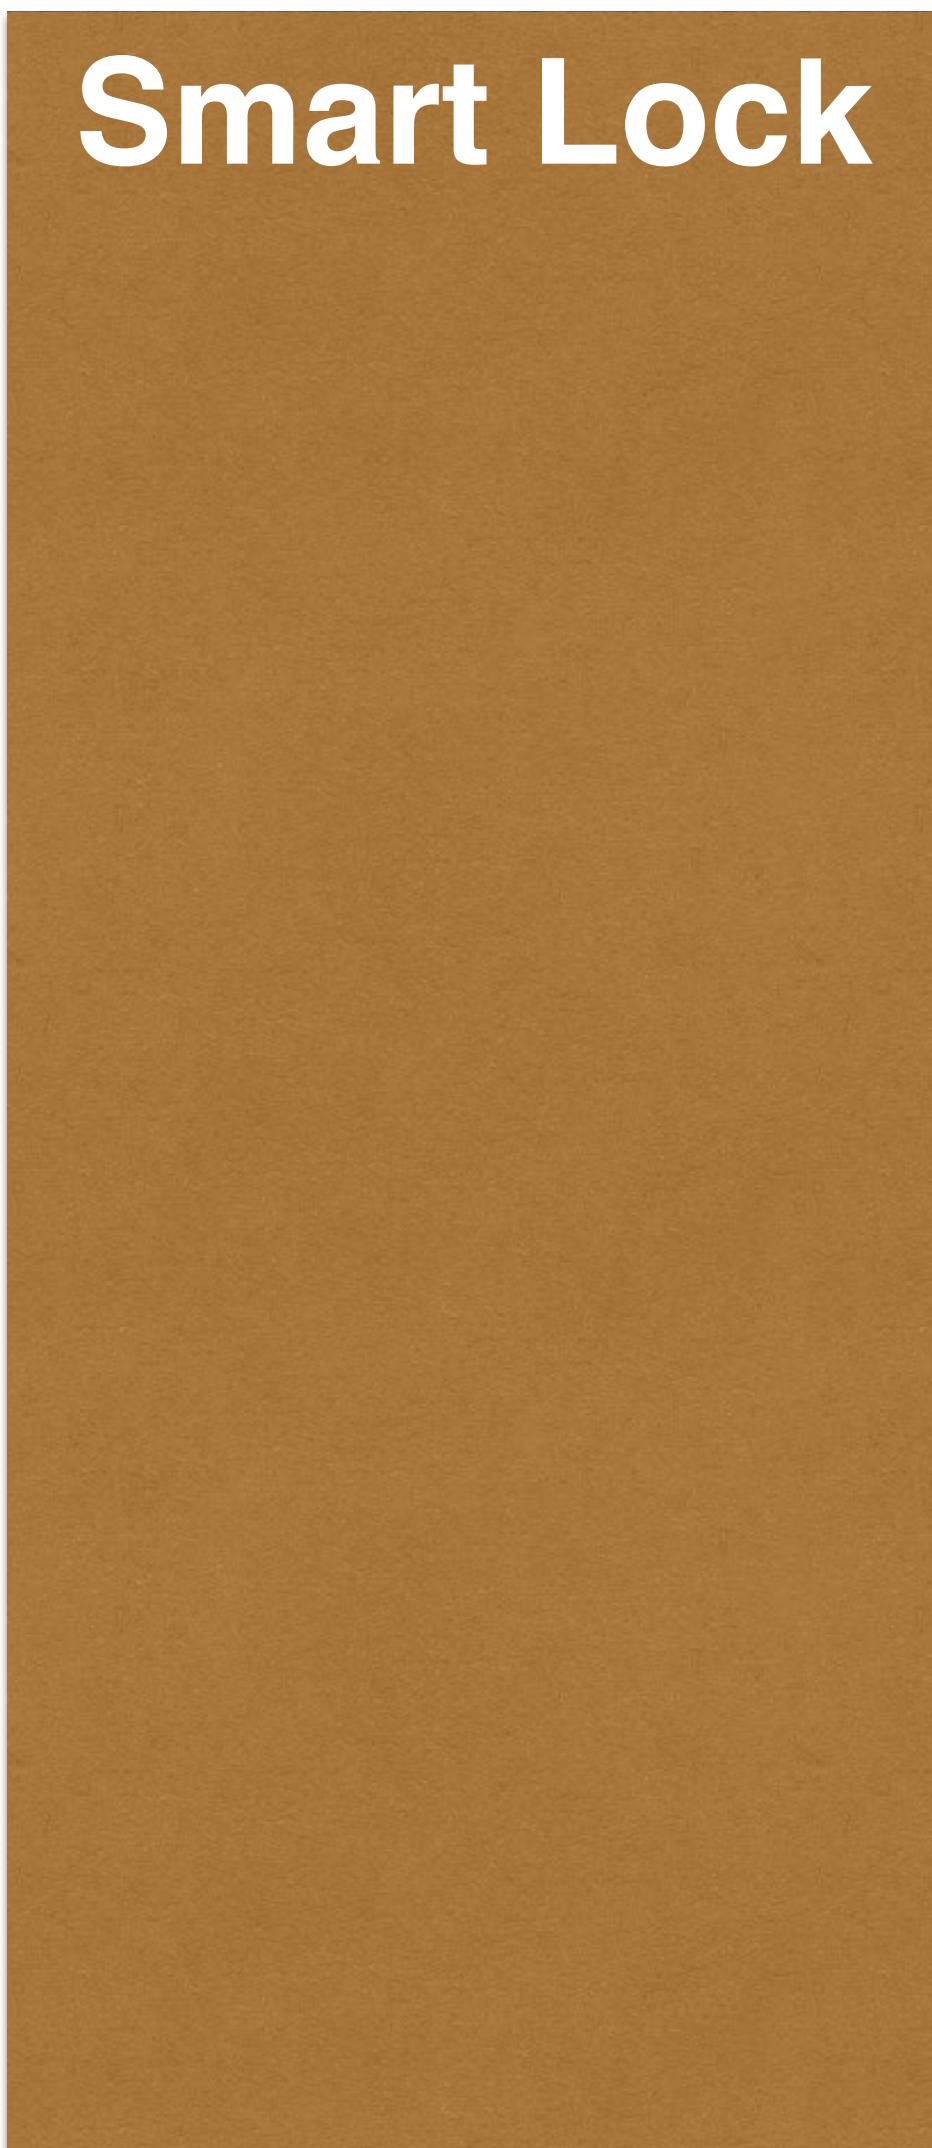
Heartbeat

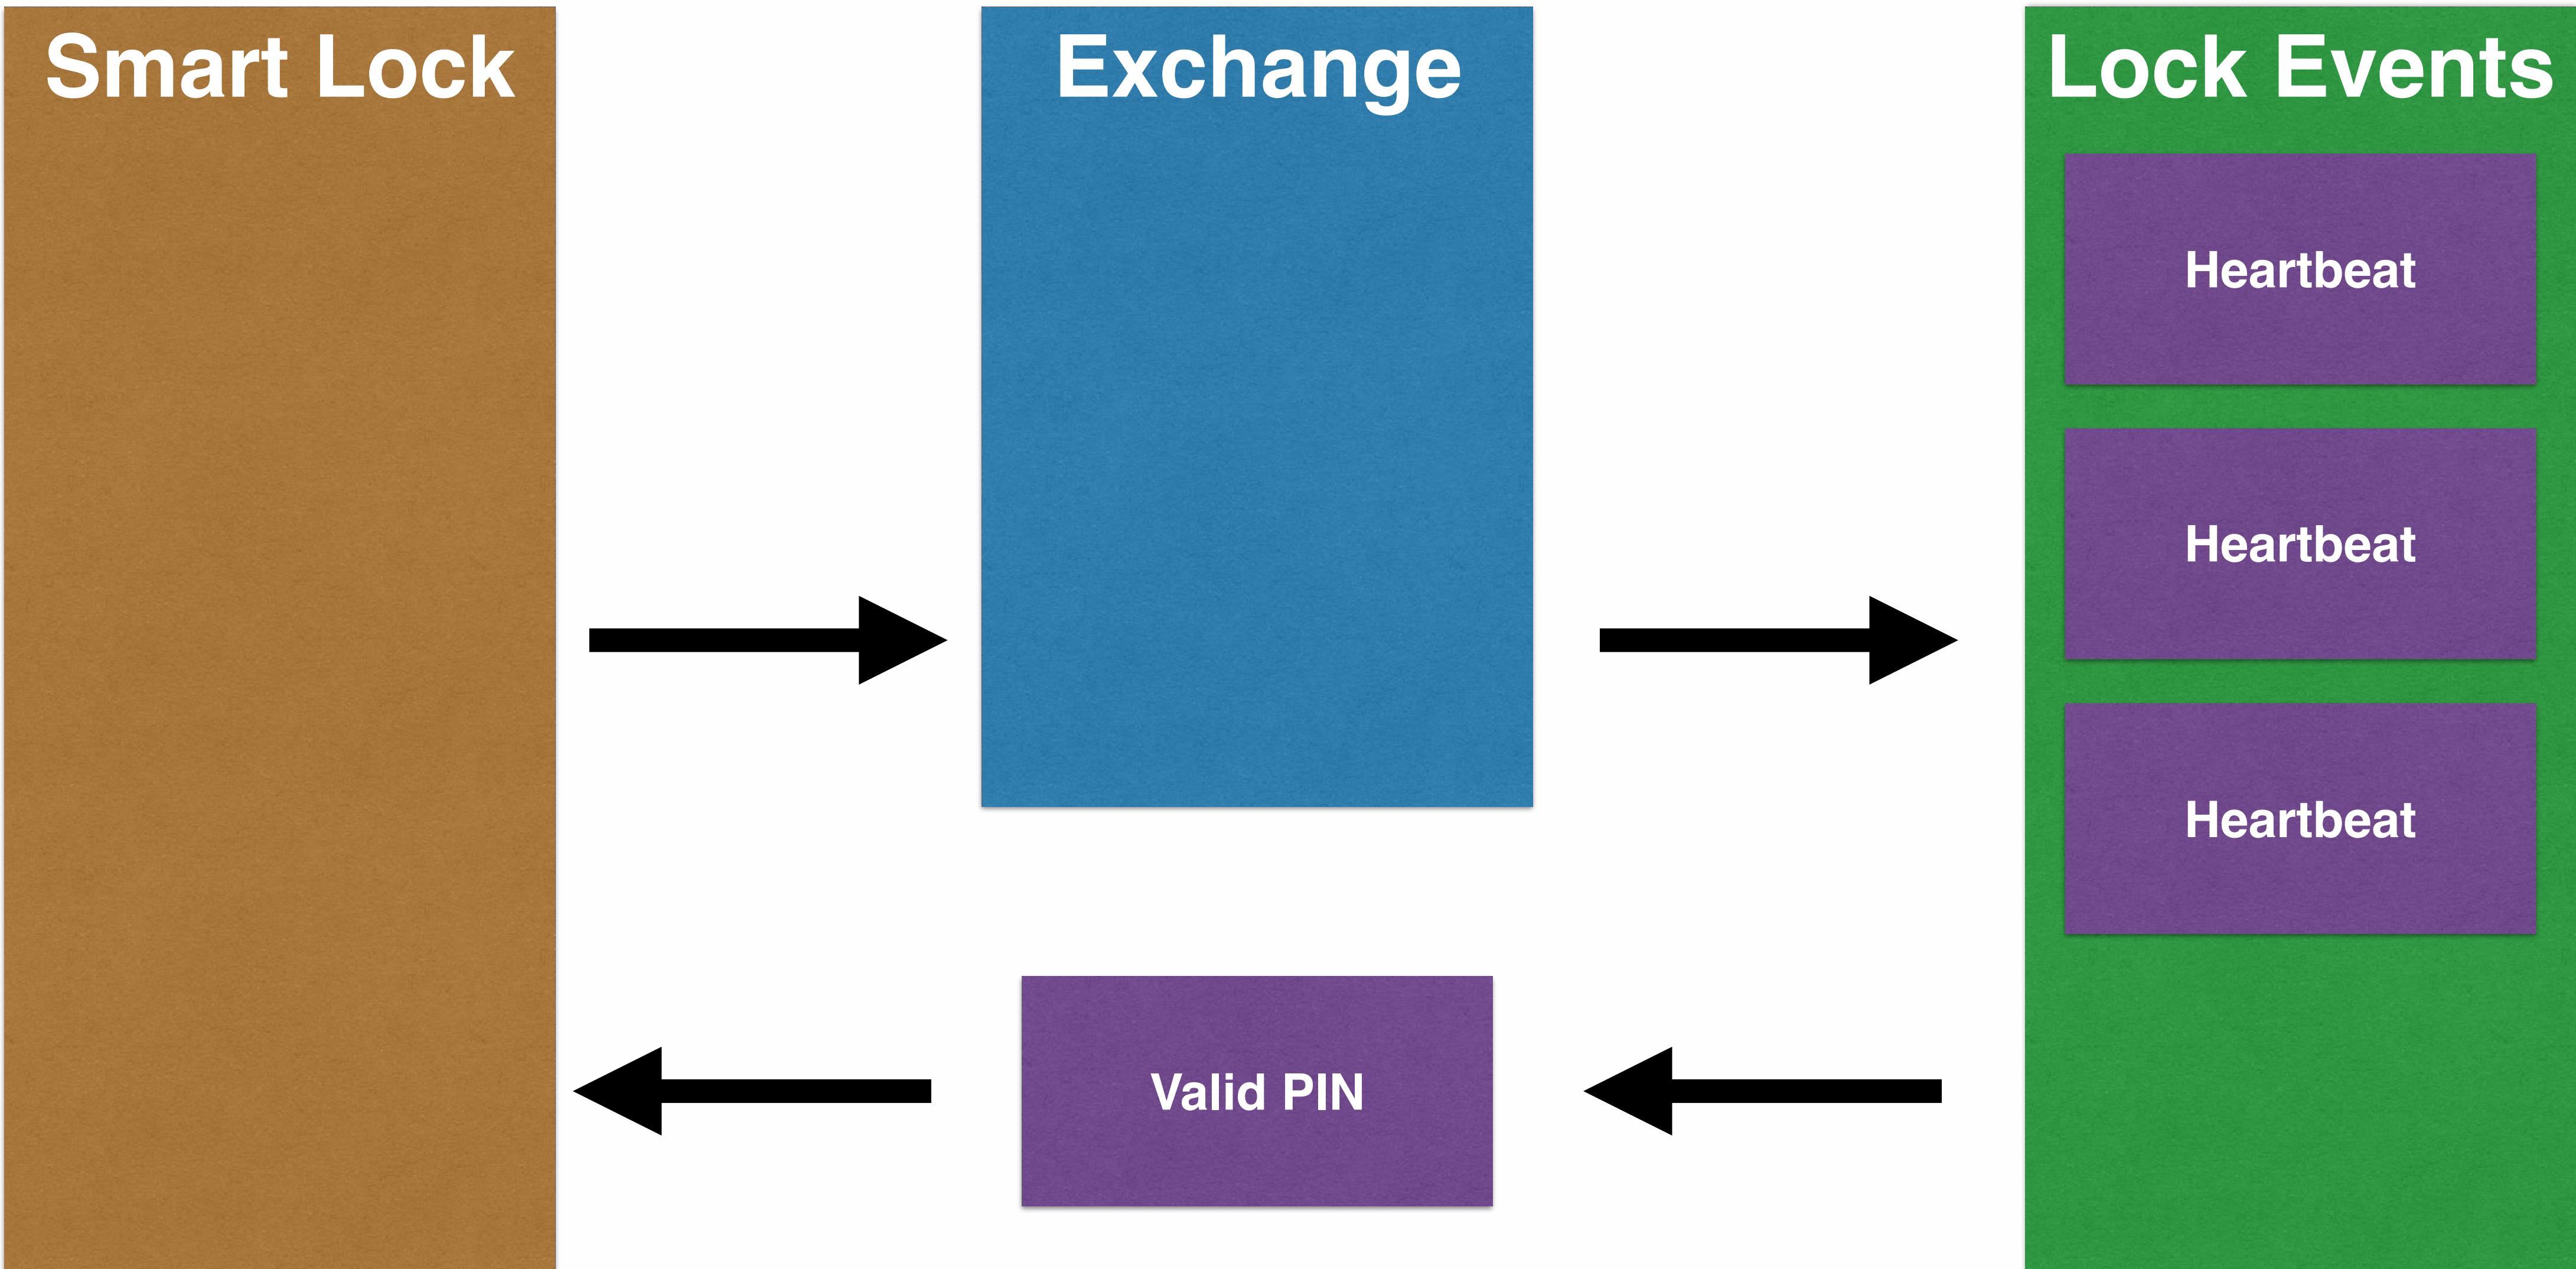


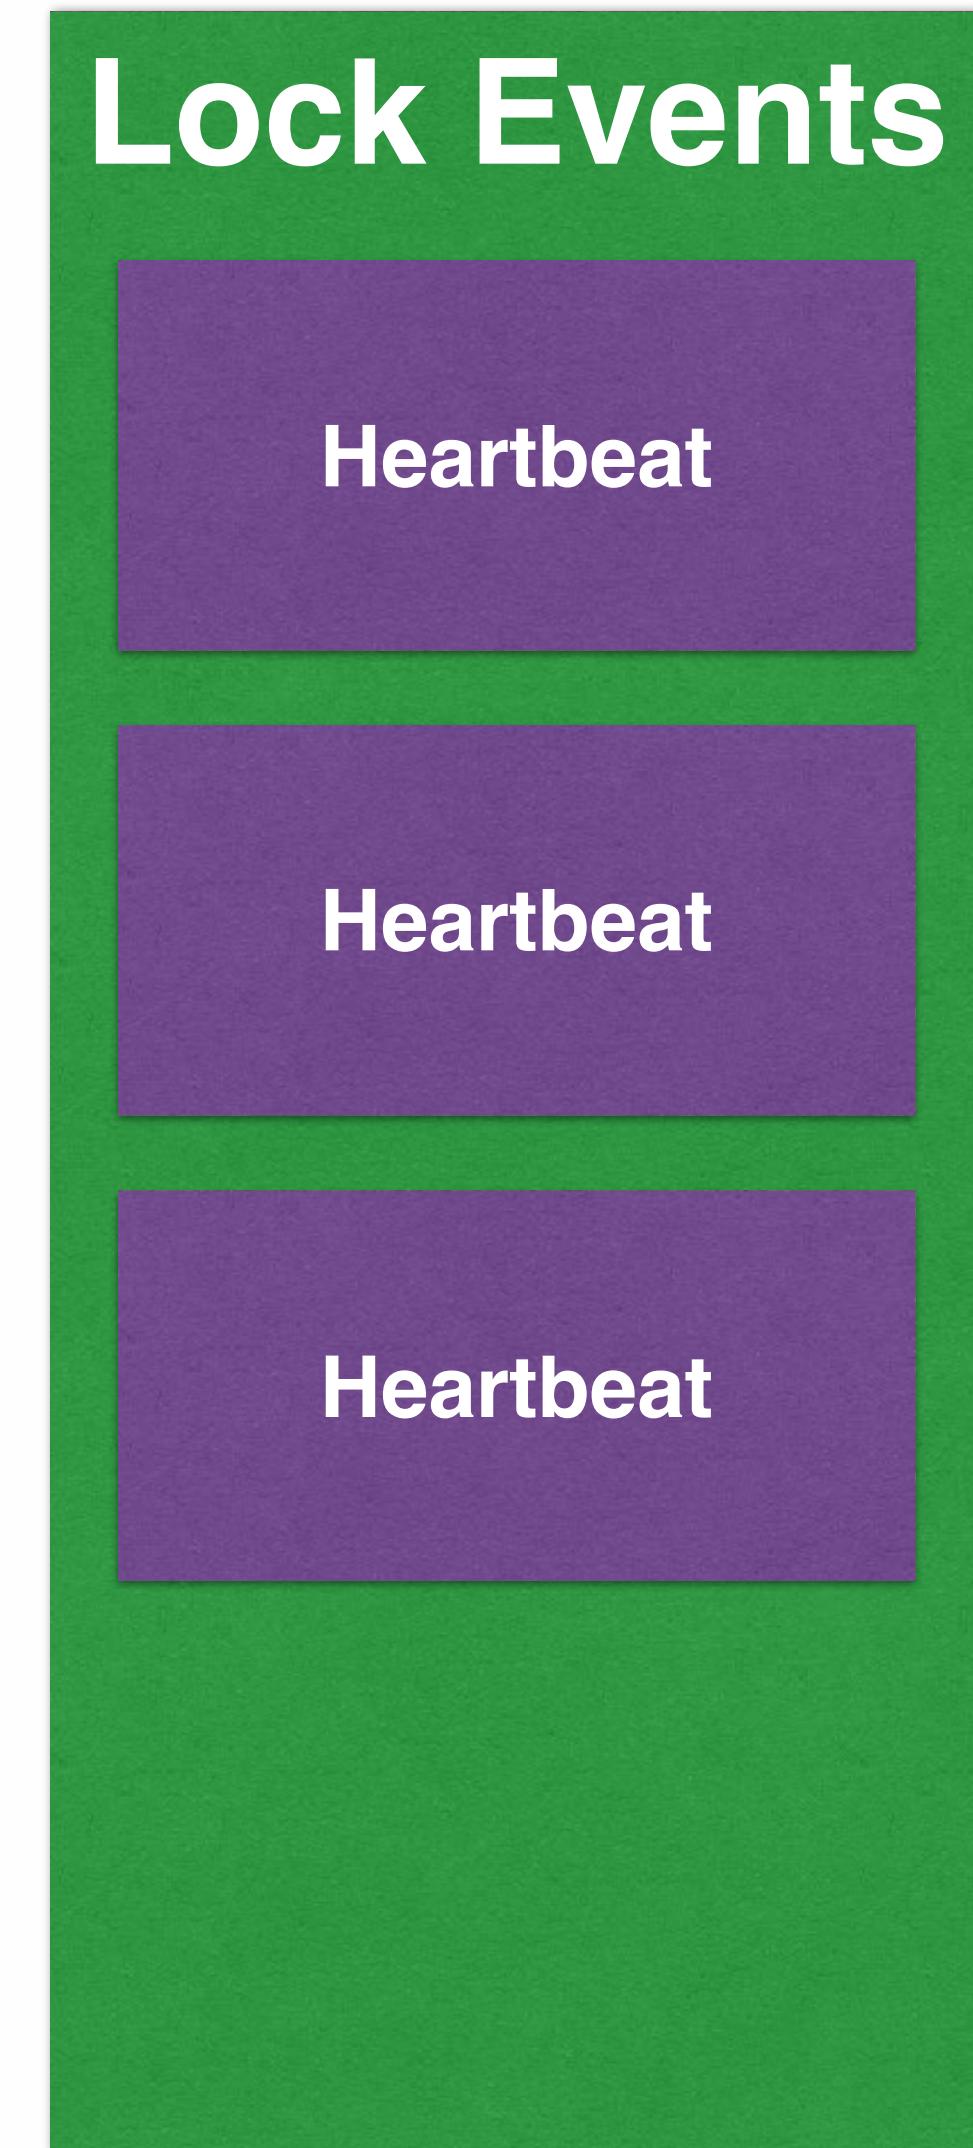
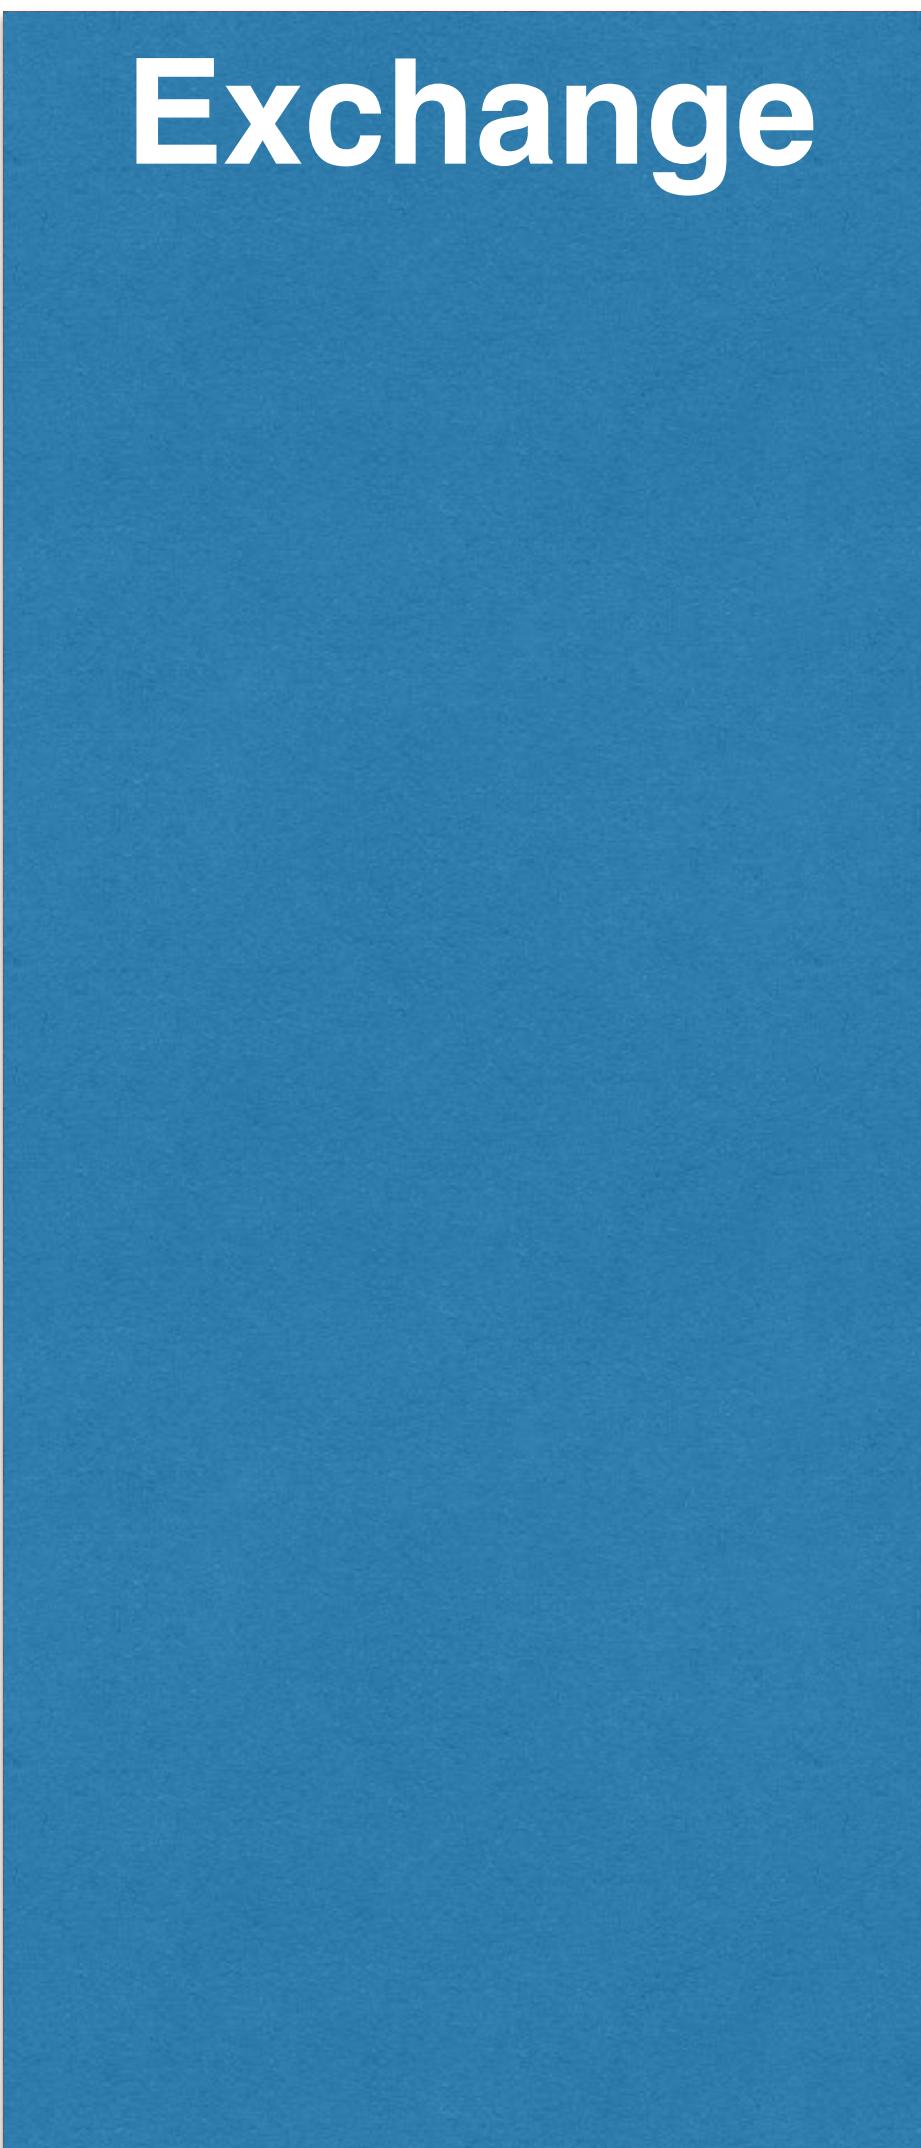












Dead lettering

Heartbeat

Heartbeat

Heartbeat

Heartbeat

Retry- Heartbeat

Heartbeat

Heartbeat

Heartbeat

Heartbeat

Retry- Heartbeat

Heartbeat

Heartbeat

Heartbeat



**Retry-
Heartbeat**

Heartbeat

Alternate exchanges

Lock Events Exchange

Check user PIN



Lock Events

Heartbeat

Heartbeat

Heartbeat

Lock Events Exchange

Check user PIN



Lock RPC Exchange

Lock

Unlock

Lock RPC Exchange

Check user PIN

Lock



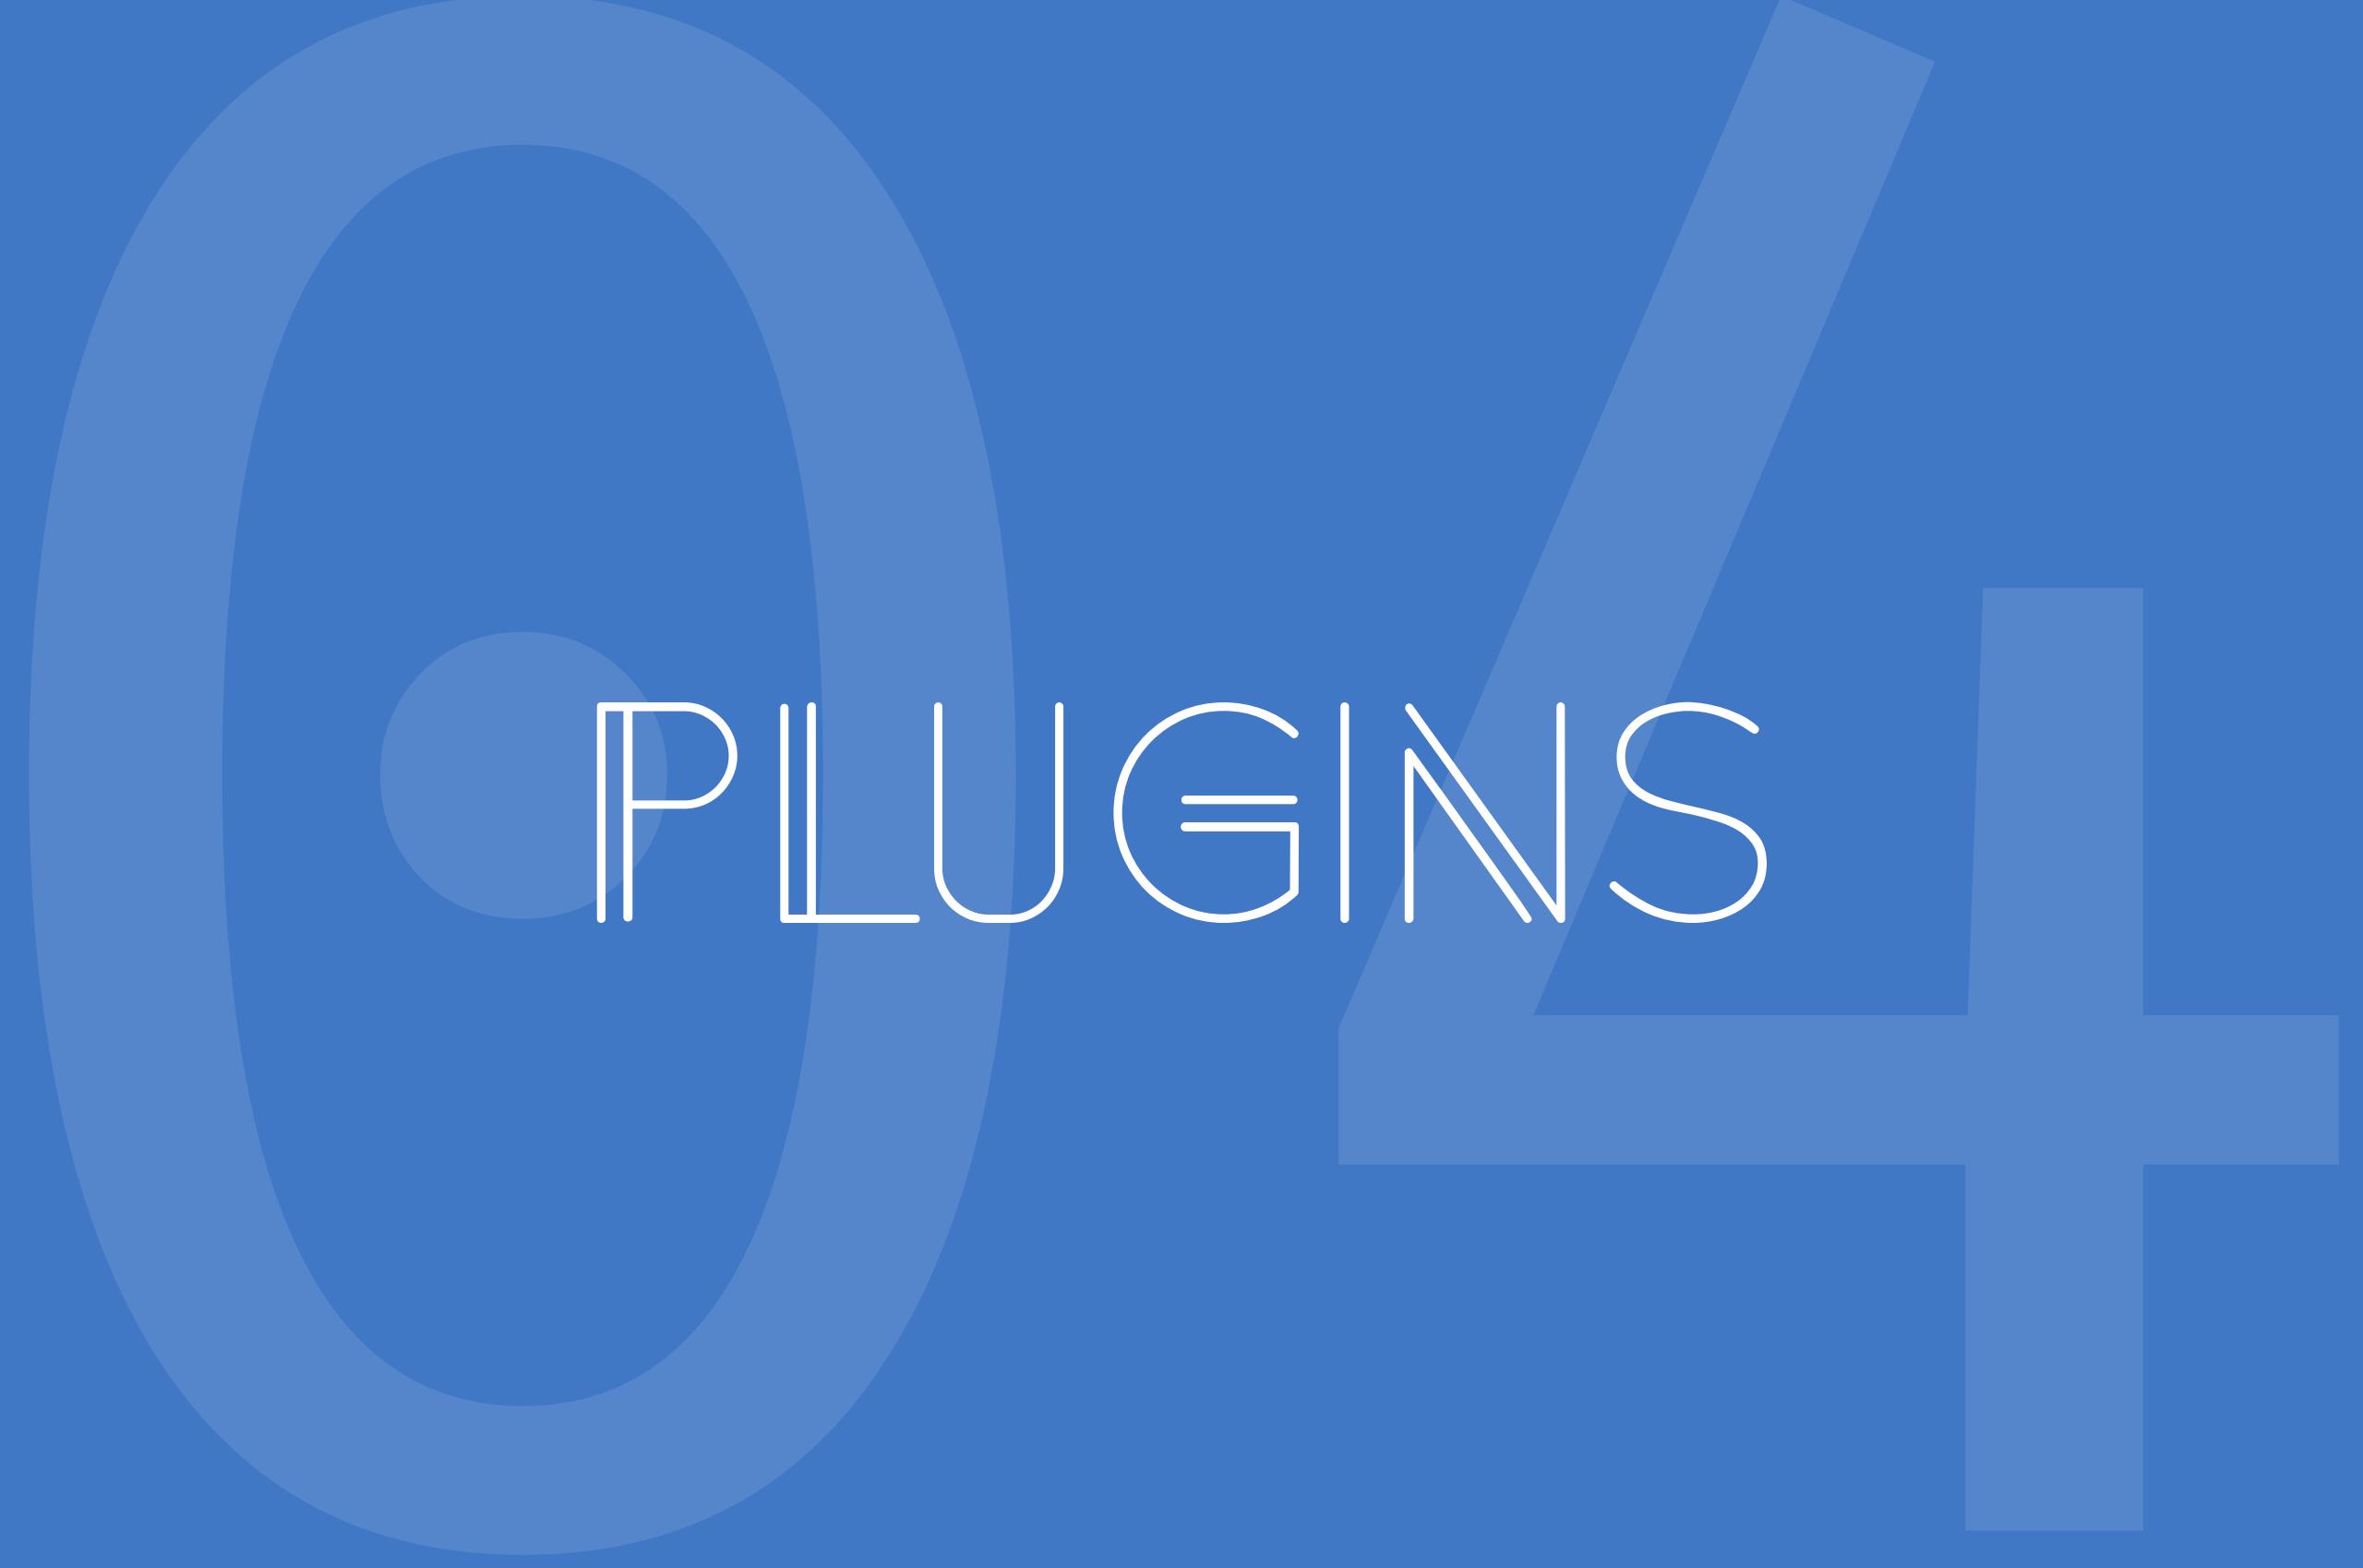
Lock RPC

Unlock

Priority consumers

Priority queues

TTL



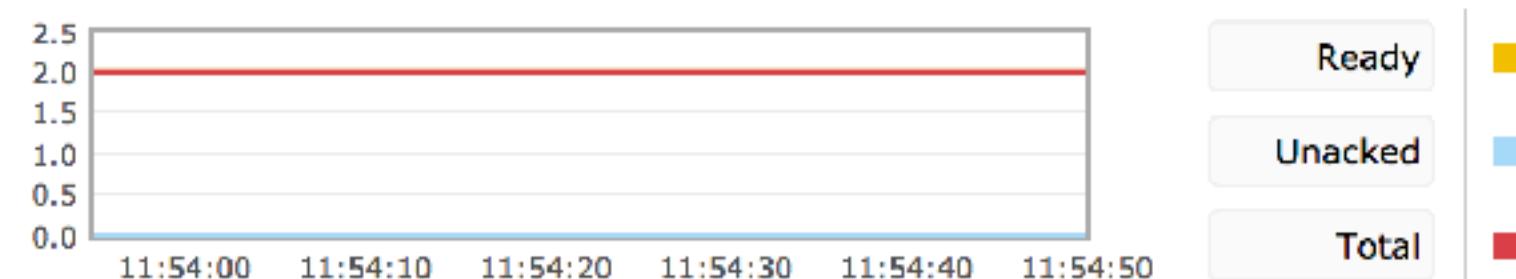
PLUGINS

Management

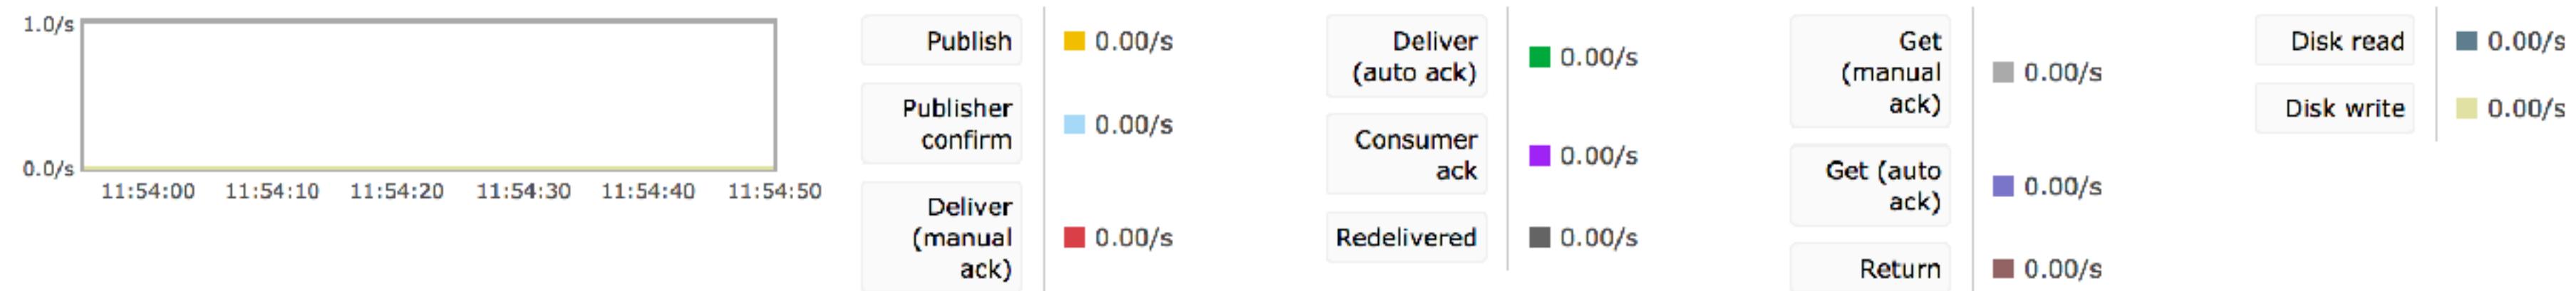
Overview

Totals

Queued messages (chart: last minute) (?)



Message rates (chart: last minute) (?)



Global counts (?)

Connections: 12

Channels: 12

Exchanges: 17

Queues: 8

Consumers: 12

Node

Node: rabbit@46e1483dee20 ([More about this node](#))

File descriptors (?)	Socket descriptors (?)	Erlang processes	Memory	Disk space	Rates mode	Info	Reset stats DB	+/-
39 1048576 available	12 943626 available	474 1048576 available	77MB 800MB high watermark	25GB 48MB low watermark	basic	Disc 1	Reset	

[Reset stats on all nodes](#)

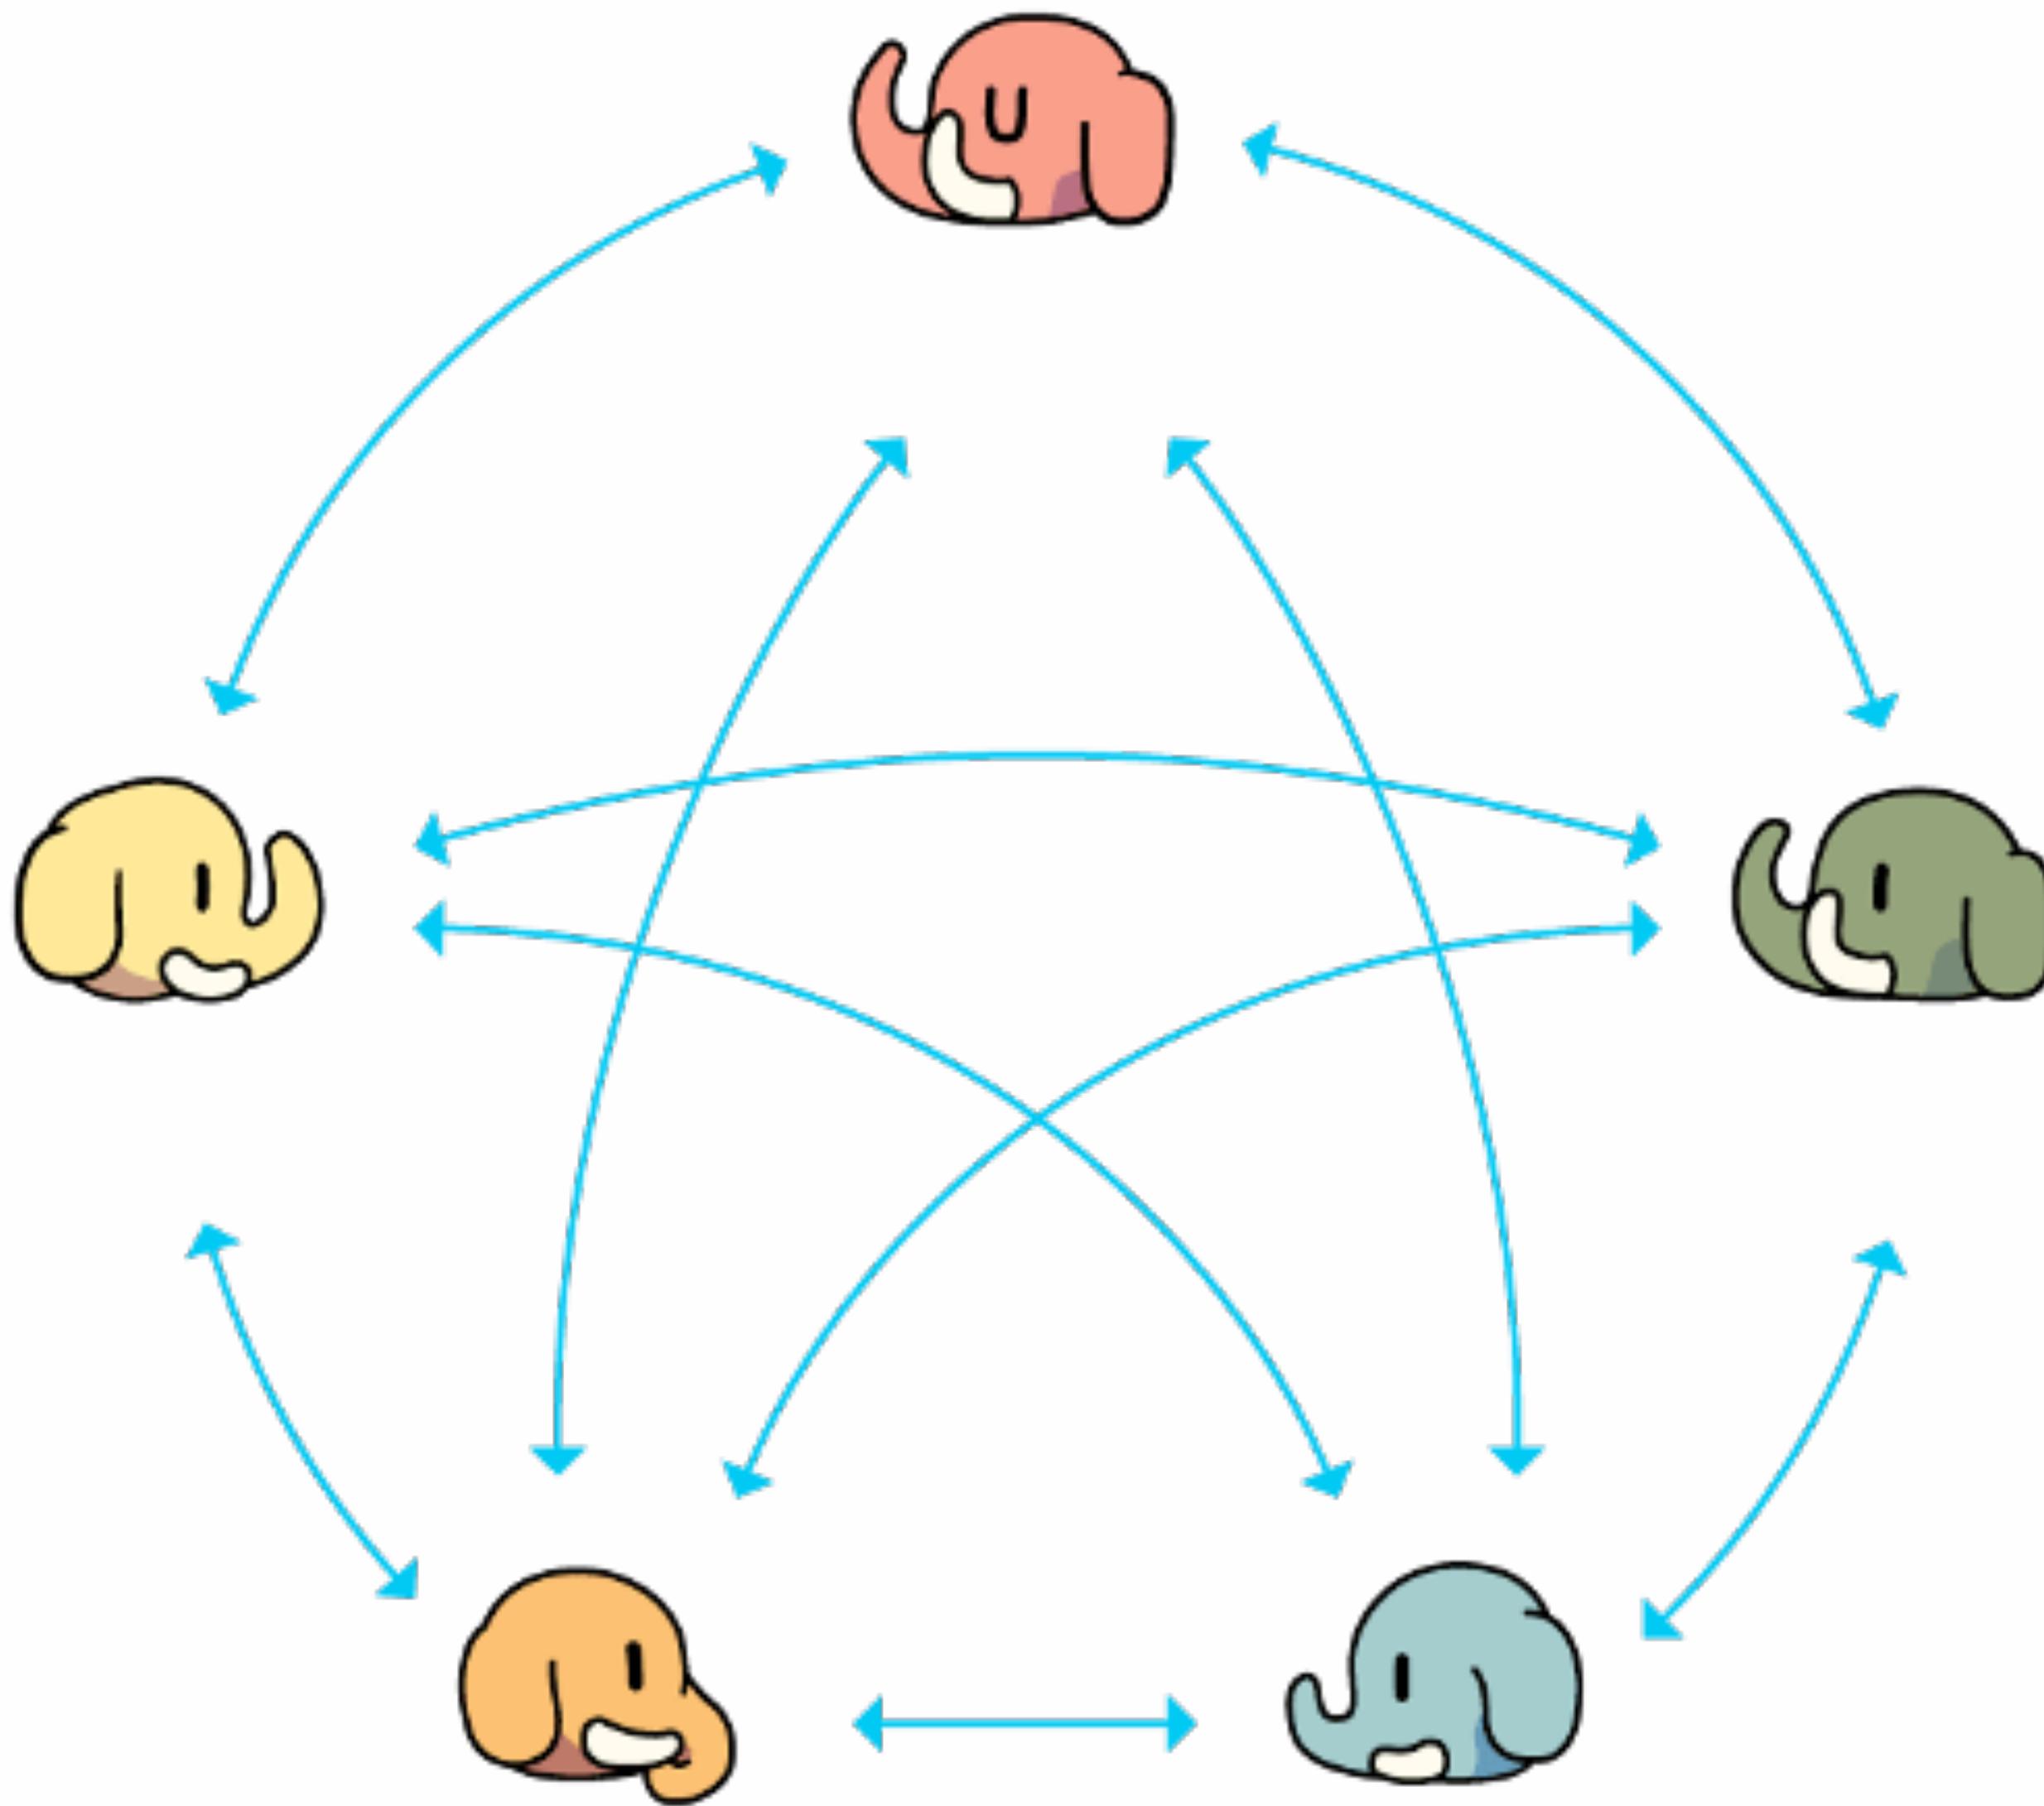
Paths

Config file	/etc/rabbitmq/rabbitmq.config
Database directory	/var/lib/rabbitmq/mnesia/rabbit@46e1483dee20
Log file	tty
SASL log file	tty

[▶ Ports and contexts](#)

MQTT, STOMP and WebSockets

Federation





CONCLUSION

Debugging

Load balancing / Scaling

Memory consumption

Should I use it on my next
project?

Monolith?
Probably not...

Services?
YES!

QUESTIONS



github.com/stankec



[@monorkin](https://twitter.com/@monorkin)



hi@stanko.io